


RQ1a: leakier splits should score higher. Only the bottom row exists.



FABRICATED PLACEHOLDER: NOT MEASURED

IN FLIGHT: the `--split_by random` and `--split_by recipe` training runs are on the cluster now (the design-disjoint run already exists). Identical encoder, budget and seed; only the split changes. When they land, run `src/test.py` for each, delete this block and rerun `analysis.make_all`.