

Training throughput across reduction methods (error bars: ± 1 s.d. of step time over graphs)
Training step time (dashed: full-graph baseline) Speedup under-delivers relative to compression

