


Full-graph inference throughput, GPU vs. CPU



No CPU evaluation pass survives in the exported results:
every inference_results CSV is device=cuda. Re-run
src/shell/test_cpu.sh to get the 'train cheap, infer full on
modest hardware' half of RQ5.