


Full-graph inference throughput, accelerator vs. CPU



No CPU evaluation pass survives in the exported results: every inference_results CSV is device=cuda. Re-run `src/shell/test_cpu.sh` to get the 'train cheap, infer full on modest hardware' half of RQ5.