


Cost of serving a query, by evaluation state

Inference throughput



Inference peak memory

