

MSC ARTIFICIAL INTELLIGENCE
MASTER THESIS

Reducing And-Inverter Graphs for Scalable Optimizability Prediction

by
Isabella V. GARDNER
314159265

2026/01/05–2026/08/31
36 ECTS

Supervisor: Michael COCHEZ

Examiner: Pascal METTES



UNIVERSITEIT VAN AMSTERDAM
Instituut voor Informatica

Reducing And-Inverter Graphs for Scalable Optimizability Prediction

This work by Isabella V. GARDNER is licensed under **CC-BY-NC**.

<https://creativecommons.org/licenses/by-nc/4.0>



This license requires that reusers give credit to the creator. It allows reusers to distribute, remix, adapt, and build upon the material in any medium or format, for noncommercial purposes only.

Abstract

Problem

Logic synthesis restructures a circuit’s And-Inverter Graph (AIG) representation for area, power, and delay, but assembling an effective sequence of transformations remains a heuristic trial-and-error process. Predicting a circuit’s optimizability beforehand, using a Graph Neural Network (GNN) trained directly on its structure, could replace that search with a data-driven choice. Industrial AIGs reach millions of nodes, however, exceeding the memory of a single device.

Gap

Graph reduction (partitioning, sparsification, and summarization) is the standard remedy for graphs that exceed device memory, well studied on citation and social networks but never systematically evaluated on AIGs. Cutting an edge there can sever a causal cone and destroy the structural signal a predictor depends on.

Approach

This thesis benchmarks generic and AIG-adapted reductions from all three families for optimizability regression, weighing predictive retention against efficiency relative to an unreduced baseline. One summarization variant merges only nodes a bounded-depth encoder cannot distinguish, a coarsening with zero accuracy cost by construction. A further test asks whether a model trained solely on reduced graphs still predicts accurately on full, unreduced circuits.

Findings

[TODO: State the RQ2/RQ3 headline trade-off (memory reduction at matched predictive retention) and the RQ5 transfer result once measured.]

Keywords: electronic design automation; logic synthesis; graph neural network; And-Inverter Graph; graph reduction; graph coarsening

Contents

List of Figures	viii
List of Tables	x
List of Symbols	xi
List of Abbreviations	xiii
1 Introduction	1
1.1 Motivation	1
1.1.1 Logic Synthesis and Optimization Script Design	1
1.1.2 Machine Learning for Optimizability Prediction	1
1.1.3 The Scale Bottleneck	1
1.1.4 Limitations of Generic Graph Reduction	1
1.1.5 Problem Statement and Research Gap	1
1.2 Research Questions	1
1.2.1 RQ1: Task Feasibility and Baseline	2
1.2.2 RQ2: Reduction Efficiency	2
1.2.3 RQ3: Predictive Retention and Trade-offs	2
1.2.4 RQ4: Value of Domain-Informed Adaptation	2
1.2.5 RQ5: Cross-State Inference	2
1.2.6 Scope and Delimitations	2
1.3 Contributions	3
1.3.1 Summary of Results	3
1.4 Thesis Outline	3
2 Preliminaries & Related Work	5
2.1 Preliminaries	5
2.1.1 Electronic Design Automation and Logic Synthesis	5
2.1.2 And-Inverter Graphs	5
2.1.3 Synthesis Algorithms	7
2.1.4 Graph Neural Networks	7
2.1.5 Graph Reduction: Definitions	8
2.1.6 Problem Formalization	9
2.2 Related Work	10
2.2.1 Machine Learning for Logic Synthesis	10
2.2.2 Scaling Graph Neural Networks to Large Graphs	11
2.2.3 Graph Reduction Techniques	11
2.2.4 Structure-Preserving and Domain-Aware Reduction	13
2.2.5 Size Generalization and Cross-State Transfer	13
2.2.6 Research Gap	13
3 Methodology	15
3.1 Architecture	15
3.1.1 Encoder Formulation	15
3.1.2 Partitioned Inputs	16
3.1.3 Sparsified Inputs	16
3.1.4 Summarized Inputs	16
3.1.5 Encoder Variant for Exact Colour Refinement	16
3.1.6 Baselines for Optimizability Prediction	17
3.2 Graph Reduction	18
3.2.1 Partitioning	18
3.2.2 Sparsification	19

3.2.3	Summarization and Coarsening	20
3.3	Data	21
3.3.1	Source Circuits	22
3.3.2	Dataset Generation	23
3.3.3	Tiered Dataset Structure	23
3.3.4	Target Label: Optimizability	25
3.3.5	Graph Representation	25
3.3.6	Dataset Analysis	25
3.3.7	Held-Out Hyperparameter Tuning Subset	35
3.4	Experiments	35
3.4.1	Research Design	35
3.4.2	Setup	35
3.4.3	Evaluation Metrics	37
3.4.4	Reproducibility	39
4	Results	41
4.1	Baseline Predictive Performance on Full AIGs (RQ1)	41
4.1.1	Comparison Against Naive Baselines	41
4.1.2	Predictive Accuracy Metrics	43
4.1.3	Ranking Efficacy	45
4.1.4	Baseline Model Outcomes	46
4.1.5	Error Analysis	48
4.1.6	Protocol Sensitivity: How Much of the Score Is Design Recognition? (RQ1a)	52
4.1.7	Cost of the Full-Graph Baseline	54
4.2	Efficiency Profile of Graph Reduction Techniques (RQ2)	57
4.2.1	Graph Reduction Offline Profile	57
4.2.2	Training Memory Dynamics and Compute Efficiency	60
4.2.3	Throughput and Training Speedup	62
4.2.4	Efficiency Ranking Across Methods	65
4.3	Predictive Retention and Performance Trade-offs (RQ3)	67
4.3.1	Accuracy Degradation	67
4.3.2	Ranking Preservation	72
4.3.3	Accuracy Loss Attributable to Reduction	75
4.3.4	The Pareto Front (Trade-Off Analysis)	78
4.4	Value of Domain-Informed Adaptation (RQ4)	80
4.4.1	Matched-Compression Pairings	80
4.4.2	Retention Gap at Matched Compression	80
4.4.3	Structural Interpretation	80
4.4.4	Adequacy of the Heuristics Tested	80
4.5	Cross-Structural Generalization (RQ5)	82
4.5.1	Reduced-to-Full Inference Accuracy	82
4.5.2	Zero-Shot Ranking	82
4.5.3	Matched-State vs. Cross-State Drop-Off	82
4.5.4	Generalization by Reduction Family	82
4.5.5	Inference Cost	82
4.6	Summary of Findings	88
4.6.1	Consolidated Summary	91
5	Discussion	93
5.1	Comparison with Related Work	93
5.1.1	Positioning Against Optimizability Prediction Work	93
5.1.2	Positioning Against Graph Reduction Work	93
5.1.3	Where Direct Comparison Is Not Possible	93
5.2	Interpretation of Findings	93
5.2.1	Feasibility of Optimizability Prediction	93
5.2.2	What Reduction Buys and What It Costs	94
5.2.3	Why Domain-Informed Adaptation Did or Did Not Help	94
5.2.4	Cross-State Transfer and Deployment	94
5.2.5	Unexpected Results	94
5.3	Practical Implications	94
5.4	Limitations	94

5.4.1	Reproducibility	94
5.4.2	Scalability	94
5.4.3	Generalizability	94
5.4.4	Reliability and Validity	95
5.4.5	Limitations of the Reduction Methods Tested	95
5.5	Outlook	95
5.6	Ethical Considerations	95
6	Conclusion	97
6.1	Answers to the Research Questions	97
6.1.1	RQ1: Task Feasibility and Baseline	97
6.1.2	RQ2: Reduction Efficiency	97
6.1.3	RQ3: Predictive Retention and Trade-offs	97
6.1.4	RQ4: Value of Domain-Informed Adaptation	97
6.1.5	RQ5: Cross-State Transfer	97
6.2	Contributions Revisited	97
6.3	Future Work	97
6.3.1	Extending Beyond a Single Synthesis Algorithm	97
6.3.2	Stronger Domain-Aware Reduction	97
6.3.3	Scaling to Industrial AIGs	97
6.3.4	From Prediction to Script Generation	97
A	First Appendix	103
A.1	Synthesis Algorithm Configurations	103
A.2	Reduction Method Parameters	103
A.3	Baseline Hyperparameters	103
A.4	Experimental Configuration	103
A.5	Extended Results	103

List of Figures

2.1	An example And-Inverter Graph	6
3.1	Graph scale distribution	27
3.2	Structural corpus statistics	27
3.3	Optimizability label distribution	28
3.4	Zero inflation by design and scale	28
3.5	Graph size against label	29
3.6	Label by source synthesis script	30
3.7	Test-set composition by design	31
3.8	Validation against test distribution	31
4.1	Baseline group comparison	42
4.2	Full-graph parity plot	43
4.3	Baseline calibration	43
4.4	Baseline performance by optimizability band	44
4.5	Variance decomposition of the baseline score	44
4.6	Ranking efficacy of the full-graph baseline	45
4.7	Residuals by circuit scale	49
4.8	Error by optimization tier	50
4.9	Per-design error on unseen designs	51
4.10	Split protocol sensitivity	53
4.11	Training trajectories	55
4.12	Training budget consumed	55
4.13	Hyperparameter sweep	56
4.14	Hyperparameter sensitivity	56
4.15	Node and edge retention	58
4.16	Retention distributions	58
4.17	Offline cost per method	59
4.18	Amortisation threshold	59
4.19	Partitioner trade-off	59
4.20	Residual redundancy after strashing	59
4.21	Peak training memory	61
4.22	Cost profile against graph size	61
4.23	Training throughput	63
4.24	Paired per-graph savings	63
4.25	Matched-state accuracy	68
4.26	RMSE against explained variance	69
4.27	Matched-state accuracy by tier	70
4.28	Matched-state accuracy by source script	71
4.29	Matched-state ranking under label strata	73
4.30	Cross-state ranking under label strata	74
4.31	Accuracy against achieved compression	76
4.32	Accuracy retained under reduction	76
4.33	All configurations by optimizability band	77
4.34	Effective receptive field	77
4.35	H1: receptive field against accuracy	77
4.36	Pareto front	79
4.37	Validation-to-test gap	79
4.38	Paired domain-vs-generic gaps	81
4.39	Pairing quality	81
4.40	Gaps against run-to-run noise	81
4.41	Unpaired domain-vs-generic view	83
4.42	Matched-state against full-graph evaluation	84

4.43	Cost of the structural shift	84
4.44	Transfer quadrant	85
4.45	Inference cost by evaluation state	86
4.46	GPU against CPU inference	87
4.47	RQ5 positive control	87
4.48	Summarization landscape	89
4.49	Residual redundancy after strashing	90
4.50	Effective receptive field	91
4.51	H1: receptive field against accuracy	91

List of Tables

3.1	Encoder configuration	15
3.2	The summarization candidates, the equivalence each merges by, and whether its compression can be set. Only a method taking a target ratio can be matched to another directly; the rest are matched through the random control of 3.2.3.	20
3.3	Where the 55 source designs come from, and under what terms. The license column gives each suite’s own terms; the ISCAS-85 and MCNC circuits predate modern license practice and carry no formal terms. The terms governing the copies actually used are stated in the text.	22
3.4	Source circuit characteristics	24
3.5	Held-out test corpus. The label is concentrated near zero, which is the denominator every R^2 in this chapter is measured against.	32
3.6	Composition of the evaluation corpus over the two stored tiers. Validation and test differ substantially in label and size distribution, which is expected under a design-level split and is the direct explanation for the validation-to-test gap.	32
3.7	Structural corpus statistics. The AND-gate share is recovered from the AND-gate-only node masks over 10,000 graphs: that method drops exactly the primary inputs and outputs, so its node retention is the AND-and-constant share. This is what fixes that method’s compression, rather than a parameter that could be dialled.	33
3.8	Whole-corpus statistics by tier.	34
3.9	The four experimental phases, what each answers, and what each costs. Only Phases 1 and 3 consume training compute.	35
3.10	The three train/test protocols implemented in this work, ordered from leakiest to strictest. Only the design-disjoint protocol is used for the reported results; the other two are run once each at the same configuration to quantify how much of a reported score is design recognition rather than structural reading.	36
4.1	RQ1 baseline comparison on the design-disjoint test split. Trivial predictors are fitted on validation and scored on test. Upstream R^2 figures are NOT comparable: OpenABC-D z-scores its targets per design, which removes between-design variance from the denominator. Only the sign transfers.	47
4.2	Per-design metrics on the design-disjoint test set. A pooled score over a design-level split is an average over this handful of whole circuits, not over a random sample.	49
4.3	RQ1a protocol sensitivity. Identical encoder, budget and seed; only the split changes. Random is the common default in circuit representation-learning evaluations, recipe-disjoint is OpenABC-D Variant 1, design-disjoint is Variant 2 and is what this thesis reports.	52
4.4	Final configuration, as recorded in the run log of the reported run. Fields the run did not log are marked rather than omitted.	54
4.5	The four partitioners keep every node and differ only in which edges they cut, at the same k . That makes them the cleanest domain-informed / generic pairs in the thesis.	59
4.6	Per-graph savings against the full-graph baseline, paired by graph id, with percentile bootstrap intervals and a Wilcoxon signed-rank test on the paired deltas. Savings are not normally distributed, so bare means would be misleading.	64
4.7	RQ2 efficiency profile. Node and edge retention are reported separately because several methods preserve node counts exactly. Savings are per-graph means paired against the full-graph baseline.	66
4.8	RQ3 matched-state accuracy: models trained and tested under the same reduction. RMSE, R^2 and Spearman are reported together throughout, because a reduction can hold mean error steady while destroying explained variance.	68
4.9	Matched-state accuracy partitioned two ways: by dataset tier, and by the synthesis script that produced the graph. Both partition the same test split, so the columns within each cut are disjoint. The source script separates the corpus far more sharply than the tier does.	69
4.10	Matched-state accuracy re-scored on four subsets of the same test split, from the persisted per-graph predictions. The pooled label is 49% exactly zero, so a pooled R^2 is carried by a minority of graphs; a conclusion that survives all four columns is a property of the reduction rather than of the label distribution.	73
4.11	Mean prediction and mean absolute error within bands of the true optimizability. R^2 is not reported: inside a narrow band there is almost no label variance to divide by. Compare each method’s predicted mean against the true mean in the first row. A row that stays flat across the bands is a model that has collapsed to a constant.	76
4.12	Accuracy against efficiency, ranked by matched-state R^2	78

4.13	RQ4 pairings. Δ is domain-informed minus generic, so a positive ΔR^2 favours the domain heuristic. The fourth pairing is NOT matched: random edge dropout at its configured rate keeps 69.7% of edges against spanning forest's 58.1%, so its gap conflates the heuristic with how much each arm removed.	83
4.14	RQ5: the same weights evaluated under the reduction they were trained with and on unreduced graphs. For every method here a full graph needs no conversion, it already is a valid input. The exact colour-refinement track is the exception and has not been run.	89
4.15	Every reduction method in one row: what it removes, what it costs offline, what it saves in training, what accuracy survives, and whether the model transfers to full graphs. Generic vs. domain-informed is a column, not a split, because it is one attribute of a method rather than the organising axis.	92
A.1	Optimization invocations, one per algorithm, as issued by the generation pipeline.	103

List of Symbols

Symbol	Description
Graph and structural notation	
$G = (\mathcal{V}, \mathcal{E})$	And-Inverter Graph (AIG): a directed acyclic graph over the basis $\{\wedge, \neg\}$, with node (gate) set $\mathcal{V}$ and edge set $\mathcal{E}$ (2.1.2).
τ	Node-role labelling, $\tau: \mathcal{V} \rightarrow \{\text{const}, \text{PI}, \text{AND}, \text{PO}\}$ (2.3).
inv	Edge-inversion flag, $\text{inv}: \mathcal{E} \rightarrow \{0, 1\}$, marking a connection as regular or inverted (2.4).
$\text{fi}(v), \text{fo}(v)$	Fanin set and fanout set of node v (2.2).
$\ell(v), \ell(G)$	Topological level of node v , and depth of G : $\ell(G) = \max_{v \in \mathcal{V}} \ell(v)$ (2.5).
$\mathcal{C}(v), \mathcal{F}(v)$	Fanin (causal) cone and fanout cone of node v (2.6).
$\text{MFFC}(v)$	Maximal Fanout-Free Cone of v : the gates that feed v exclusively, removed with v under restructuring (2.8).
Task, model, and reduction operators	
S	Deterministic synthesis algorithm (Orchestrate); $S(G)$ is the function-equivalent, optimized circuit (2.1.3).
$y(G)$	Optimizability of G under S : the fraction of nodes S removes, $y(G) = 1 - \mathcal{V}(S(G)) / \mathcal{V}(G) \in [0, 1]$ (2.24).
f_θ	Trained GNN encoder, approximating y from an input AIG (2.1.6).
R	Reduction operator; $R(G)$ is the reduced graph, and $R = \text{id}$ denotes the unreduced baseline (2.1.5).
Measuring reduction	
$\eta_{\mathcal{V}}(G), \eta_{\mathcal{E}}(G)$	Node and edge retention (compression) ratio of a reduction: $ \mathcal{V}(R(G)) / \mathcal{V}(G) $ and $ \mathcal{E}(R(G)) / \mathcal{E}(G) $ (2.23).
n^*	Amortisation threshold: the number of epochs beyond which applying a reduction is cheaper than training on the unreduced graph (3.40).

List of Abbreviations

Abbreviation	Expansion
AIG	And-Inverter Graph
EDA	Electronic Design Automation
FRAIG	Functionally Reduced AIG
GNN	Graph Neural Network
LS	Logic Synthesis
MFFC	Maximal Fanout-Free Cone
PI	Primary Input
PO	Primary Output

Introduction

Hardware complexity long ago outpaced manual design. Electronic Design Automation (EDA) compiles a behavioural description of a circuit into a manufacturable layout through a sequence of stages, and logic synthesis is the stage that fixes how many gates the result contains. It operates on the And-Inverter Graph (AIG), a Directed Acyclic Graph (DAG) whose internal nodes are two-input conjunctions and whose edges carry an optional inversion. Machine learning is now applied across the flow to predict what an expensive tool invocation would return, and at this stage one obstacle recurs: industrial AIGs are orders of magnitude larger than the graphs these models were built for. What follows treats the size of that input, rather than the capacity of the model, as the variable to change.

1.1 Motivation

1.1.1 Logic Synthesis and Optimization Script Design

Logic synthesis optimizes a circuit by applying a sequence of rewriting passes to it, and the sequence matters as much as the passes themselves: the same alphabet of transformations yields different area, power and delay depending on the order it is applied in. Choosing that sequence, a synthesis *script*, for a given circuit is heuristic trial and error over a tool that has to run to completion before its result is known.

1.1.2 Machine Learning for Optimizability Prediction

A model that predicted a circuit’s *optimizability*, the fraction of nodes a given script removes from it, would replace much of that search with a forward pass. Graph Neural Networks (GNNs) are the natural family for the task: an AIG is a graph, and the quantity to be predicted is structural. Scale is what stops them. Industrial AIGs reach millions of nodes, at which point training a GNN on them exceeds the memory available on a single device and drives training time beyond what is practical.

1.1.3 The Scale Bottleneck

The usual responses are to cap model capacity, accepting a weaker architecture than the task warrants, or to distribute training across many devices. Neither removes the constraint. A single AIG can exhaust one device on its own, at which point the graph has to be split whether or not the model is.

1.1.4 Limitations of Generic Graph Reduction

Graph reduction, the partitioning, sparsification and summarization methods that shrink a graph before a model reads it, is

the standard remedy on citation and social graphs. AIGs are not those graphs. Their edges carry a Boolean function, and their long paths carry the logic depth a synthesis pass has to work through, so a cut chosen to balance partition sizes or to preserve a spectrum can sever a causal cone and remove the very structure the prediction rests on.

1.1.5 Problem Statement and Research Gap

The field scales GNNs on AIGs in two ways today. The first is to train on extracted subcircuits of at most a few thousand gates, a crude sampling that discards global structure and is the de-facto reduction in most circuit-GNN work, including the DeepGate line [Li et al., 2022]. The second is to scale the *architecture*: DeepGate3 [Shi and TODO, 2024] pools over subcircuits with a transformer, and DeepGate4 [Zheng et al., 2025] uses sparse attention for sub-linear memory. Both change the model, not the input.

Graph reduction shrinks the *input* instead, and is well studied on citation, social and knowledge graphs [Hashemi et al., 2024]. It has never been systematically applied to whole AIGs as a preprocessing step for GNN training, and no work measures which reduction preserves a downstream synthesis label. The two literatures do not meet: generic summarization work [Bollen et al., 2023, Generale et al., 2022, Loukas, 2019] does not know about structural hashing or AIG semantics, and EDA work [Mishchenko et al., 2005, Li et al., 2022] does not frame its reductions as GNN-input summarization.

The gap this thesis addresses is therefore twofold: (i) there is no systematic evaluation of graph reduction for AIGs, and (ii) there is no evidence on whether AIG-specific adaptation of those reductions is worth its cost. Section 2.2.6 returns to both once the literature has been laid out, and qualifies the first of them.

1.2 Research Questions

Throughout, *optimizability* is the fraction of AIG nodes a given synthesis algorithm removes from a circuit (formalised in 2.1.6 and 3.3.4), and *compression* is reported as the node and edge retention ratios η_V and η_E of (2.23). Keeping the two apart is not cosmetic: several reduction methods hold $\eta_V = 1$ and lower η_E alone, so a single compression figure is ambiguous and RQ4 depends on the definition being pinned down first. Memory and time differences are taken per graph against the full-graph baseline and carry a 95% bootstrap interval and a two-sided paired rank test, whereas accuracy differences are single-run point estimates (3.4.3). Hypotheses are stated for RQ3 and RQ4 alone, the two questions whose direction is predicted in advance and can still be overturned by the result.

1.2.1 RQ1: Task Feasibility and Baseline

How well can a Graph Neural Network predict the optimizability of a circuit directly from its full, unreduced AIG representation, relative to trivial predictors and published circuit-representation models?

RQ1a: Protocol Sensitivity

How much of the RQ1 result survives a stricter split? The same model is trained under three train/test protocols differing in nothing but how graphs are assigned: a random split, a recipe-disjoint split, and the design-disjoint split behind every result reported here (3.4.2). The spread across the three separates structural inference from design recognition, and the per-design breakdown under the strictest protocol shows whether the residual error is spread evenly over held-out designs or carried by a few (4.1.6).

1.2.2 RQ2: Reduction Efficiency

Applied to AIGs, how do partitioning, sparsification, and summarization/coarsening differ in shrinkage rate, peak memory, training and inference speed, and the offline cost of computing the reduction itself?

1.2.3 RQ3: Predictive Retention and Trade-offs

To what extent do models trained on reduced AIGs retain predictive accuracy (measured via Smooth L1 Loss, RMSE, and R^2) and circuit-ranking ability (Spearman correlation) relative to the full-graph baseline established in RQ1, and which techniques offer the best accuracy-to-efficiency trade-off?

Hypothesis H1 (path contraction). At matched retention, summarization retains more accuracy than sparsification, because contraction raises the effective receptive field of (3.42) where edge deletion cannot: $\bar{c}_L(R(G)) > \bar{c}_L(G)$ under summarization and $\bar{c}_L(R(G)) \leq \bar{c}_L(G)$ under sparsification, at the encoder depth L . AIGs are deep relative to L , so the encoder is receptive-field starved, and over-squashing is the failure H1 expects contraction to relieve [Alon and Yahav, 2021].

1.2.4 RQ4: Value of Domain-Informed Adaptation

Do domain-informed reductions, which exploit topological level, reconvergence structure, and gate roles, retain more predictive accuracy than their generic counterparts at matched retention?

Hypothesis H2 (weak heuristics). The domain-informed reductions do not win. Edge-span weighting, level slicing, gate-role filtering and reconvergence-bounded merging are simple enough that at matched η_V and η_E they are not expected to separate from their generic counterparts on R^2 or Spearman correlation. Failure to separate them places the predictive signal across the AIG rather than in the causal cones and long paths these heuristics protect (4.4.4).

1.2.5 RQ5: Cross-State Inference

Can a model trained exclusively on reduced AIGs, whether partitioned, sparsified or summarized, be queried on full ones at inference time without losing predictive accuracy?

Inference is the regime where this pays off. Training stores activations and gradients for every node in a batch, whereas forward-only inference stores neither. A model that had to be trained on reduced graphs to fit in memory may therefore still serve full graphs on substantially more modest hardware. The practical claim under test is *train cheap, infer full*.

One reduction in the comparison is exact rather than approximate. Colour refinement at an unbounded count cap is provably lossless for this encoder (3.1.5), so its cross-state score must match the full-graph baseline, which is what separates a genuine cross-state failure from a broken evaluation path.

1.2.6 Scope and Delimitations

Optimizability prediction is the instrument, not the objective: it is the task the reductions are compared on, chosen for a label that is expensive to obtain and sensitive to structure. The delimitations below were fixed in advance. Limitations found during the work are separated out into 5.4.

Target synthesis algorithm. Training and evaluation target Orchestrate alone. The Deepsyn, Syn4 and C2RS graphs sharing the corpus are unused, since cross-algorithm variance would not isolate what a reduction costs (6.3.1).

Definition of optimizability. Optimizability is relative *node* reduction only. Area, delay and power, the objectives synthesis actually optimizes for, are not modelled, and node count stands in for area (5.4.4).

Graph scale. The corpus is bounded at an intermediate rather than industrial scale (3.4.2). RQ1 needs a full-graph baseline, and a baseline that does not fit in memory cannot be measured. The AIGs of millions to billions of nodes invoked in the motivation are out of reach here (5.4.2).

Reduction families. Partitioning, sparsification and summarization/coarsening are covered, the taxonomy of Hashemi et al. [2024] with partitioning substituted for condensation. Condensation is excluded because it establishes no node correspondence, leaving RQ5 undefined (2.2.3).

Model architecture. One encoder, an edge-aware GCN+ (3.1), held fixed so that outcome differences are attributable to the reduction. Published circuit models enter only as RQ1 baselines (3.1.6), so whether the ranking of reductions transfers is untested.

Prediction rather than script generation. Optimizability is predicted, not synthesis scripts. Spearman correlation ranks *circuits* under one fixed script, whereas selecting a script would re-

quire ranking *algorithms* for one circuit, which a single-algorithm model cannot do (6.3.4).

1.3 Contributions

This thesis reduces the input rather than the model. Its contributions are (i) a corpus of AIGs tiered by optimization state, extending OpenABC-D [Chowdhury et al., 2021] (3.3), (ii) a systematic benchmark of partitioning, sparsification and summarization as a preprocessing step for GNN training on AIGs, measured on shrinkage, peak memory, speed and predictive retention, (iii) AIG-specific adaptations of each family, built from topological level, reconvergence structure and gate role and compared against their generic counterparts at equal compression, (iv) a coarsening that is provably lossless for this encoder [Bollen et al., 2023], fixing a point of zero accuracy cost against which every lossy method can be read (3.1.5), (v) an evaluation of cross-state inference, in which one set of weights is trained on reduced graphs and queried on full ones, and (vi) a controlled comparison of three train/test protocols, measuring how much of a reported score is design recognition rather than structural inference (3.4.2).

The corpus is the enabling contribution, and it departs from OpenABC-D on three axes. It draws on 55 source designs against 29, adding the deep arithmetic of the EPFL, ISCAS-85 and MCNC suites, which OpenABC-D excludes, together with eight randomly generated AIGs. Every base graph carries a label under each of four synthesis scripts, and the outputs of three of them form a second tier, labelled again under the target script. OpenABC-D varies the transformation recipe applied to one starting circuit, whereas this construction also varies the optimization state of the circuit that script is applied to, which is what RQ3 and RQ5 require (2.2.1).

1.3.1 Summary of Results

1.4 Thesis Outline

Chapter 2 introduces the necessary background (logic synthesis and AIGs, graph neural networks, and the three families of graph reduction), states the prediction problem in the notation used throughout, and then reviews the literature the gap in 1.1.5 sits in. Chapter 3 describes the encoder and the baselines it is measured against, the reduction methods implemented for each family, the dataset and its labels, and the experimental protocol. Chapter 4 reports results per research question. Chapter 5 interprets them against the literature and states the limitations, and Chapter 6 answers each research question in turn and outlines future work. Command templates, reduction parameters and baseline hyperparameters are recorded in the appendix.

Preliminaries & Related Work

2.1 Preliminaries

2.1.1 Electronic Design Automation and Logic Synthesis

Electronic Design Automation

Electronic Design Automation (EDA) is the tool chain that compiles a functional description of a digital circuit into a manufacturable layout [De Micheli, 1994]. Its optimization stages fix the area, power and delay of the fabricated chip. Every label used in this thesis is the output of one such tool, and the cost of obtaining it is what motivates prediction.

Position of Logic Synthesis in the Design Flow

The flow is a sequence of refinements: a Register-Transfer-Level (RTL) description is compiled into a network of Boolean gates, Logic Synthesis (LS) restructures that network against cost objectives, technology mapping binds the result to the cells of a target library, and placement and routing assign physical positions and wires. LS is pre-physical. Its network has no coordinates, no wire lengths and no layout-extracted timing, which is what makes the spatial coarsening methods of 2.2.4 inapplicable to it.

Synthesis Scripts, Recipes, and Steps

A step is one function-preserving transformation of the network, and a synthesis script is an ordered composition of steps,

$$S = t_m \circ t_{m-1} \circ \dots \circ t_1. \quad (2.1)$$

The t_i do not commute. What each one gains depends on the structure its predecessors produced, so the same script yields different gains on different designs and script choice cannot be fixed once and reused. A recipe is a script whose steps are drawn at random from a fixed alphabet of passes, which is how the corpus of 3.3.2 is built, one stored graph per prefix.

2.1.2 And-Inverter Graphs

Definition and Structure

An And-Inverter Graph (AIG) is a Directed Acyclic Graph (DAG) $G = (\mathcal{V}, \mathcal{E})$ representing a Boolean network over the basis $\{\wedge, \neg\}$ [Mishchenko et al., 2006]. Every internal node is a two-input AND gate, and every edge carries an optional inversion, so conjunction is the only operation that costs a node and negation costs none. The basis is functionally complete, which is what allows an arbitrary Boolean network to be expressed in one uniform structure rather than in a gate library.

Completeness does not buy uniqueness. One Boolean function admits many AIGs, and which one an optimizer is handed bounds what it can remove, a sensitivity known as structural bias [Gardner et al., 2025]. Optimizability is therefore a property of a graph rather than of the function it computes, which is what makes it predictable from structure at all. Structural bias also makes corpus construction consequential, since a corpus assembled by running synthesis scripts carries whatever structure those scripts favour (3.3.6).

The vocabulary used throughout follows from the edge set. The fanin and fanout of a node are

$$\text{fi}(v) = \{u \in \mathcal{V} : (u, v) \in \mathcal{E}\}, \quad \text{fo}(v) = \{w \in \mathcal{V} : (v, w) \in \mathcal{E}\}, \quad (2.2)$$

with $|\text{fi}(v)| = 2$ at every AND gate, $\text{fi}(v) = \emptyset$ at every Primary Input (PI) and at the constant, and $|\text{fi}(v)| = 1$ at every Primary Output (PO). Acyclicity yields a topological order $\prec$ on $\mathcal{V}$ with $u \prec v$ for every $(u, v) \in \mathcal{E}$, so signals flow from the PIs to the POs and never return.

The POs are materialised as nodes here, rather than as pointers into the AND network as is also common. That choice is why the node role defined below takes four values instead of three, and it puts the circuit interface inside $|\mathcal{V}|$, the quantity the target of 3.3.4 is measured on.

Node Types and Inverted Edges

Two labellings are carried on every graph. A node role

$$\tau : \mathcal{V} \rightarrow \{\text{const}, \text{PI}, \text{AND}, \text{PO}\} \quad (2.3)$$

separates the constant, the PIs, the internal AND gates and the POs, and an inversion flag $\text{inv} : \mathcal{E} \rightarrow \{0, 1\}$ marks each edge as a regular connection or an inverted one. Inversion is semantic rather than decorative. Writing $b_u \in \{0, 1\}$ for the value produced at u , the value an edge $e = (u, v)$ delivers to v is

$$b_u \oplus \text{inv}(e), \quad (2.4)$$

so two edges agreeing in their endpoints and differing in inv carry complementary signals. The one-hot encodings of τ and inv are the four-valued node feature and two-valued edge feature of 3.3.5, and are not restated there. A small AIG carrying both labellings is drawn in 2.1.

Structural Properties: Levels, Logic Cones, Fanout

The topological level of a node is its longest distance from the input frontier,

$$\ell(v) = \begin{cases} 0 & \text{if } \text{fi}(v) = \emptyset, \\ 1 + \max_{u \in \text{fi}(v)} \ell(u) & \text{otherwise,} \end{cases} \quad (2.5)$$




Figure 2.1: An AIG over three primary inputs x_1, x_2, x_3 and two primary outputs y_1, y_2 , with signals flowing upward. Circles are two-input AND gates, solid edges are regular, and dashed edges are inverted in the sense of (2.4). Height is the topological level of (2.5), so the graph has depth $\ell(G) = 5$. The two gates immediately above the inputs share both fanins and differ only in inversion, which is what keeps them distinct under structural hashing.

and the depth of the graph is $\ell(G) = \max_{v \in \mathcal{V}} \ell(v)$. The drawn height of a gate in 2.1 is exactly ℓ . Level is an absolute coordinate that an AIG has and a general graph does not, and four constructions in this thesis read it: the positional encoding of 3.3.5, the level slicing and span-weighted partitioners of 3.2.1, and the level bands of the cone candidate in 3.2.3.

Writing $u \rightsquigarrow v$ for the existence of a directed path, of length zero when $u = v$, the fanin cone (equivalently the logic or causal cone) and the fanout cone of v are

$$\mathcal{C}(v) = \{u \in \mathcal{V} : u \rightsquigarrow v\}, \quad \mathcal{F}(v) = \{w \in \mathcal{V} : v \rightsquigarrow w\}. \quad (2.6)$$

Both cones contain v itself. The fanin cone is the set of gates on which v functionally depends, and (2.5) measures the longest path through it. The claim of 1.1.4, that generic reduction severs causal structure, is a claim about $\mathcal{C}$: deleting one edge on a long path removes gates from the cones of everything downstream of it, whether or not the result resembles the original under an undirected measure. The depth-bounded form $\mathcal{C}_k$ of (3.41) is what makes the effect measurable.

Fanout degree $|\text{fo}(v)|$ separates gates whose output is consumed once from gates whose output is shared. Chains of single-fanout gates carry no sharing information; multi-fanout gates are the points at which logic is reused, and they are what constrains rewriting. Reconvergence is the same property seen from below: v is a reconvergence point of u when two internally disjoint directed paths run from u to v . AIGs are dense in reconvergence, because structural hashing (2.1.3) shares every duplicated subexpression instead of replicating it [Kuehlmann et al., 2002].

Where a fanout cone reconverges is a question of post-dominance [Lengauer and Tarjan, 1979]. Append a virtual sink t succeeding every PO, and extend $\rightsquigarrow$ to it. The immediate post-dominator of v is the earliest gate every path from v to t must cross,

$$\text{ipdom}(v) = \min_{\rightsquigarrow} \{u \neq v : \text{every } v \rightsquigarrow t \text{ path meets } u\}, \quad (2.7)$$

which is the gate at which $\mathcal{F}(v)$ reconverges, or t when that

cone never reconverges. Grouping gates by ipdom is what the cone candidate of 3.2.3 uses.

The Maximal Fanout-Free Cone (MFFC) of a root gate v collects the gates feeding v and nothing else [Mishchenko et al., 2006],

$$\text{MFFC}(v) = \{u \in \mathcal{C}(v) : \text{every path } u \rightsquigarrow t \text{ passes through } v\}, \quad (2.8)$$

so removing v from the network removes exactly $\text{MFFC}(v)$ and nothing more. That property makes the MFFC the natural unit of local restructuring, and refactoring collapses and re-expresses one MFFC at a time (2.1.3). It is also the unit merged by the second domain candidate of 3.2.3, which is why that candidate carries an argument about label leakage and the others do not.

Why AIGs Differ from General Graphs

Five properties separate an AIG from the citation, social and knowledge graphs on which graph reduction was developed, and each costs a generic method something specific.

(i) An AIG is directed and acyclic. Machinery built for undirected graphs does not transfer unchanged, because a quotient of a DAG need not be acyclic, and a reduction computed on the undirected projection discards the orientation that carries the logic. The spanning forest of 3.2.2 is the concrete case.

(ii) Depth, not width, is the binding dimension. The critical path is a longest directed path from a PI to a PO, and its length is the depth $\ell(G)$ of (2.5): the quantity balancing exists to shorten (2.1.3) and the quantity an encoder has to span. After L message-passing layers a node representation depends on its L -hop neighbourhood and on nothing else (2.1.4), so a gate reads the whole of the logic it depends on, $\mathcal{C}_L(v) = \mathcal{C}(v)$, only when

$$L \geq \ell(v). \quad (2.9)$$

Corpus depth is 35 levels at the median and 24,802 at the maximum (3.3.1), against an encoder depth in single digits, so (2.9) holds only near the input frontier. Receptive-field starvation is the normal condition on this data, and it is the starvation that hypothesis H1 (1.2.3) proposes coarsening relieves.

(iii) Edges carry function. The inversion in (2.4) is part of what the graph represents, so a merge that averages or discards inv changes the Boolean function rather than approximating it. Every summarization method here is therefore constrained to keep normal and inverted connections distinguishable, a constraint no general method imposes on itself.

(iv) The interface is fixed. The PIs and POs are the circuit's contract with the rest of the design, and merging them away changes what the graph is rather than coarsening it. The boundary is enforced uniformly in 3.2.3.

(v) Node count is at once the cost and the scale of the label. Fanin cardinality is fixed by the basis, so

$$|\mathcal{E}| = 2|\tau^{-1}(\text{AND})| + |\tau^{-1}(\text{PO})| < 2|\mathcal{V}|, \quad (2.10)$$

and one message-passing layer therefore costs $\Theta(|\mathcal{V}|)$ at fixed width. An AIG carries no surplus density, so removing edges alone cannot take that cost below the same order at any retention ratio (2.1.5). The same count is the denominator of the optimizability target (2.24), so $y(G)$ is quantised in steps of $1/|\mathcal{V}(G)|$

and is coarsest on the smallest circuits, which is where the corpus piles up at exactly zero (3.3.6).

A reduction operator can respect or ignore the first four at no syntactic cost, since the graph it produces is well formed either way. Whether respecting them buys accuracy is the question RQ4 (1.2.4) measures.

2.1.3 Synthesis Algorithms

Rewriting, Refactoring, and Balancing

Every script in this thesis is assembled from four local, function-preserving transformations [Mishchenko et al., 2006]. A cut of a gate v is a node set through which every PI-to- v path passes, and it is K -feasible when it has at most K members. Rewriting replaces the logic inside a K -feasible cut with the smallest precomputed equivalent, accepting the replacement only on a net gain [Bjesse and Borälv, 2004, Mishchenko et al., 2006]. Refactoring makes the same exchange with the MFFC of (2.8) as the unit, and resubstitution re-expresses a gate through gates already present in the network, removing the MFFC left redundant [Mishchenko and Brayton, 2006]. Balancing restructures AND trees under associativity to reduce the depth $\ell(G)$ of (2.5) at an unchanged function.

Structural hashing carries a signature on every AND gate [Kuehlmann et al., 2002],

$$\kappa(v) = \{ (u, \text{inv}(u, v)) : u \in \text{fi}(v) \}, \quad (2.11)$$

and maintains $\kappa(u) \neq \kappa(v)$ for distinct AND gates on construction, so one-hop structural redundancy is gone before any summarization runs (3.2.3).

Functional equivalence is the stronger relation,

$$u \equiv v \iff b_u = b_v \text{ under every PI assignment}, \quad (2.12)$$

with b_u as in (2.4). A Functionally Reduced AIG (FRAIG) merges every such pair [Mishchenko et al., 2005], whereas (2.11) decides (2.12) only in the special case $\kappa(u) = \kappa(v)$. Every summarization method here stays on the structural side of that line, because a merge by (2.12) performs part of the optimization the model is trained to predict (5.4.5).

Optimization Objectives: Area, Power, Delay

Area, power and delay are the costs synthesis is judged on, and none of the three is physical before technology mapping, so each is tracked through a structural proxy on the AIG: node count $|\mathcal{V}(G)|$ for area and depth $\ell(G)$ of (2.5) for delay. The ideal is the smallest network computing the same function,

$$G^* \in \arg \min_{H \equiv G} |\mathcal{V}(H)|, \quad (2.13)$$

which no practical script attains. A script instead returns some $S(G) \equiv G$ with $|\mathcal{V}(S(G))| \leq |\mathcal{V}(G)|$. Both proxies are inherited downstream: technology mapping binds fewer cells from a smaller network and reports a shorter critical path from a shallower one.

The Orchestrate Script

Orchestrate, the target algorithm, moves the choice of transformation into the graph [Li et al., 2024]. At each AND gate, rewriting, refactoring and resubstitution are attempted, the one with the best local area-delay score is committed, and the sweep repeats for a bounded number of passes. Two properties make it the target: (i) it subsumes the orderings of the fixed scripts built from the same primitives, so the label measures optimization potential rather than the accident of one command sequence, and (ii) it is deterministic, so $S(G)$ is a single graph rather than a distribution over runs and optimizability under Orchestrate is well defined as the $y(G)$ of 3.3.4. The exact invocation is recorded in A.1.

The Non-Target Scripts: Deepsyn, Syn4, and C2RS

Three further scripts of the synthesis system were applied during dataset generation [Brayton and Mishchenko, 2010], and every tier-1 graph that enters training is the output of one of them (3.3.3). Deepsyn searches: randomized restructuring passes are applied under a wall-clock budget and the best network encountered is returned. Its seed was drawn afresh for every graph and never recorded, so that portion of the corpus cannot be regenerated (5.4.1). Syn4 is a fixed deterministic flow over the same primitives. C2RS is a chain of balancing, refactoring, rewriting and resubstitution steps [Mishchenko and Brayton, 2006] whose resubstitution cut widens from $K = 6$ to $K = 12$, so successive steps reason over progressively larger windows. The invocations are recorded in A.1.

2.1.4 Graph Neural Networks

Message Passing Framework

A Graph Neural Network (GNN) computes a representation $h_v^l \in \mathbb{R}^w$ for every node v by repeated neighbourhood aggregation [Kipf and Welling, 2017, Hamilton et al., 2017]. Writing $\mathcal{N}(v)$ for the neighbours of v and e_{uv} for the attribute of edge (u, v) , one layer computes

$$h_v^l = \phi^l \left(h_v^{l-1}, \bigoplus_{u \in \mathcal{N}(v)} \psi^l(h_u^{l-1}, e_{uv}) \right), \quad (2.14)$$

where ψ^l builds a message per edge, $\bigoplus$ is a permutation-invariant aggregation such as a sum, and ϕ^l updates the node state. The encoder of this thesis is one instantiation of (2.14), written out in 3.1.1.

After L layers, h_v^L is a function of exactly the L -hop neighbourhood of v and of nothing outside it. On an AIG that neighbourhood is a depth- L prefix of the cones of (2.6), so an encoder of depth L reads L levels of logic around each gate. Every depth-coupling argument in this thesis rests on this one fact: the refinement depth of 3.2.3 is set to L , the receptive-field metric of 3.4.3 counts the gates inside L hops, and hypothesis H1 (1.2.3) is a claim about what path contraction does to those L hops.

Graph-Level Readout and Pooling

A graph-level prediction pools the node representations into one vector, here by a mean over nodes, optionally after summing the per-layer outputs so that shallow and deep features both reach the readout (jumping knowledge, Xu et al., 2018). Pooling imposes the constraint that shapes the whole experimental setup: every node of a graph must be present in the same forward pass, so a graph is the indivisible unit of batching and cannot be split across steps. Batches are therefore assembled under a node budget rather than as a fixed graph count (3.4.2), and the memory cost of one step is set by the largest graphs, which is exactly the cost reduction is meant to lower.

Positional and Structural Encodings

Message passing sees only relative structure: two nodes whose L -hop neighbourhoods are isomorphic receive identical representations wherever they sit in the graph. A Positional Encoding (PE) breaks that indifference by attaching a coordinate to each node. On an undirected graph no canonical coordinate exists and encodings are built from random walks or spectra; on a DAG the topological level $\ell(v)$ of (2.5) is an absolute coordinate that the graph carries for free. The level-based encoding used here (3.3.5) is therefore a DAG-native choice rather than a generic one.

Expressivity and the Weisfeiler-Leman Hierarchy

Colour refinement, equivalently the one-dimensional Weisfeiler-Leman algorithm (1-WL), iteratively refines a node colouring by the multiset of neighbour colours,

$$c_v^{(t)} = \text{hash}\left(c_v^{(t-1)}, \left\{ \left\{ c_u^{(t-1)} : u \in \mathcal{N}(v) \right\} \right\}\right), \quad (2.15)$$

starting from initial labels $c_v^{(0)}$. A partition of $\mathcal{V}$ is *equitable* when every node of a class has the same number of neighbours in every class; run to a fixed point, (2.15) computes the coarsest equitable partition refining the initial colouring [Grohe et al., 2014]. The connection to (2.14) is an upper bound: a message-passing network is at most as discriminative as 1-WL, so two nodes sharing a colour after L rounds receive identical representations from every L -layer encoder [Xu et al., 2019].

Replacing the multiset in (2.15) by a set yields bisimulation, which ignores how many neighbours of a colour exist and asks only whether one does. A count c interpolates between the two, truncating each multiplicity at c : the cap $c = 1$ is bisimulation and $c = \infty$ is colour refinement [Bollen et al., 2023]. The bound above is what makes quotienting by such a partition lossless rather than approximate: merging nodes the network provably cannot distinguish removes no information the network could have used, and the automorphism orbits of a graph form an equitable partition, so genuinely symmetric structure is always compressible this way. AIGs have such structure in quantity. The two fanins of an AND gate are interchangeable, datapath logic replicates one bit slice across the word width, and standard sub-circuits recur throughout a design, each of which creates orbits that a random graph does not have. How much of this redundancy survives structural hashing is measured in 3.2.3.

Training Cost, Memory, and Inference

Training and inference differ in what must be held in memory. Backpropagation through L applications of (2.14) retains the node states of every layer and the messages materialised on every edge, whereas forward-only inference keeps one layer resident and discards it. Above a graph-independent floor of parameters and optimizer state, the memory a graph itself costs is therefore

$$\begin{aligned} \Delta m_{\text{train}} &= \Theta(L(|\mathcal{V}| + |\mathcal{E}|)w), \\ \Delta m_{\text{infer}} &= \Theta((|\mathcal{V}| + |\mathcal{E}|)w), \end{aligned} \quad (2.16)$$

and the factor L between them is the whole asymmetry. The largest graph in the corpus sets the peak, since pooling forbids splitting a graph across steps (2.1.4).

The factor L in (2.16) is the mechanism RQ5 (1.2.5) depends on. Training on $R(G)$ and querying on G stay inside one memory budget when

$$|\mathcal{V}(G)| + |\mathcal{E}(G)| \leq L(|\mathcal{V}(R(G))| + |\mathcal{E}(R(G))|), \quad (2.17)$$

so any reduction retaining more than a $1/L$ share of G admits full-graph inference on the hardware its own training required. Training is the binding constraint, not inference.

Four techniques recur in Chapter 3 and are defined here once. Reported memory is *allocated* memory, the tensors actually held, as opposed to the larger *reserved* pool the allocator keeps (3.4.3). Mixed precision stores activations in half-width floating point and roughly halves their cost. Gradient checkpointing discards intermediate activations and recomputes them during the backward pass, trading compute for memory, and is what makes the largest baseline runnable at all (3.1.6). Gradient accumulation sums gradients over several small steps before updating, so the effective batch size exceeds what fits in one step. All four trade something for memory at a *fixed* input size, which is what separates them from the input reduction this thesis studies (2.2.2).

Depth has a cost of its own. A fixed-width representation cannot carry the information of a receptive field that grows exponentially with depth, so signal from distant nodes is compressed away, a failure known as over-squashing [Alon and Yahav, 2021]. Contracting paths shortens the distance signal must travel and is a recognised remedy [Dwivedi et al., 2022]. That contraction is the mechanism behind hypothesis H1 (1.2.3), which expects summarization to help where sparsification can only remove.

2.1.5 Graph Reduction: Definitions

A reduction operator R maps a graph G to a smaller graph $R(G)$ intended to stand in for it. The taxonomy follows the survey of Hashemi et al. [2024], with one substitution: partitioning replaces condensation as the third family, because partitioning is what a distributed setup forces on a graph that does not fit one device (2.2.2) and condensation is excluded from this study outright (2.2.3).

Partitioning

Partitioning splits $\mathcal{V}$ into k disjoint blocks and deletes every edge whose endpoints fall in different blocks. Every node survives, so

whatever it compresses is the edge set alone.

Sparsification

Sparsification removes edges, or nodes together with their incident edges, while preserving some structural property of the original, such as pairwise distances up to a stretch factor. Edge-removing and node-removing variants differ in whether $|\mathcal{V}|$ moves at all, which is what forces the two-ratio convention of 2.1.5.

Summarization and Coarsening

Summarization merges nodes into *super-nodes*. A coarsening is specified by a merge map, a surjection $\mu: \mathcal{V} \rightarrow \mathcal{V}'$ onto the cluster set, and the resulting quotient graph connects two clusters whenever any of their members were connected. The number of member edges a super-edge stands for is its multiplicity,

$$w_{u'v'} = |\{(u, v) \in \mathcal{E} : \mu(u) = u', \mu(v) = v'\}|, \quad (2.18)$$

and a quotient that keeps attribute-distinct super-edges parallel is a multigraph. Because μ is a function, a coarsening is a partition of $\mathcal{V}$, so a tolerance band such as “within one level” cannot define one: the relation it induces is not transitive (3.2.3).

Exact Summarization

A coarsening is *exact* for an encoder when the encoder cannot tell the quotient from the original. Writing f_θ for an L -layer instance of (2.14) with a graph-level readout, and R for the coarsening with merge map μ , exactness is

$$f_\theta(R(G)) = f_\theta(G) \quad \text{for every } \theta. \quad (2.19)$$

Quantifying over θ is what separates exactness from an approximation that happens to be accurate at the weights that were trained. The conditions under which (2.19) holds [Bollen et al., 2023] constrain three separate objects, and a method meeting two of the three is not exact.

(a) The quotient. Merged nodes must be indistinguishable to L rounds of colour refinement (2.1.4),

$$\mu(u) = \mu(v) \implies c_u^{(L)} = c_v^{(L)}, \quad (2.20)$$

so the partition induced by μ refines the L -round colour partition, whose classes are the coarsest admissible merge. Two side conditions come with (2.20). The initial colouring $c^{(0)}$ must separate distinct feature rows, and (2.15) must run over the same neighbourhood and the same edge attributes the encoder aggregates over: refining by fan-out while aggregating over fan-in merges nodes the encoder can still tell apart. The count cap must also match what the aggregation can count. Capping below that makes the colour a coarser invariant than the aggregate the encoder computes, which is why bisimulation is exact for a set-valued aggregation and lossy for a counting one.

(b) The aggregation. Collapsing $w_{u'v'}$ parallel edges into one super-edge replaces $w_{u'v'}$ identical messages by a single message, so $\bigoplus$ must read the multiplicity of (2.18) rather than the

set of distinct neighbours. For the summing encoder used here that multiplicity enters as a coefficient on the message:

$$\bigoplus_{u \in \mathcal{N}(v)} \psi(h_u, e_{uv}) = \sum_{u' \in \mathcal{N}'(\mu(v))} w_{u'\mu(v)} \psi(h_{u'}, e_{u'\mu(v)}), \quad (2.21)$$

where $\mathcal{N}'$ is the neighbourhood taken in the quotient. Where the multiplicity enters decides whether (2.21) holds. The message function ψ is nonlinear, so $\psi(h, w_{u'v'}e) \neq w_{u'v'} \psi(h, e)$ and a multiplicity carried as an edge attribute breaks the equality that a coefficient outside ψ preserves (3.1.5).

(c) The readout. Conditions (a) and (b) give $h_{\mu(v)}^l = h_v^l$ at every layer, so each super-node carries the representation its members would have received. Pooling over $\mathcal{V}'$ must then weight each super-node by its member count $m_{v'} = |\mu^{-1}(v')|$ to reproduce pooling over $\mathcal{V}$:

$$\frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} h_v^L = \frac{\sum_{v' \in \mathcal{V}'} m_{v'} h_{v'}^L}{\sum_{v' \in \mathcal{V}'} m_{v'}}. \quad (2.22)$$

An unweighted mean over super-nodes violates (2.22), giving a super-node of fifty members the weight of one. The member counts already carried by a summarized feature schema are what make the weighting available.

Colour refinement at unbounded count cap is the one reduction in this study that satisfies all three (3.2.3). Sparsification and partitioning produce no quotient and delete edges the encoder aggregates over, so (2.19) fails for them by construction, and exactness is unavailable to those families rather than merely unachieved.

Measuring Reduction: Compression Ratio

Reduction is reported as retention, kept over original, separately per node and per edge:

$$\eta_{\mathcal{V}}(G) = \frac{|\mathcal{V}(R(G))|}{|\mathcal{V}(G)|}, \quad \eta_{\mathcal{E}}(G) = \frac{|\mathcal{E}(R(G))|}{|\mathcal{E}(G)|}, \quad (2.23)$$

with reduction $1 - \eta$. One number cannot carry both: partitioning and edge-removing sparsification hold $\eta_{\mathcal{V}} = 1$ while lowering $\eta_{\mathcal{E}}$, and node-removing methods move the two together. How the ratios of (2.23) are measured on this corpus, and how methods whose compression is an outcome rather than a parameter are brought to a matched ratio, is in 3.4.3.

2.1.6 Problem Formalization

The objects fixed above assemble into one problem statement. Let G be an AIG with the node roles of (2.3) and edge inversions of (2.4) as its features, and let S be a deterministic synthesis algorithm (2.1.3) mapping G to the function-equivalent, optimized $S(G)$. The prediction target is the *optimizability* of G under S , the fraction of nodes the algorithm removes,

$$y(G) = 1 - \frac{|\mathcal{V}(S(G))|}{|\mathcal{V}(G)|} \in [0, 1]. \quad (2.24)$$

Obtaining $y(G)$ means running S , which is the expensive step prediction is meant to avoid. An encoder f_θ , a GNN in the sense

of 2.1.4 with a graph-level readout, is trained to approximate y . A reduction operator R , drawn from the three families of 2.1.5, shrinks the encoder’s input.

RQ5 (1.2.5) is well posed only under a compatibility requirement: $R(G)$ and G must occupy the same input space, so that one set of weights θ ingests both. Every summarization method here therefore produces features of which the full-graph features are the single-member special case (3.2.3). A reduction that changed the schema would make cross-state inference undefined rather than merely inaccurate.

2.2 Related Work

Two literatures bear on this thesis. Graph reduction is developed enough to have surveys of its own [Hashemi et al., 2024, Shabani et al., 2023]. Circuit representation learning is developed enough to support a growing model family. The subsections below take each in turn, then the two nearest attempts at a bridge between them, before 2.2.6 states what still separates them.

2.2.1 Machine Learning for Logic Synthesis

Quality-of-Result and Optimizability Prediction

The direct precedents predict synthesis Quality of Result (QoR) from a circuit before running the tool. OpenABC-D [Chowdhury et al., 2021] is the reference dataset for the task and publishes, among its graph-level labels, the percentage of nodes a recipe optimizes away, essentially the target (2.2.4); that overlap is what makes its accompanying SynthNet a fair baseline here rather than a repurposed model (2.2.1). Recipe-conditioned models jointly encode the circuit and the synthesis sequence [TODO, 2022, Wu et al., 2022]. One terminological collision is worth heading off: LOSTIN [Wu et al., 2022] attaches a super-node that encodes the synthesis *sequence*, a temporal use of the same device this thesis uses structurally, and the two share nothing beyond the name.

Learned Synthesis Script Generation

A separate line learns to construct or select synthesis scripts, by reinforcement learning or sequence prediction over the same primitive transformations. It is the downstream application the motivation invokes, and it is out of scope here (1.2.6): script generation consumes exactly the kind of cheap optimizability estimate this thesis builds, which is why the prediction task is worth solving on its own. DRILLS [Hosny et al., 2020] is representative, an RL agent choosing among the transformations of 2.1.3 at each step rather than committing to a fixed script in advance.

Circuit Representation Learning

The DeepGate line [Li et al., 2022, Shi and TODO, 2024, Zheng et al., 2025] learns general-purpose gate embeddings on AIGs, and PolarGate [Liu et al., 2024] identifies the handling of inverted connections as the representational bottleneck of AIG encoders, external support for treating inversion as a relation

summarization must preserve rather than average away (3.2.3). Two facts from this line shape the present study. First, DeepGate3 and DeepGate4 answer scale by growing the *model*, with sparse attention and larger training corpora, where this thesis shrinks the *input*. DeepGate4 doubles as a baseline (2.2.1), so the contrast is measured rather than rhetorical. Second, these models train on extracted subcircuits of tens to a few thousand gates. Extraction is a reduction, applied at corpus construction and never named or evaluated as one, and that unexamined step is half of the gap this chapter closes in 2.2.6.

Benchmarks and Datasets for ML in EDA

OpenABC-D [Chowdhury et al., 2021] remains the largest published synthesis-prediction corpus: 29 designs, 1500 recipes each, labels per design-recipe pair. The dataset built for this thesis (3.3) differs in axis rather than only in size: it is tiered by optimization state, pairing each base graph with already-optimized variants of itself (3.3.3), because the reduction questions RQ3 and RQ5 need graphs at several points along the optimization trajectory, not several recipes on one starting point. OpenLS-DGF [Ni et al., 2025] takes a third axis again, generating designs adaptively rather than fixing a benchmark suite, and Chowdhury et al. [2023] argue the position this dataset acts on: ML4EDA lacks the standard, prototypical datasets that drove progress in computer vision.

Evaluation Protocols and the Unseen-Design Gap

“Unseen” means three different things in this literature: an unseen recipe, an unseen design, or an unseen design-recipe pair whose parts were both seen, and the third is far easier than the second. Generalization to unseen designs is the stated motivation of much of the field, yet matched seen and unseen results are rare: of the papers surveyed for this thesis, 4 of 63 report both sides, OpenABC-D [Chowdhury et al., 2021], LOSTIN [Wu et al., 2022], LSOformer [Karimi et al., 2024], and the XGBoost AIG-timing predictor of Jiang et al. [2025], while HOGA [Deng et al., 2024] motivates on the gap without publishing a seen-design baseline, so its degradation cannot be computed. The canonical unseen-IP split trains on 16 small designs and tests on 8 large ones, confounding design novelty with size extrapolation. Reported error moves more with the benchmark than with the architecture, from a few percent on small homogeneous suites to above twenty percent on OpenABC-D, and within one test set the error concentrates in a few designs rather than spreading uniformly. Those three observations motivate RQ1a’s matched three-protocol comparison and the per-design reporting used throughout (1.2.1).

Models Used as Baselines

Three published models are used for comparison here, and 3.1.6 states only how each was adapted, so what each one *is* belongs here.

SynthNet. SynthNet [Chowdhury et al., 2021] is the reference architecture released with OpenABC-D, the dataset’s own

demonstration of a learnable, recipe-conditioned QoR predictor. A graph convolutional encoder produces node states h_v as in (2.14), a second encoder ρ reads the synthesis recipe s as a sequence of transformation choices, and the two are concatenated ahead of a regression head:

$$\hat{y}(G, s) = \text{MLP}\left(\pi(\{h_v\}_{v \in \mathcal{V}(G)}) \parallel \rho(s)\right), \quad \pi(\cdot) = [\max(\cdot) \parallel \text{mean}(\cdot)]. \quad (2.25)$$

Conditioning the prediction on s is the distinctive move: one trained model answers a QoR query for any recipe in the search space instead of for one fixed script, turning QoR prediction into an oracle a script search can call cheaply and repeatedly.

HOGA. HOGA [Deng et al., 2024] reports a large reduction in QoR estimation error over conventional GNNs and is motivated by generalization to unseen designs, the architecture-side answer to the scalability limit this thesis addresses from the input side. Hop-wise features are propagated once before training and combined by a gated self-attention module:

$$\begin{aligned} x_v^{(k)} &= \hat{A}^k x_v^{(0)}, & k &= 0, \dots, K, \\ h_v &= \sum_{k=0}^K \alpha_v^{(k)} x_v^{(k)}, & \alpha_v^{(k)} &= \text{softmax}_k(a^\top \tanh(W x_v^{(k)})), \end{aligned} \quad (2.26)$$

with $\hat{A}$ the normalized adjacency. Only the $K + 1$ vectors $x_v^{(k)}$ reach the learned model, never $\mathcal{E}(G)$ itself, so training is attention over fixed-size vectors, decoupled from $|\mathcal{V}(G)|$ and from (2.14)’s coupling between depth and receptive field. Because the propagation step runs once and is never revisited, nodes can be shuffled or split arbitrarily downstream of it, which is what makes training distributable across devices.

DeepGate4. DeepGate4 [Zheng et al., 2025] is the largest model in the DeepGate line and reports the strongest circuit representation results in it, a graph transformer built to stay attention-based at sizes where dense attention no longer fits in memory. Full self-attention over $\mathcal{V}(G)$ costs $\Theta(|\mathcal{V}(G)|^2)$ per layer; DeepGate4 instead restricts every query to a distance-bounded virtual edge set,

$$\bar{\mathcal{E}}_k = \{(u, v) : \text{dist}(u, v) \leq k\}, \quad (2.27)$$

paired with an update rule the authors report as adaptable to any graph transformer, which keeps activation memory sub-linear in graph size at a fixed radius k . The restriction is sparse relative to full attention, not relative to the graph: DeepGate4 answers the memory problem by growing the receptive field efficiently, where the reduction studied here shrinks $\mathcal{V}(G)$ before any attention or message passing runs.

2.2.2 Scaling Graph Neural Networks to Large Graphs

Sampling-Based Approaches

The oldest answer to a graph that does not fit is to sample it: neighbour sampling caps the fanout of each aggregation [Hamilton et al., 2017], and subgraph sampling trains on induced pieces. Both are per-epoch stochastic operations built for node-level

targets, and neither survives the move to this task: graph-level regression needs every node in one forward pass (2.1.4), and the reduction under study must be a deterministic, cacheable artifact computed once, not redrawn every epoch.

Partitioning and Distributed Training

The standard industrial answer is to partition the graph across devices and exchange messages at the boundaries, or drop them. Cutting edges is then forced rather than chosen, which is why partitioning is carried as a reduction family here (2.1.5) despite removing no nodes: the accuracy cost of a severed spanning edge is exactly what a distributed practitioner pays, measured on one device. DistDGL [Zheng et al., 2020] represents the practice at billion-node scale.

Memory-Efficient Training Techniques

Gradient checkpointing, mixed precision and gradient accumulation, defined in 2.1.4, are used here rather than studied here, but a reader will still ask why reduction is needed when checkpointing exists. The answer is composition. Each technique trades compute, precision or latency for memory at a fixed input size; reduction changes the input size itself, and the two stack. The baseline configurations depend on the first fact, DeepGate4 being runnable only under checkpointing (3.1.6), and the contribution depends on the second.

2.2.3 Graph Reduction Techniques

Each family below is fixed by the quantity its operator controls, and the methods inside a family differ only in how that quantity is chosen.

Graph Partitioning

A k -way partition $\pi: \mathcal{V} \rightarrow \{1, \dots, k\}$ deletes every edge whose endpoints land in different blocks, so a partitioner controls the cut and nothing else. The reference objective is the balanced minimum cut,

$$\min_{\pi} |\{(u, v) \in \mathcal{E} : \pi(u) \neq \pi(v)\}| \quad \text{subject to} \quad \max_i |\pi^{-1}(i)| \leq (1 + \varepsilon) \quad (2.28)$$

with imbalance tolerance ε . The problem is NP-hard, and METIS is the standard approximation to it [Karypis and Kumar, 1998]: the graph is coarsened by successive matchings, the smallest graph is partitioned directly, and the assignment is projected back with local refinement at each level. Every cut edge in (2.28) counts one, whatever it connects. That indifference is the only term the domain-aware variant of 3.2.1 changes.

Two further constructions are references rather than competitors. Hash partitioning draws $\pi(v)$ uniformly and independently of $\mathcal{E}$, which is the default in distributed GNN systems, where partition quality is traded against the cost of computing it [Zheng et al., 2020]; each edge then survives with probability $1/k$, and a structure-aware objective earns its cost only by retaining more than that fraction. Ordering-based slicing sorts

$\mathcal{V}$ by a scalar the graph already carries and cuts the sorted sequence into k equal blocks, so balance is exact and the cut follows entirely from the ordering (3.2.1).

Graph Sparsification

Sparsification keeps a subset $\mathcal{E}' \subseteq \mathcal{E}$, or an induced subgraph on a node subset $\mathcal{V}' \subseteq \mathcal{V}$, and the methods differ in what they promise the survivor preserves. The strongest such promise available at near-linear cost is a bound on stretch. A t -spanner is a subgraph $H \subseteq G$ with

$$d_H(u, v) \leq t \cdot d_G(u, v) \quad \text{for every } u, v \in \mathcal{V}, \quad (2.29)$$

so an edge may be discarded only when an alternative path of length at most t survives [Peleg and Schäffer, 1989], and randomised constructions reach the optimal size-stretch trade-off in linear time [Baswana and Sen, 2007]. The guarantee is paid for out of dense short-range redundancy. A graph that lacks it returns a spanner nearly as large as itself, which is the negative result reported in 3.2.2.

The remaining methods trade the guarantee for a knob or for simplicity. A minimum-weight spanning forest keeps $|\mathcal{V}| - c$ edges on c components [Kruskal, 1956]: a connectivity backbone with every cycle destroyed, at a compression that is a property of the graph rather than a setting. Independent edge removal keeps each edge with probability $1 - p$ and was introduced as a per-epoch regulariser for deep GNNs [Rong et al., 2020]; a mask drawn once instead of per epoch is a reduction, and its ratio is freely tunable, which is what lets it serve as a matched-compression control. Centrality pruning keeps the $\lceil \kappa |\mathcal{V}| \rceil$ highest-scoring nodes under PageRank [Brin and Page, 1998, Langville and Meyer, 2003] and returns the induced subgraph, so every edge incident on a discarded node goes with it, and κ is again free. The two ratios of (2.23) separate at exactly this point: the first two methods hold $\eta_{\mathcal{V}} = 1$, the last moves both.

Graph Summarization and Coarsening

Coarsening merges nodes into super-nodes under the merge map of 2.1.5, and the strands differ in what they treat as grounds for merging two nodes.

The exact strand merges only nodes that are provably indistinguishable. Colour refinement (2.15) computes the coarsest equitable partition [Grohe et al., 2014], and Bollen et al. [2023] turn that partition into a training procedure: a graph is quotiented by its d -step refinement classes, and a network trained on the quotient produces exactly the output it would have produced on the original. Exactness is conditional, and the conditions are what make the result usable. The model must be message passing of depth at most d over an aggregator that reads the multiset rather than the set of incoming messages, so that neighbour counts still reach it. The graph must be presented as a multigraph whose super-edges carry those multiplicities, and its initial features must be constant on each initial colour class. A count cap c truncates the multiset, with $c = 1$ collapsing it to a set and giving bisimulation: coarser, and no longer exact for a counting model. Section 2.1.5 states the condition

formally. Both requirements are met by construction here, since the encoder sums over neighbours and the summarized schema records super-edge multiplicities (3.2.3), so refinement at $c = \infty$ is lossless for this encoder and is the study’s positive control rather than one more contender (3.1.5). k -SNAP groups by attribute and neighbour class [Tian et al., 2008] and is the depth-one case of the same construction, so it is cited and not run.

The spectral and classical strand descends from scientific computing. Local Variation coarsening contracts node sets under a restricted spectral approximation bound [Loukas, 2019], the lineage is surveyed by Chen et al. [2022], and Kron reduction has a directed extension [Sugiyama and Sato, 2023]. Its guarantees are stated for the spectrum of a symmetric Laplacian, so none of them survives the move to a directed graph unexamined, and the strand is surveyed here rather than run.

Convolution matching is the GNN-aware strand [Dickens et al., 2024]: nodes merge when the graph convolution itself cannot separate them, rather than when some graph property cannot, and about 95% of full-graph performance is reported at 1% of the node count on node classification. Merging under the operator the model actually applies is the strongest general-purpose criterion published, which makes it the bar a domain-specific method has to clear (3.2.3). Universal Graph Coarsening merges nodes that collide under locality-sensitive hashing, in time linear in the size of the graph [TODO, 2024], and occupies the cheap end of the same axis. Shabani et al. [2023] survey the use of both strands with GNNs.

Uniform merging to a stated target is the control the rest are read against. Compression is a free parameter for only some coarseners, so no general procedure makes two real methods meet at one ratio; a random merge reaches any ratio exactly, and each method is therefore compared against a random control run at that method’s own achieved compression (3.2.3).

Condensation is the family this thesis excludes (1.2.6). GCond and its descendants synthesise a new small graph by gradient matching instead of merging the real one [Jin et al., 2022], so (i) no correspondence exists between synthetic and real nodes, (ii) the method is label-dependent by construction, (iii) gradient matching ties the result to one architecture, and (iv) with no real nodes there is nothing to run cross-state inference on.

Reduction as Preprocessing for GNN Training

Two systems already train on a reduced graph and then serve the original. SCAL coarsens the training graph and applies the same weights to the full graph at inference [Huang et al., 2021], the shared-weights mechanism RQ5 tests. Generale et al. [2022] instead transfer weights back through the node-to-super-node mapping for a node-level task on knowledge graphs, and their ablation finds the *content* carried on a super-node to be the deciding factor, which is why member counts and level statistics travel with super-nodes here (3.2.3). The distance between the two is the distance between their tasks. A node-level prediction needs a mapping back to individual nodes, whereas a graph-level regression needs only an input space that both structural states inhabit (3.2.3), so shared weights suffice and no transfer step exists to go wrong.

2.2.4 Structure-Preserving and Domain-Aware Reduction

Reduction on Directed Acyclic Graphs

Direction changes more than notation. DAG-native architectures process nodes in topological order and read direction as causality [Thost and Chen, 2021, Zhang et al., 2019]. On the reduction side only isolated pieces exist, such as the directed extension of Kron reduction [Sugiyama and Sato, 2023]. Coarsening theory is otherwise stated for undirected graphs: its guarantees concern spectra of symmetric Laplacians, and none of them carry to a quotient of a DAG, which need not even be acyclic (2.1.2). No guarantee therefore transfers, and every method run here is compared against a random control at matched compression (3.2.3) instead of being trusted on a guarantee stated for a different object.

Reduction in Circuit and Netlist Domains

AIG-native reduction predates the summarization literature and ignores it. Structural hashing merges gates that agree on the fanin signature (2.11) at the moment of construction [Kuehlmann et al., 2002]. FRAIG merges functionally equivalent gates, refuting candidates by simulation and proving the survivors by SAT [Mishchenko et al., 2005]. Maximal fanout-free cone decomposition (2.8) is how refactoring already carves a network into mergeable cones [Mishchenko et al., 2006]. Contracting a chain of single-fanout gates is its chain-shaped special case, and the post-dominator of 2.1.2 [Lengauer and Tarjan, 1979] bounds a cone whose interior reconverges. Read through the lens of 2.1.5, all three are equivalence-based coarsenings under another name, and connecting them to the exactness framework of 2.1.4 is part of the contribution claimed here.

That connection carries a hazard the generic coarsening literature never meets. The label regressed here is the fraction of nodes synthesis removes, and the reductions just listed are themselves synthesis operations, so a coarsener drawn from that family performs part of the optimization it is supposed to predict and shrinks the quantity being measured. The hazard is graded rather than binary. Merging by functional equivalence (2.12) is what a synthesis pass does. Merging by cone structure is what a synthesis pass operates on. Merging by an invariant the encoder cannot see touches neither. No method run here occupies the first position, and the domain-specific candidate is the only one at the second, which is why it alone has to argue that intra-cone wiring is noise with respect to the label (3.2.3, 5.4.5).

CTS-Bench [Khadka et al., 2026] is the nearest published competitor. It benchmarks coarsening trade-offs for GNNs on post-placement gate-level netlists for clock-skew prediction, with a single bespoke heuristic that consumes physical coordinates, so it cannot reach a pre-physical AIG at all (2.1.1), and it compares no second method. Its findings support the premise of this thesis rather than competing with it: up to $17.2\times$ memory reduction and $3\times$ training speedup, accuracy collapsing to negative R^2 under zero-shot cross-state evaluation, and an explicit call for domain-aware coarsening. Two of its numbers are carried forward. The negative zero-shot R^2 is the standing warning against assuming that a model trained on reduced graphs

will serve full ones, and the dissociation between its error metrics, MAE nearly unmoved at 0.16 to 0.17 while R^2 fell below zero, is the citation behind reporting RMSE, R^2 and Spearman correlation together rather than any one alone (3.4.3).

2.2.5 Size Generalization and Cross-State Transfer

Size and Distribution Shift in GNNs

Message-passing weights are shared across nodes and layers, so a trained model accepts a graph of any size, a point the inductive line established in principle [Hamilton et al., 2017]. Accepting an input is not generalising to it. Buffelli et al. [2022] measure accuracy falling as test graphs grow past the sizes seen in training, and regularise toward size-invariant representations to recover part of the loss. A reduction moves size and structure together, so the shift it induces between training and inference is bounded below by the one they report.

Train-Test Structural Mismatch

The direct evidence is thin and it points both ways. SCAL trains on coarsened graphs and evaluates on the originals with usable node-level accuracy under a single set of weights [Huang et al., 2021]. Generale et al. [2022] reach the same conclusion on knowledge graphs, but only after mapping super-node representations back onto their members, a step no graph-level target provides. CTS-Bench runs the analogous evaluation on netlists and lands below $R^2 = 0$ (2.2.4), the one published measurement on circuit data [Khadka et al., 2026]. One success on a different task, one success behind a transfer step, and one failure in the closest domain leave the question open on graph-level circuit regression. A study that asks it therefore needs a compression whose losslessness is established in advance (3.1.5); without one, a poor result across structural states cannot be separated from a broken evaluation path.

2.2.6 Research Gap

The two halves surveyed above do not meet. The openings are specific, and each is narrower than a claim that nobody has coarsened a circuit graph.

- **Unseen designs are seldom the tested condition.** Prediction on designs the model has never seen is the field’s stated motivation. Matched seen and unseen-design numbers appear in 4 of the 63 papers surveyed here (2.2.1). Under a random or recipe-disjoint split every test design also appears in training under another recipe, so an unmeasured share of the reported score is design recognition rather than structural inference. The standard design-disjoint split on the reference corpus trains on 16 small designs and tests on 8 large ones [Chowdhury et al., 2021], which confounds novelty with size, and a model motivated on the gap need not publish the seen-design number its degradation would be measured against [Deng et al., 2024].

- **Circuit models already reduce their input and never measure it.** The representation-learning line trains on extracted subcircuits of tens to a few thousand gates [Li et al., 2022, Zheng et al., 2025]. Extraction is a reduction, applied at corpus construction, and it is neither named as one nor evaluated as one (2.2.1). What it costs is unknown.
- **Graph reduction has not been evaluated on AIGs.** The taxonomy is developed and surveyed [Hashemi et al., 2024, Shabani et al., 2023], and no method in it has been measured against a logic-synthesis target. Transfer cannot be assumed, for a mechanical reason: a generic method cuts or merges under an undirected criterion, and a single severed edge removes gates from the fanin cone (2.6) of everything downstream of it (2.1.2).
- **Coarsening results in this domain carry no compression-matched control.** CTS-Bench reports memory and speed gains for one bespoke heuristic and compares no second method [Khadka et al., 2026]. Absent a control at the same ratio, a gain cannot be attributed to the merge criterion rather than to the compression.
- **The two equivalence literatures are not connected.** Structural hashing, FRAIG and maximal fanout-free cone decomposition are coarsenings by merge map [Kuehlmann et al., 2002, Mishchenko et al., 2005, 2006], and the exactness framework that decides which of them an encoder can absorb without loss [Bollen et al., 2023] has never been turned on them.

After CTS-Bench, being first to coarsen in EDA cannot be claimed unqualified, and the claim made here is narrower. This thesis is the first to (i) benchmark multiple coarsening families, exact refinement-based, GNN-aware, and domain-specific, against a compression-matched random control for GNN training in EDA, (ii) do so on AIGs and optimizability regression, (iii) bring provably lossless compression into an EDA GNN setting, (iv) connect AIG-native equivalence reduction to the generic summarization framework, and (v) run a controlled three-protocol split comparison with per-design error distributions on this task. Spectral and hashing-based coarsening are surveyed in 2.2.3 but not run, and are no part of the claim. The eventual results are compared against SynthNet [Chowdhury et al., 2021], HOGA [Deng et al., 2024] and DeepGate4 [Zheng et al., 2025] on the prediction task, and against CTS-Bench [Khadka et al., 2026] on the shape of the memory, speed and accuracy trade-off.

Methodology

3.1 Architecture

[TODO: make more concise also to evaluate what to put in the appendix and what to keep here write that as a permanent note so my thesis supervisor can see so not a ocmment] The encoder follows GNN+ [Luo et al., 2025]; hyperparameters come from the literature and from the tuning sweep of 3.4.2. One architecture is used for unreduced, sparsified, partitioned and summarized graphs alike, with two exceptions stated below: partitioning changes the pooling path (3.1.2), and the exact-compression track changes the input schema (3.1.5).

The encoder is an edge-aware graph convolutional network over AIGs. A node carries a one-hot type indicator over [constant, PI, AND, PO], so an AND gate is [0, 0, 1, 0], and an edge carries a one-hot indicator of whether it is inverted. The one departure from that schema is the summarized track, where a node is a super-node and its vector counts the members it absorbed: a super-node holding 2 primary inputs, 15 AND gates and 1 primary output is [0, 2, 15, 1], and a super-edge counts the inverted and non-inverted wires merged into it. A one-hot vector is the count vector of a single-member super-node, so both graph types occupy the same four-dimensional input space and are read by the same encoder weights (3.2.3). The remaining settings are given in 3.1 and derived in 3.1.1.

Table 3.1: Encoder configuration, used unchanged for unreduced, sparsified, partitioned and summarized graphs. The two exceptions are the pooling path under partitioning and the input schema of the exact-compression track, described in Sections 3.1.2 and 3.1.5.

Component	Setting
Node features	4-dimensional type indicator (member counts when summarized)
Edge features	2-dimensional inversion indicator
Positional encoding	Log logic level, projected to 32 dimensions
Input width	160 (128-dimensional node projection concatenated with edge features, normalised over each graph before the first layer)
Depth	4 message-passing layers
Hidden width	128, constant across layers
Activation	LeakyReLU in the encoder blocks, ReLU in the regression head
Feed-forward block	128 → 256 → 128
Normalisation	Layer normalisation in graph mode, bracketing the feed-forward block
Residuals	Skip connection added after message passing in every layer
Dropout	0.15 in the encoder blocks, 0.3 in the regression head
Degree normalisation	Disabled
Readout	Jumping knowledge summed over the 4 layers with mean pooling
Regression head	128 → 64 → 1 with a terminal sigmoid

3.1.1 Encoder Formulation

Nothing in this subsection is a contribution of this thesis. The departures below are GNN+’s [Luo et al., 2025], each cited to the equation of that paper it comes from; they are written out

because the reduction claims of 3.1.4 and 3.1.5 are claims about *this* forward pass, and cannot be checked against a bulleted configuration. The settings specific to this work are collected at the end of the subsection. The classic GCN layer [Kipf and Welling, 2017] that GNN+ modifies is

$$h_v^l = \sigma \left(\sum_{u \in \mathcal{N}(v) \cup \{v\}} \frac{1}{\sqrt{\hat{d}_u \hat{d}_v}} h_u^{l-1} W^l \right). \quad (3.1)$$

Throughout, h_v^l is the representation of node v after layer l , e_{uv} the attribute of edge (u, v) , σ the activation, $\parallel$ concatenation, and biases are suppressed for readability.

Edge-conditioned messages. Each edge attribute is projected by its own weight W_e^l and fused with the neighbour representation, so the message a node receives depends on the wire it arrived on:

$$m_v^l = \sum_{u \in \mathcal{N}(v)} c_{uv} \sigma(h_u^{l-1} W^l + e_{uv} W_e^l). \quad (3.2)$$

Placing σ on each message rather than on the aggregate follows the GNN+ reference implementation; its Eq. 6 writes the activation outside the sum. The per-edge coefficient c_{uv} is the one place (3.2) is configurable rather than fixed. With symmetric degree normalisation enabled it is the GCN coefficient $1/\sqrt{\hat{d}_u \hat{d}_v}$ of (3.1), computed once per edge when graphs are cached rather than at every forward pass; disabled, it is 1 and no edge-weight tensor is carried at all. It is disabled here. The same edge-weight channel is what the exact-compression track of 3.1.5 reuses: setting c_{uv} to the edge multiplicity w_{uv} scales the message *outside* σ , which is what makes one merged message equal the sum of the messages it replaces.

Normalisation, dropout, residual and feed-forward block. GNN+ adds four of its six techniques inside the layer block:

$$\tilde{h}_v^l = h_v^{l-1} + \text{Drop}(\sigma(m_v^l)), \quad (3.3)$$

$$\bar{h}_v^l = \text{Norm}(\tilde{h}_v^l), \quad (3.4)$$

$$\bar{h}_v^l = \text{Norm}(\bar{h}_v^l + \text{FFN}(\bar{h}_v^l)), \quad (3.5)$$

$$\text{FFN}(h) = \text{Drop}(\text{Drop}(\sigma(h W_1^l)) W_2^l), \quad (3.6)$$

with W_1^l : 128 → 256 and W_2^l : 256 → 128. The published layer carries one normalisation that (3.3) does not. Its Eqs. 7–9 normalise the aggregated message *before* the activation and before the residual is added, and its Eq. 10 closes the FFN by normalising that block’s own residual sum; the GNN+ reference implementation keeps both and inserts a third ahead of the FFN, the bracketing arrangement of GraphGPS [Rampásek et al., 2022]. Only the message normalisation is dropped here. The two dropouts of (3.6) likewise follow the reference implementation, which Eq. 10 does not show.

Positional encoding. A per-node structural scalar, here the log-compressed logic level ℓ_v , is projected and concatenated to the node features rather than summed into them:

$$h_v^0 = \text{GraphNorm}(x_v W_{\text{in}} \parallel \sigma(\ell_v W_{\text{PE}})), \quad (3.7)$$

with $x_v \in \mathbb{R}^4$, $W_{\text{in}}: 4 \rightarrow 128$ and $W_{\text{PE}}: 1 \rightarrow 32$, so $h_v^0 \in \mathbb{R}^{160}$. GNN+ holds its concatenation at the model width, projecting the joined vector back down in Eq. 12 and narrowing the node projection beforehand in its implementation. Keeping all 160 dimensions here means the first layer maps $160 \rightarrow 128$ and is therefore the one layer whose residual in (3.3) is skipped, its input and output widths being unequal.

Configuration specific to this work. The encoder is GNN+; the following are this project’s configuration of it, and none of them change the technique:

- **Degree normalisation off.** $c_{uv} = 1$ in (3.2), where GCN+ carries the $1/\sqrt{\hat{d}_u \hat{d}_v}$ of (3.1) into its own layer. Fan-in in an AIG is fixed by node type, so the coefficient, taken over fan-in and without the self-loop, is constant within each type pair and is zero on every edge leaving a primary input. GNN+ supplies no evidence either way: its ablations [Luo et al., 2025, Tables 5–6] cover the six techniques it adds and not the GCN coefficient it inherits, and its own edge-aware layer disables both that coefficient and the self-loop, so this setting agrees with the reference implementation where the implementation departs from Eq. 6. The normalised variant is implemented and reachable through the model configuration, not exposed as a training flag.
- **No self-loop.** Aggregation in (3.2) runs over $\mathcal{N}(v)$ without the $\cup\{v\}$ of (3.1); a node’s own state re-enters through the residual of (3.3) only.
- **LeakyReLU for σ** in the encoder, and layer normalisation in graph mode for Norm, taking statistics over all nodes of one graph, rather than the batch normalisation used throughout the GNN+ ablations.
- **Level-based positional encoding** in (3.7), a single learned projection of log logic depth, where GNN+ uses a Random Walk Structural Encoding (RWSE). Level comes from a single topological pass, where RWSE costs one sparse matrix product per walk step on graphs of this size and adds the walk length to a tuning budget already committed to the reduction parameters.
- **Readout and head.** GNN+ reads out from the last layer alone; here the layer outputs are summed first (jumping knowledge over all $L = 4$ layers), then mean-pooled and passed through a bounded head:

$$h_G = \frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} \sum_{l=1}^L h_v^l, \quad (3.8)$$

$$\hat{y} = \text{Sigmoid}(\text{Drop}_{0.3}(\text{ReLU}(h_G W_3)) W_4), \quad (3.9)$$

with $W_3: 128 \rightarrow 64$ and $W_4: 64 \rightarrow 1$. The terminal sigmoid matches the $[0, 1]$ range of the optimizability target, and is the reason the head uses ReLU and a heavier dropout ($p = 0.3$) than the $p = 0.15$ of the encoder blocks.

3.1.2 Partitioned Inputs

Partitioning is the one reduction family that requires an architectural change. Graphs are partitioned offline and the edges spanning partitions are removed, which leaves the node set intact but disconnects the graph into k components. Pooling directly to the graph level would then average over components the encoder never allowed to communicate, so the readout is made hierarchical:

- GNN encoder (applied on intra-partition edges only)
- Mean pooling from node representations to partition representations
- Mean pooling from partition representations to a single graph representation
- Prediction head applied to the graph representation to produce the output $\hat{y}$

Feature normalization is also made partition-aware: the input GraphNorm normalizes within each partition rather than within each graph, so statistics do not leak across a boundary the message passing itself cannot cross.

3.1.3 Sparsified Inputs

Sparsification alters only the data pipeline: a sparsified AIG preserves the feature schema of the original, so the encoder of 3.1 ingests it unmodified.

Reductions are stored as node and edge masks applied at load time rather than materialised into the corpus, so a single corpus serves every sparsification setting. Edge masks (random edge dropout, spanning forest) leave the node count unchanged, whereas node masks (PageRank pruning, AND-gate-only) reduce it. The node budget governing batch assembly (3.4.2) is therefore measured on the post-mask count.

3.1.4 Summarized Inputs

Summarized graphs are also fed to the unmodified encoder, which is possible only because the merge preserves the input schema as a superset: node and edge features become member counts rather than one-hot indicators, and a super-node containing a single node reproduces its original feature row exactly (3.2.3). Positional encodings are pooled by member minimum, so a super-node enters the circuit at the level of its earliest member.

A full graph is consequently a valid input to a summary-trained model without preprocessing, being a graph of size-1 super-nodes. The cross-state inference of RQ5 therefore evaluates one set of weights on both representations.

3.1.5 Encoder Variant for Exact Colour Refinement

Colour refinement at count-cap ∞ (3.2.3) is exact rather than approximate, but the guarantee holds only under a modified encoder.

When k nodes merge, parallel edges collapse into one super-edge whose message must equal the sum of the k it replaces. Multiplicity w_{uv} carried as an edge attribute enters σ in (3.2), which does not decompose over sums. Setting $c_{uv} = w_{uv}$ instead scales the message outside σ and gives equality. Edge attributes are therefore dropped, and inversion moves onto the nodes: fan-in is fixed by node type in an AIG (zero for constants and primary inputs, two for AND gates, one for primary outputs), so a node’s type and its count of inverted incoming edges determine which fan-ins are inverted, up to order. Which specific fan-in is inverted is lost, which a graph-level target does not need.

Coarsening to depth d equal to the number of message-passing layers then yields outputs identical to the uncoarsened graph [Bollen et al., 2023]: compression that costs no accuracy by construction, and the positive control for RQ5 (3.4.1). The variant pays for this in schema compatibility. A full graph must be converted before this encoder can be queried on it, whereas the other methods support cross-state inference with no such step.

[TODO: go over this when comparins to exact sumam-rizaiont related work]

3.1.6 Baselines for Optimizability Prediction

RQ1 is a comparative claim and needs a reference. Two groups supply one: trivial predictors, which fix what a given R^2 is worth on this label, and three published circuit models ported to this task. A third group of standard graph-level encoders (GCN, GraphSAGE, GIN) was considered and may be added if time permits.

Trivial Predictors

Three predictors, each fitted on the validation split and scored on test, so that every baseline is scored on data it has not seen. Two are constant,

$$\hat{y}_i^{\text{mean}} = \bar{y}_{\text{val}}, \quad \hat{y}_i^{\text{med}} = \text{med}(y_{\text{val}}), \quad (3.10)$$

and the third is an ordinary least squares fit on graph size alone,

$$\hat{y}_i^{\text{size}} = \text{clip}_{[0,1]}(\beta_0 + \beta_1 \log_{10} |\mathcal{V}(G_i)| + \beta_2 \log_{10} |\mathcal{E}(G_i)|). \quad (3.11)$$

Equation (3.10) fixes the zero point of (3.32). The test-set mean is the least-squares optimal constant, so any other constant scores $R^2 \leq 0$, with equality only at $\bar{y}_{\text{val}} = \bar{y}_{\text{test}}$. Equation (3.11) is the weakest non-trivial hypothesis about the task, that optimizability is a function of size alone, and any claim that the encoder reads circuit *structure* has to clear it.

[TODO: evaluate which of these i will actualy do shoudl be the base scores the means and all that for comparatability and compared to random and CIs] [TODO: is ranodmly intialized GNN a baseline? i think so but i need to check the literature]

Adapting Published Circuit Models

SynthNet [Chowdhury et al., 2021], HOGA [Deng et al., 2024] and DeepGate4 [Zheng et al., 2025] are introduced in 2.2.1. None was published for this task, so each had to be adapted, and what follows states those adaptations rather than the models.

Shared porting procedure. Each paper supplies an encoder g_ϕ carrying an AIG to node states $\{h_v\}_{v \in \mathcal{V}(G)}$, together with a head supervising that paper’s own target. The encoder is reused unmodified and the head is replaced, so every port has the form

$$\hat{y}(G) = \text{Sigmoid}(\text{MLP}(\pi(\{h_v\}_{v \in \mathcal{V}(G)}))), \quad (3.12)$$

with π a permutation-invariant readout and the terminal sigmoid matching the $[0, 1]$ range of the target of 2.1.6. Only π , the head and the input adapter differ across the three, and the trunk of (3.12) is in every case the published one.

Whether π is inherited or chosen differs by model, and the distinction bears on what each comparison licenses. SynthNet already predicts a graph-level quantity, so π is its own concatenated maximum and mean pooling. HOGA classifies nodes and DeepGate4 pools over cones; neither defines a whole-circuit readout, so for both π is mean pooling and is this work’s choice, not the paper’s.

Adaptations required by a graph-level target. Three, each a consequence of the target being one scalar per circuit under one fixed script:

- **The recipe branch is removed.** SynthNet predicts $\hat{y}(G, s)$, conditioning on a synthesis recipe s through a separate encoder. Training here fixes the script, so s is constant across the corpus and that branch emits one vector for every sample, contributing a bias to the head and nothing else. It is dropped rather than left to learn that constant.
- **Batches are bounded by nodes, not graphs.** A graph-level readout requires all of $\mathcal{V}(G)$ in one forward pass, so a graph cannot be split across batches, and neither published batch size transfers: HOGA minibatches nodes drawn from one graph, DeepGate4 minibatches cones. Batches are assembled to a node budget B and paired with an accumulation count a , and it is the product that matters. One graph contributes one label regardless of its size, so the number of losses averaged per update is $a \cdot \mathbb{E}[B/|\mathcal{V}(G)|]$, held near the primary model’s value. Budget and accumulation are tuned and reported as a pair.
- **Activations are recomputed rather than retained.** DeepGate4 attends over a virtual edge set $\bar{\mathcal{E}}_k = \{(u, v) : \text{dist}(u, v) \leq k\}$, which at the published radius $k = 8$ measures roughly 180 edges per circuit node against an AIG fanin of at most two. Each of its attention layers materialises a message tensor over $\bar{\mathcal{E}}_k$, so recomputing them during the backward pass, which leaves outputs and gradients unchanged, is what makes the baseline runnable at all.

Corrections to the released implementations. Two departures are corrections rather than adaptations, and are stated because they change published behaviour. HOGA’s released attention normalises its scores over the head axis rather than the key axis, so the weights never form a distribution over hop neighbours; the port normalises over keys, following Eq. 5 of the paper. Its optional reverse-propagation branch applies the forward adjacency to the reverse features, which duplicates the forward branch rather than reversing it, and the port applies the reverse adjacency. A third difference is a convention, not an error:

OpenABC-D orients edges toward the primary inputs and this corpus orients them the other way, so each node summarises its fanout cone under one convention and its fanin cone under the other. Both orientations are run and the pair reported. **[TODO: i will provably not report both just the proper one the correct one that was originally used in paper]**

Comparability. Only the data pipeline is shared, since a different split would make the comparison meaningless. Hyperparameters, optimizer, loss and schedule take each paper’s published values. Settings a paper leaves unpublished are recorded as assumptions in A.3.

Scores printed in those papers are a different matter and cannot be placed beside the ones here. OpenABC-D standardises its targets within each design d , so its reference spread is the within-design sum of squares where the denominator of (3.32) is the total,

$$\underbrace{\sum_i (y_i - \bar{y})^2}_{\text{here}} = \underbrace{\sum_i (y_i - \bar{y}_{d(i)})^2}_{\text{upstream}} + \sum_d n_d (\bar{y}_d - \bar{y})^2. \quad (3.13)$$

The two differ by the between-design term, which on a design-disjoint split (3.4.2) is the variation the task is about. Only the sign of the score transfers (4.1.4). **[TODO: do i need this?]**

[TODO: this should really be mostly on how i adapted my baselines of the task. so recipe was not inputted i cut off that branch]

3.2 Graph Reduction

This section describes the reduction methods evaluated in this thesis. Their objective is to relieve the memory bottleneck that prevents training on unreduced AIGs and to shorten training time, while preserving the structural features on which accurate optimizability prediction depends.

The three families are defined in 2.1.5 and their generic members are surveyed in 2.2.3. What follows states what changes when those methods meet an AIG: the adaptation each one required, and the value every free parameter was given together with what fixes it. The parameters are collected in A.2.

3.2.1 Partitioning

Applied to an AIG $G = (\mathcal{V}, \mathcal{E})$, an assignment $\pi: \mathcal{V} \rightarrow \{1, \dots, k\}$ with blocks $\mathcal{V}_i = \pi^{-1}(i)$ keeps exactly the edges internal to a block,

$$R_\pi(G) = (\mathcal{V}, \mathcal{E} \setminus \mathcal{E}_{\text{cut}}(\pi)), \quad \mathcal{E}_{\text{cut}}(\pi) = \{(u, v) \in \mathcal{E} : \pi(u) \neq \pi(v)\}, \quad (3.14)$$

the disjoint union of the k induced subgraphs. Node retention is $\eta_{\mathcal{V}} = 1$ by construction, so the only quantity a partitioner controls is the edge retention $\eta_{\mathcal{E}} = 1 - |\mathcal{E}_{\text{cut}}(\pi)|/|\mathcal{E}|$ of (2.23). All blocks of a graph are kept and enter the model together, through the hierarchical readout of 3.1.2.

The number of blocks is set per graph from a target block

size s ,

$$k(G) = \min\left\{k_{\max}, \max\{k_{\min}, \lfloor |\mathcal{V}|/s \rfloor\}\right\}, \quad (3.15)$$

with $s = 10,000$ and $[k_{\min}, k_{\max}] = [2, 32]$. The two bounds are fixed by the corpus rather than chosen. The floor binds below $2s = 20,000$ nodes, which covers the median graph of 3.3.6, and it guarantees that every graph is genuinely partitioned instead of passing through unreduced. The cap is reached only above $s k_{\max} = 320,000$ nodes, beyond the largest graph in the evaluation splits, so it acts as a guard against k growing without bound and never as an operating point. The target block size itself was fixed a priori and never swept, which leaves the sensitivity of every partitioning result to s unmeasured.

The four methods below receive the same $k(G)$ on every graph and differ only in the assignment π , which is what makes them directly comparable.

Random Hashing

Each node is assigned a block independently and uniformly at random, $\pi(v) \sim \text{Uniform}\{1, \dots, k\}$, from a fixed seed. The assignment ignores the edge set entirely, so each edge survives with probability $1/k$ and

$$\mathbb{E}[\eta_{\mathcal{E}}(R_\pi(G))] = \frac{1}{k}, \quad (3.16)$$

independently of the graph. This is the floor of the family: a structure-aware method justifies its cost only by retaining more than this fraction. Blocks are balanced only in expectation, but the deviation is negligible at the block sizes implied by (3.15).

METIS

METIS [Karypis and Kumar, 1998] approximates the balanced minimum-cut partition

$$\min_{\pi} |\mathcal{E}_{\text{cut}}(\pi)| \quad \text{subject to} \quad \max_i |\mathcal{V}_i| \leq (1+\varepsilon) \left\lceil \frac{|\mathcal{V}|}{k} \right\rceil, \quad (3.17)$$

by the multilevel heuristic of 2.2.3. It is applied unmodified, at the balance tolerance ε its implementation defaults to, which is what makes it the generic arm of the comparison rather than a method this thesis tuned.

Two consequences of applying it to an AIG matter downstream. It operates on the undirected projection, so edge direction plays no part in where the cuts fall. Its objective also counts every cut edge equally, so severing a connection between adjacent levels costs the same as severing an edge that bridges half the circuit’s depth. That is precisely the assumption the domain-informed variant below modifies.

Level Slicing

Nodes are ordered by topological level (2.5) and the ordering is cut into k blocks of equal cardinality. Writing $\text{rank}(v) \in \{0, \dots, |\mathcal{V}| - 1\}$ for the position of v in a stable sort by $\ell(v)$,

$$\pi(v) = 1 + \left\lfloor \frac{k \cdot \text{rank}(v)}{|\mathcal{V}|} \right\rfloor. \quad (3.18)$$

Balance is exact by construction, and blocks are contiguous in depth: $\ell(u) < \ell(v)$ implies $\pi(u) \leq \pi(v)$. Because every edge runs from a lower level to a strictly higher one, an edge is cut exactly when its endpoints straddle one of the $k-1$ block boundaries. The local directed structure of each depth band therefore survives intact. The price is that every path from the inputs to the outputs crosses every boundary and is severed $k-1$ times. The construction introduces no parameter of its own: k comes from (3.15) and the ordering is fixed by the graph.

Span-Weighted METIS

The domain-aware adaptation changes only the objective of (3.17). Each edge $e = (u, v)$ has a span

$$\text{sp}(e) = \ell(v) - \ell(u) \geq 1, \quad w(e) = 1 + \gamma \cdot \text{sp}(e), \quad (3.19)$$

the number of levels it crosses, and the weighted cut $\sum_{e \in \mathcal{E}_{\text{cut}}(\pi)} w(e)$ is minimised under the same balance constraint. A long edge is a shortcut between distant levels, and cutting it removes gates from the fanin cones (2.6) of everything downstream, which is exactly the damage the uniform objective is indifferent to. Under (3.19) the cost of a cut grows linearly with the length of the chain it breaks, so the minimiser pushes cuts toward span-one edges and the longest causal chains stay whole.

The weight $\gamma = 10$ follows from a limit rather than from a sweep. A minimiser is unchanged by a common scaling of all weights, so as γ grows, w and γsp pose the same problem, and that limiting problem is minimisation of the total span severed. Every weight sits within $1/(1 + \gamma \text{sp}(e)) \leq 1/11$ of the limit at $\gamma = 10$, so the objective is already a small perturbation of pure span minimisation and larger values would move the partition little. Smaller values would not: at $\gamma = 1$ the unit term is half the cost of a span-one cut.

3.2.2 Sparsification

Two of the four methods below act on the edge set alone, replacing $\mathcal{E}$ with a subset $\mathcal{E}' \subseteq \mathcal{E}$ and leaving $\eta_{\mathcal{V}} = 1$. The other two select a node subset $\mathcal{V}' \subseteq \mathcal{V}$ and return the induced subgraph $R(G) = G[\mathcal{V}']$, so every edge with an endpoint outside $\mathcal{V}'$ is removed as well. That difference determines which retention figure is meaningful for each method, and it is why 2.1.5 insists on reporting node and edge retention separately rather than as one compression number. Only one of the four has a compression parameter at all; for the other three the achieved retention is a measurement, reported in 4.2.1 rather than here.

Random Edge Dropout

This naive baseline uniformly drops a fixed random fraction of edges across the graph, acting as a control for structural degradation. Each edge draws an independent uniform variate and survives when the draw clears the drop rate p ,

$$\mathcal{E}' = \{e \in \mathcal{E} : U_e \geq p\}, \quad U_e \stackrel{\text{i.i.d.}}{\sim} \text{Unif}[0, 1], \quad (3.20)$$

so each edge is kept independently with probability $1 - p$ and $\mathbb{E}[\eta_{\mathcal{E}}] = 1 - p$. Random edge removal is established as a per-epoch regulariser for deep Graph Neural Network (GNN) training [Rong et al., 2020]. Here the mask of (3.20) is instead drawn

once and fixed, which makes it a reduction rather than a regulariser.

The drop rate is the only freely tunable compression parameter in the family, so its value decides which parameter-free method it can be matched against, and it was set for a pairing other than the one it is used in. It is fixed at $p = 0.3$, chosen to place $\mathbb{E}[\eta_{\mathcal{E}}]$ near the edge retention of the AND-gate method of 3.2.2. The pairing it actually serves is against the spanning forest of 3.2.2, which requires $p = 1 - \eta_{\mathcal{E}}^{\text{forest}} \approx 0.419$ instead. The reported runs use 0.3, so that pairing removes visibly less than its partner and no accuracy difference inside it can be read as a difference between the two methods (4.4.1, 4.39).

AND-Gate Retention

This domain-specific method removes the circuit interface and keeps the logic. In the node role notation of (2.3),

$$\mathcal{V}' = \{v \in \mathcal{V} : \tau(v) \notin \{\text{PI}, \text{PO}\}\}, \quad (3.21)$$

so all Primary Input (PI) and Primary Output (PO) nodes are dropped with every edge that touches them, and the internal AND network survives intact. The rationale is that the PI/PO boundary is fixed by the circuit's interface and cannot be restructured by synthesis, so the AND network carries the optimizable structure. The method is parameter-free. Its node retention is exactly the AND-gate share of the circuit, a corpus statistic (3.3.6) rather than a value that can be dialled, and it varies from graph to graph.

Random Spanning Forest

This method extracts a sparse connectivity backbone. Let $\tilde{G}$ be the undirected projection of the AIG, and assign each edge an independent weight $w_e \sim \text{Unif}[0, 1]$. The method keeps the minimum-weight spanning forest

$$F = \arg \min_{F' \in \mathfrak{F}(\tilde{G})} \sum_{e \in F'} w_e, \quad |\mathcal{E}(F)| = |\mathcal{V}| - c, \quad (3.22)$$

where $\mathfrak{F}(\tilde{G})$ is the set of spanning forests of $\tilde{G}$ and c the number of connected components, computed with Kruskal's algorithm [Kruskal, 1956]. The weights are almost surely distinct, so each draw has a unique minimiser and the randomised weights select a random forest rather than a deterministic one. A forest contains no cycles. Every reconvergent path therefore loses an edge, while connectivity is kept intact. Because F is computed on $\tilde{G}$, edge direction plays no part in which edges survive. That is a real loss of information on a DAG, and one reason to expect this method to damage the predictive signal.

The edge count in (3.22) also makes the method parameter-free: $|\mathcal{V}| - c$ edges is a property of the graph rather than a setting, so its retention follows from the density of the corpus (4.2.1).

This method replaced a stretch-bounded spanner [Peleg and Schäffer, 1989], which was implemented and abandoned. A t -spanner keeps a subgraph $H \subseteq \tilde{G}$ with $d_H(u, v) \leq t \cdot d_{\tilde{G}}(u, v)$ for every node pair, so an edge can be discarded only when an alternative path of length at most t survives. Spanners exploit dense short-range redundancy, and AIGs have little of it: the

reconvergence they are rich in runs through long paths, so at stretch $t = 3$ the randomised construction of [Baswana and Sen \[2007\]](#) returned the graph nearly unchanged. That stretch is the smallest the construction admits above the exact case $t = 1$, so the failure is not one a weaker setting could have avoided. The negative result is worth reporting: it is a small, concrete instance of the thesis’s general claim that generic graph machinery does not transfer to AIGs unexamined.

PageRank Node Pruning

This centrality-based approach scores every node by PageRank and retains only the highest-scoring fraction, implicitly dropping all edges incident to the removed low-centrality nodes. The score is the stationary solution of

$$\text{pr}(v) = \frac{1 - \alpha}{|\mathcal{V}|} + \alpha \sum_{u \in \text{fi}(v)} \frac{\text{pr}(u)}{|\text{fo}(u)|}, \quad (3.23)$$

computed on the directed graph with fanin and fanout as in (2.2) and the mass of the POs, whose fanout is empty, redistributed uniformly [\[Brin and Page, 1998, Langville and Meyer, 2003\]](#). The kept set $\mathcal{V}'$ collects the $\lceil \kappa |\mathcal{V}| \rceil$ nodes of highest score. The damping factor is left at the $\alpha = 0.85$ of the original formulation [\[Brin and Page, 1998\]](#) and was not tuned, since it governs how far centrality propagates and not how much the method removes.

The retention fraction κ is the method’s compression knob and is set to $\kappa = 0.8$, calibrated to the AND-gate share of 3.3.6 so that the two form a generic and domain-informed pair at approximately equal node retention (4.4.1). The equality is weaker than the shared figure suggests. Here κ fixes node retention exactly on every graph, whereas the AND-gate share varies by circuit, so the pair meets at the corpus mean and not graph by graph.

3.2.3 Summarization and Coarsening

A coarsening is a quotient. A merge map $\mu: \mathcal{V} \rightarrow \mathcal{V}'$ sends each node to the super-node containing it, and the reduced graph is the quotient of G under μ (2.1.5). Unlike the other two families, which only delete, a coarsening must also decide what a super-node is. The rewrite that turns a merge map into a graph is therefore fixed once (3.2.3), and every candidate below reduces to its choice of μ .

Four candidates span domain-specific, exact and general-purpose, with a random control that is what makes matched-compression comparison possible (Table 3.2). No summarization configuration has been trained, so what follows states the intended construction rather than a measured method.

Merge Map and Super-Node Schema

One rewrite turns any merge map into a graph, so that methods differ in what they merge and in nothing else. A super-node carries the counts of its members by role, and a super-edge the

Table 3.2: The summarization candidates, the equivalence each merges by, and whether its compression can be set. Only a method taking a target ratio can be matched to another directly; the rest are matched through the random control of 3.2.3.

Method	Merge criterion	Role	Control
Cone or MFFC contraction	fanout-free cones	domain-specific	graph
Graded colour refinement	d -round colour classes	exact anchor	graph
Convolution matching	convolution equivalence	general bar	target
Random within-type	uniform within type	control	exact

counts of the normal and inverted connections it stands for,

$$x_{u'} = \left(|\{ u \in \mu^{-1}(u') : \tau(u) = r \}| \right)_{r \in \{\text{const}, \text{PI}, \text{AND}, \text{PO}\}}, \quad a_{u'v'} = \quad (3.24)$$

with τ the role of (2.3) and inv the inversion of (2.4). Super-edges are kept only for $u' \neq v'$, so intra-cluster edges become self-loops and are dropped; their count is retained as a graph attribute rather than lost silently. The two entries of $a_{u'v'}$ sum to the multiplicity $w_{u'v'}$ of (2.18). Positional encodings are pooled by member minimum, so a super-node sits at the level of its earliest member.

Boundary preservation. Every candidate merges within one node role,

$$\mu(u) = \mu(v) \implies \tau(u) = \tau(v). \quad (3.25)$$

The PI/PO interface is the circuit’s contract with the rest of the design, and dissolving it changes what the graph represents rather than coarsening it. Imposing the restriction uniformly also makes a comparison between methods a comparison of *which* nodes are merged, not of whether the interface survived.

Schema compatibility. The count representation of (3.24) is a strict superset of the one-hot one: a super-node with a single member reproduces its original feature row exactly, and a single edge reproduces its own inversion indicator. An unreduced AIG is therefore already a valid input in the summarized schema, without conversion, and that is what makes RQ5 testable with one set of weights. It is a design constraint on every method rather than a convenience. A merge producing features outside this space would leave summary and full graphs in different input spaces and make cross-state inference undefined.

Domain-Specific Coarsening

Two AIG-native candidates, with the choice between them open because neither one’s compression on this corpus is known.

MFFC contraction sends every gate to the root of the maximal fanout-free cone containing it, $\mu(u) = v$ for the v with $u \in \text{MFFC}(v)$ of (2.8), iterated to a fixed point. What survives is the multi-fanout sharing skeleton together with the fixed interface. It takes no parameter, so its compression is a property of the circuit.

Cone coarsening merges along two axes. Chains of single-fanout gates contract into one super-node up to a capped chain

length, which reduces depth. Gates inside a bounded band of topological levels that share an immediate post-dominator (2.1.2) merge together, which compresses width while leaving critical-path depth, and with it the level encoding, exact. The chain cap and the band width are the only compression controls a domain-specific candidate offers here, and neither has been given a value.

The two are related rather than rival. Chain contraction is the capped, chain-shaped special case of MFFC absorption, which additionally captures cones whose interior reconverges. Both preserve acyclicity, since a cluster is absorbed only into the single cluster receiving all of its outgoing edges, which cannot close a cycle in a DAG. For cone coarsening the guarantee holds at band width zero, where the width axis merges only within a level; wider bands give it up along with the exactness of the level encoding.

Merging MFFCs invites an objection, and the answer is the thesis claim in miniature. Refactoring collapses and re-expresses one MFFC at a time, so merging them might be expected to destroy exactly the signal the model must read. The claim here is the opposite, and it is falsifiable: intra-cone wiring is what local rewriting is *free* to change, so it is noise with respect to the label, whereas the sharing structure that *constrains* what rewriting can achieve is what survives the merge.

Graded Colour Refinement

The exact anchor of the study, and the one candidate whose merge map is defined by what the encoder can see rather than by a property of the graph. Nodes merge when they agree after d rounds of the refinement (2.15), and two substitutions in that recurrence carry the whole adaptation to AIGs,

$$c_v^{(0)} = (\tau(v), \ell(v)), \quad \{ \{ c_u^{(t-1)} : u \in \mathcal{N}(v) \} \} \mapsto \{ \{ c_u^{(t-1)}, \text{inv}(u, v) \} : u \in \text{fi}(v) \} \} \quad (3.26)$$

Initialising from the role and the level (2.5) rather than uniformly keeps nodes at different depths in different classes, so a super-node's pooled level is exact. Carrying *inv* inside the neighbour's identity keeps a normal and an inverted connection out of the same class; without it, logically distinct subgraphs merge. The neighbourhood is the fanin (2.2), so refinement runs backward from the primary outputs. Direction is itself a parameter, and it changes both how much compresses and which structure survives.

The subscript in (3.26) is the count cap of 2.1.4, and only its two named endpoints are run: $c = \infty$, lossless for a summing GNN, and $c = 1$, which is bisimulation, coarser and lossy for a counting model. Intermediate caps are omitted because the endpoints bracket them and each value costs a training run. The depth is set to $d = L$, the encoder's message-passing depth, and that equality is the justification rather than a convenience. A smaller d merges less and gains nothing; a larger one merges nodes the encoder can separate, and the losslessness argument fails.

At $c = \infty$ the merge is orthogonal to the optimization being predicted by construction: it removes only what the model provably cannot see, so it cannot erase the label. That is what makes it the positive control for RQ5 rather than merely another contender, and realising the guarantee end to end requires the encoder variant of 3.1.5.

Structural hashing has already removed the one-hop equivalences (2.1.3), since every corpus graph is strashed. What this method finds is therefore *residual* redundancy beyond one hop. How much of it exists is not yet known, and the answer bounds what the exact track can compress.

Convolution Matching

The strongest general-purpose competitor. ConvMatch [Dickens et al., 2024] merges nodes that are equivalent with respect to the *graph convolution operation* itself rather than with respect to a graph property, and it accepts a target compression ratio directly. The only adaptation is (3.25), imposed on the candidate set so that the convolution objective is left unmodified. It is the bar the domain-specific candidate must clear, and the right bar because it is GNN-aware and domain-blind: a win against it is a win against a method optimising the same objective, not against a strawman.

Its behaviour on a deep DAG with inverted edges is untested in the literature, and one property of its objective explains why it may differ from the published setting. The merge cost is an unweighted sum over the convolution output, so merging two low-degree nodes scores better than merging two genuine convolution twins of higher degree, regardless of similarity. Convolution equivalence therefore decides *between* pairs of comparable degree and not across them, which on the degree profile of an AIG biases merging toward the input frontier.

Random Within-Type Merging

Uniformly random nodes of the same type are merged until the node count reaches a target. The target is not a free parameter: it is set per pairing to the partner method's achieved compression. Restricting merges by (3.25) is what makes the comparison isolate *which* nodes are merged rather than confounding it with whether the interface survived.

As a naive floor, it answers the question a reader asks first: a method that does not beat random merging at the same compression bought nothing with its sophistication. It is also the counterpart of random edge dropout (3.2.2), so both halves of the study have a floor built the same way.

As a matching instrument, it is what makes RQ4 answerable for this family at all. Compression is a free parameter for only some methods: convolution matching takes a target ratio, cone coarsening offers coarse integer knobs, and MFFC contraction and colour refinement have none whatsoever. There is no general way to make two *real* methods meet at one ratio. Each method is instead compared against a random control run at that method's own achieved compression, which replaces a matching problem that has no solution.

3.3 Data

[TODO: lets]

3.3.1 Source Circuits

The corpus is built from 55 source circuits, drawn from three collections and summarised in 3.3. The characteristics of each design, measured on its untransformed base AIG, are given in 3.4.

Twenty-nine are the hardware IP designs of OpenABC-D [Chowdhury et al., 2021], which collects them in turn from OpenCores, IWLS, the MIT Lincoln Laboratory CEP and OpenPiton. They span bus and peripheral control (spi, simple_spi, i2c, sasc, ss_pcm, usb_phy, wb_dma, wb_conmax, ac97_ctrl, pci, mem_ctrl, ethernet, vga_lcd), cryptography (aes, aes_xcrypt, aes_secworks, des3_area, sha256), signal processing (fir, iir, dft, idft, jpeg), and processors and SoC fragments (tv80, tinyRock, bp_be, picosoc, fpu, dynamic_node).

All 47 non-synthetic designs entered the pipeline as BENCH netlists from the OpenABC-D distribution, obtained from the NYU UltraViolet repository record (<https://ultraviolet.library.nyu.edu/records/mw6q2-a8p15>) linked in the project’s code repository (<https://github.com/NYU-MLDA/OpenABC>). Its bench directory carries the classical benchmarks of the following paragraph alongside the 29 IP designs the dataset paper documents. Each netlist is converted to an AIG by a single ABC pass that reads it and applies structural hashing, which is the only pre-processing applied before the transformation of 3.3.2.

Eighteen more are classical combinational benchmarks that OpenABC-D does not include: the eight arithmetic circuits of the EPFL suite [Amarú et al., 2015] (div, hyp, log2, max, multiplier, sin, sqrt, square), four ISCAS-85 circuits [Brglez and Fujiwara, 1985] (c1355, c5315, c6288, c7552), and six MCNC/LGSynth circuits [Yang, 1991] (apex1, bc0, dalu, i10, k2, mainpla). These extend the corpus into deep arithmetic, which OpenABC-D deliberately excludes in favour of larger and functionally more diverse IP. The effect on the depth distribution is pronounced. Median depth across the 47 designs outside this suite is 31 levels and the deepest of them is fpu at 819, whereas the EPFL circuits reach 4,373 (div), 5,059 (sqrt) and 24,802 (hyp). The three deepest circuits in the corpus therefore all enter through this collection, and they set the bound on the precomputed depth encoding of 3.3.5.

The remaining eight are synthetic circuits, named 128, 256, 512, 1024, 2048, 4096, 8192 and 16384 after their *primary input count* rather than their size: 16384 has 16,384 inputs but 131,072 AND gates, so the name understates the circuit by an order of magnitude. They are randomly sampled AIGs, produced by the random AIG generator of mockturtle, part of the EPFL logic synthesis libraries [Soeken et al., 2018]. The generator takes a number of primary inputs, a number of AND gates and a seed, and grows the network one gate at a time. It draws two operands uniformly at random from the signals built so far, namely the constant, the primary inputs and every gate already created, inverts each with probability one half, and adds the resulting AND gate to the pool. Draws that structurally hash to an existing node or collapse to a trivial case do not count towards the gate budget and are redrawn, so every gate is distinct. Once the budget is met, every node still without a fanout is made a primary output. The gate budget is eight times the input count in every one of the eight, from 128 inputs / 1,024 gates up to 16,384 inputs / 131,072 gates, doubling at each step. The number of primary outputs is not a parameter but a consequence of the sampling

and lands between 37% and 39% of the gate count (398 for the smallest, 49,307 for the largest): it is the share of nodes never drawn as an operand.

The generation script itself was not kept, so this account is reconstructed from the eight AIGs rather than from a recipe, and they determine it tightly. Across all eight there is not one duplicated or trivial AND gate, the gate count is exactly eight times the input count, and the primary outputs are exactly the fanout-free nodes together with the unused constant, which the generator emits as the first output, a detail of that implementation and of no other. What is *not* recoverable is the seed. The consequences of that are taken up in 3.4.4.

That construction is worth stating precisely, because it explains how these eight behave later. A network whose fanins are drawn uniformly at random from the whole pool, with duplicates rejected as they are created, arrives already structurally hashed and with almost none of the repeated local subfunctions that rewrite, refactor and resub exist to collapse. Their near-zero optimizability (3.3.4) is therefore a property of how they were sampled rather than a measurement failure, and it is the largest circuits, where the pool is widest and collisions rarest, that sit closest to an exact fixed point.

[TODO: if in table text could maybe be reduced? comment out the original text and rewrite more concise outline style]

Table 3.3: Where the 55 source designs come from, and under what terms. The license column gives each suite’s own terms; the ISCAS-85 and MCNC circuits predate modern license practice and carry no formal terms. The terms governing the copies actually used are stated in the text.

Source	Designs	License
OpenABC-D [Chowdhury et al., 2021]	29	BSD 3-Clause
EPFL combinational suite [Amarú et al., 2015]	8	MIT
ISCAS-85 [Brglez and Fujiwara, 1985]	4	none stated
MCNC / LGSynth [Yang, 1991]	6	none stated
Synthetic (this work)	8	n/a
	55	

The license column of 3.3 records each suite’s own terms, but that is not the term that governs the copies actually used, because the 18 classical benchmarks were not obtained from their original distributions. Every non-synthetic netlist entered the corpus through the OpenABC-D release, so those copies were received under its BSD 3-Clause license. The EPFL suite’s MIT license applies in addition, and the ISCAS-85 and MCNC circuits attach no formal terms of their own, leaving the BSD 3-Clause license of the redistributing dataset as the only written term on the copies used. Any release of the corpus is therefore governed by BSD 3-Clause attribution requirements, retaining the MIT notice for the EPFL circuits (5.6).

The provenance is the point here, not the inventory, because the design-level split (3.4.2) rests on the assumption that these 55 are genuinely distinct circuits rather than variants of one another. Mostly they are: three different sources, four benchmark traditions, and no design derived from another. The assumption is not perfect, and the exceptions should be named rather than assumed away. The three AES variants implement the same ci-

pher; `dft` and `idft` are forward and inverse transforms of one structure, as are `fir` and `iir` to a lesser degree; and the eight synthetic circuits differ from each other only in scale. A split by design does not guarantee that such a family lands entirely on one side of it, so a small amount of family-level similarity can cross the train/test boundary even though no individual circuit does. This is a weaker form of leakage than the one design-level splitting exists to close, being a structural family resemblance rather than the same circuit in two optimization states, but it bounds how strongly an unseen-design result can be read and is taken up in 5.4.

Scale spans nearly three orders of magnitude, which is what forces the node-budget batching of 3.4.2. The base AIGs range from 403 AND nodes (`ss_pcm`) to 245,046 (`dft`) at a median of 11,464, and depth ranges from 8 levels to 24,802 at a median of 35. Transformation and optimization move these figures, so the bounds the preprocessing pipeline records over the full corpus are higher still, at 366,040 AND nodes and a depth of 24,972.

Stats

[TODO: mention all the depths, nodes distributions and etc. for the chosen starting circuits. rename this to something more formal]

3.3.2 Dataset Generation

Random Structural Transformation

Each source design is expanded by applying random synthesis recipes to it and retaining every intermediate result, following the construction of OpenABC-D [Chowdhury et al., 2021]. A recipe is a sequence of 20 transformations drawn independently and uniformly from seven area-oriented passes: balancing, and rewriting, refactoring and resubstitution each in a plain and a zero-cost variant, the latter accepting replacements that leave node count unchanged. The same 200 recipes, the first 200 of the 1,500 released with OpenABC-D and all of them distinct, are applied to every design, so recipe identity is not confounded with design identity.

The AIG is written after *every* pass rather than only at the end, so a single recipe contributes 20 graphs, one per prefix of the recipe, plus one further graph after the structural-choice pass that follows the transformation sequence: 21 in total. Together with the unmodified design itself this gives $1 + 200 \times 21 = 4,201$ base graphs per design, and $55 \times 4,201 = 231,055$ overall. The recipes continue into a technology-mapping tail, which does not execute because no standard-cell library is loaded in this flow, so no graph beyond the structural-choice pass is produced. The per-design count is verified rather than assumed, and generation is treated as failed unless exactly 4,201 AIGs are present. The pass alphabet, the recipe format and the unexecuted tail are listed in A.1.

Randomised transformation rather than plain replication is what makes this corpus a learning problem at all. Replicating a design would yield identical graphs carrying identical labels, adding weight but no information. A random recipe prefix instead yields a circuit that is *functionally* equivalent to its source,

since every pass is equivalence-preserving, but structurally different, and, critically, differently optimizable: how much a synthesis script can still remove depends on which rewriting the graph has already absorbed. The corpus therefore varies along exactly the axis the label measures while holding the function fixed.

That is also what makes the design-level split (3.4.2) necessary rather than merely conservative. The 4,201 graphs derived from one design are not independent samples: they share a Boolean function, a primary-input interface, and a large amount of common structure, and successive prefixes of one recipe differ by a single pass. A random split over graphs would place near-siblings of every test graph into training, and the resulting score would measure recall of a design the model had already seen rather than generalization. Splitting on the design partitions the corpus along the one boundary this construction never crosses.

[TODO: TODO: step-21 choice networks are not plain AIGs; the node count understates them in 43 of 200 cases. this was a mistake i dont want to retrain can this be added in limitation instead do i mention this here? i will just pull it from inference values at least]

Optimization

Four synthesis scripts were applied during dataset generation: `Orchestrate`, `Deepsyn`, `Syn4` and `C2RS`. They differ in how they search: some apply a fixed sequence of rewriting, refactoring and balancing passes, while others explore seeded or scored variations. Command templates and per-algorithm parameters are given in A.

Only `Orchestrate` is used for training and evaluation in this thesis (1.2.6); the remaining three exist on disk and are the most direct extension available (6.3.1).

3.3.3 Tiered Dataset Structure

The corpus has two tiers of *stored* graphs. Every stored graph is a training example; the optimized results used to label them are not themselves retained, for the reason given in 3.3.3.

Tier 0: Base Graphs

This originates from 55 designs randomly transformed into 231,055 base graphs (Tier-0).

Tier 1: Single-Algorithm Optimization

Applying the target synthesis algorithms to these base graphs yields exactly 231,055 primary (Tier-1) samples per algorithm. For the `Orchestrate` prediction task three of them contribute samples (`Deepsyn`, `Syn4` and `C2RS`), while `Orchestrate` itself does not: an `Orchestrate`-produced graph relabelled by `Orchestrate` is precisely the sequential step that 3.3.3 stops at, so it is a label source rather than a sample.

Table 3.4: Characteristics of the 55 source circuits before any transformation. Primary inputs (PI), primary outputs (PO), AND nodes (N), edges (E), inverted edges (I) and depth (D), measured on the base AIG of each design. Every AND node has two fanins and every output one, so $E = 2N + PO$ throughout. Rows are grouped by function and ordered by size within each group, and colour repeats that grouping: Communication and bus protocol, Controller, Cryptography, Signal processing, Processor and SoC, Arithmetic, Combinational logic, Synthetic random.

Design	Source	PI	PO	N	E	I	D	Function
ss_pcm	OpenABC-D	104	90	403	896	462	8	Single slot PCM
usb_phy	OpenABC-D	132	90	487	1,064	513	10	USB PHY 1.1
sasc	OpenABC-D	135	125	613	1,351	788	9	Simple asynchronous serial controller
simple_spi	OpenABC-D	164	132	930	1,992	1,084	12	MC68HC11E based SPI interface
i2c	OpenABC-D	177	128	1,169	2,466	1,188	15	Bidirectional serial bus protocol
spi	OpenABC-D	254	238	4,219	8,676	5,524	35	Serial peripheral interface
wb_dma	OpenABC-D	828	702	4,587	9,876	4,768	29	Wishbone DMA/bridge
pci	OpenABC-D	3,429	3,157	19,547	42,251	25,719	29	PCI controller
wb_conmax	OpenABC-D	2,122	2,075	47,840	97,755	42,138	24	Wishbone Conmax
ethernet	OpenABC-D	10,731	10,422	67,164	144,750	86,799	34	Ethernet IP core
ac97_ctrl	OpenABC-D	2,339	2,137	11,464	25,065	14,326	11	Wishbone AC97 controller
mem_ctrl	OpenABC-D	1,187	962	18,092	37,146	16,307	36	Wishbone memory controller
vga_lcd	OpenABC-D	17,322	17,063	105,334	227,731	141,037	23	Wishbone VGA/LCD controller
des3_area	OpenABC-D	303	64	4,971	10,006	4,686	30	DES3 encrypt/decrypt
sha256	OpenABC-D	1,943	1,042	15,816	32,674	18,459	76	SHA-256 hash
aes	OpenABC-D	683	529	28,925	58,379	20,494	27	AES (LUT-based)
aes_secworks	OpenABC-D	3,087	2,604	40,778	84,160	45,391	42	AES-128 (simple)
aes_xcrypt	OpenABC-D	1,975	1,805	45,840	93,485	36,180	43	AES-128/192/256
fir	OpenABC-D	410	351	4,558	9,467	5,696	47	FIR filter
iir	OpenABC-D	494	441	6,978	14,397	8,596	73	IIR filter
jpeg	OpenABC-D	4,962	4,789	114,771	234,331	146,080	40	JPEG encoder
idft	OpenABC-D	37,603	37,419	241,552	520,523	317,210	43	Inverse discrete Fourier transform
dft	OpenABC-D	37,597	37,417	245,046	527,509	322,206	43	Discrete Fourier transform
tv80	OpenABC-D	636	361	11,328	23,017	11,653	54	TV80 8-bit microprocessor
dynamic_node	OpenABC-D	2,708	2,575	18,094	38,763	23,377	33	OpenPiton NoC architecture
fpu	OpenABC-D	632	409	29,623	59,655	37,142	819	OpenSPARC T1 floating point unit
tinyRocket	OpenABC-D	4,561	4,181	52,315	108,811	67,410	80	32-bit tiny RISC-V core
bp_be	OpenABC-D	11,592	8,413	82,514	173,441	109,608	86	BlackParrot RISC-V execution engine
picosoc	OpenABC-D	11,302	10,797	82,945	176,687	106,737	43	SoC with PicoRV32 RISC-V
c6288	ISCAS-85	32	32	2,337	4,706	4,132	121	16×16 multiplier
max	EPFL	512	130	2,865	5,860	3,629	288	Maximum
sin	EPFL	24	25	5,416	10,857	6,085	226	Sine
square	EPFL	64	128	18,485	37,098	23,439	251	Square
sqrt	EPFL	128	64	24,618	49,300	36,581	5,059	Square root
multiplier	EPFL	128	128	27,062	54,252	32,086	275	Multiplier
log2	EPFL	32	32	32,060	64,152	36,784	445	Base-2 logarithm
div	EPFL	128	128	57,247	114,622	87,301	4,373	Divider
hyp	EPFL	256	128	214,335	428,798	277,636	24,802	Hypotenuse
c1355	ISCAS-85	41	32	512	1,056	658	30	32-bit error-correcting circuit
apex1	MCNC	45	45	1,578	3,201	1,082	15	PLA-derived logic
bc0	MCNC	26	11	1,592	3,195	1,958	32	PLA-derived logic
c5315	ISCAS-85	178	123	1,613	3,349	1,743	39	ALU and selector
dalu	MCNC	75	16	1,735	3,486	1,963	36	Datapath ALU
c7552	ISCAS-85	207	107	2,198	4,503	2,692	31	ALU and control
i10	MCNC	257	224	2,273	4,770	2,323	59	Combinational logic
k2	MCNC	45	45	2,290	4,625	1,784	23	PLA-derived logic
mainpla	MCNC	27	54	5,346	10,746	7,172	39	PLA-derived logic
128	Synthetic	128	398	1,024	2,446	1,024	13	Random AIG, 128 inputs
256	Synthetic	256	773	2,048	4,869	2,038	14	Random AIG, 256 inputs
512	Synthetic	512	1,572	4,096	9,764	4,122	17	Random AIG, 512 inputs
1024	Synthetic	1,024	3,085	8,192	19,469	8,216	16	Random AIG, 1024 inputs
2048	Synthetic	2,048	6,195	16,384	38,963	16,484	15	Random AIG, 2048 inputs
4096	Synthetic	4,096	12,231	32,768	77,767	32,906	16	Random AIG, 4096 inputs
8192	Synthetic	8,192	24,624	65,536	155,696	65,302	17	Random AIG, 8192 inputs
16384	Synthetic	16,384	49,307	131,072	311,451	130,909	18	Random AIG, 16384 inputs

Labelling and Tier Depth

Labelling a graph means running the target script on it and recording how many nodes it removed (3.3.4). The optimized *output* of that run is a graph too, so the construction could be iterated indefinitely; it is stopped after one step, and the outputs of the labelling runs on tier-1 graphs are discarded rather than promoted to a third tier.

Two consequences are worth stating, because both shape the results.

First, the two tiers are not independent of each other. Optimizing a tier-0 graph with algorithm S produces both that graph’s label under S and the corresponding tier-1 graph: one synthesis run serves both purposes. The tier-1 set is therefore exactly the set of labelling outputs of tier 0.

Second, this is what bounds how optimizable the corpus is. A tier-1 graph has already been through one full synthesis script, so what Orchestrate can still remove from it is small, though how small depends far more sharply on *which* script ran first than on the tier itself (3.6), to the point that the tier ordering inverts and tier 1 carries the higher mean target. That decomposition is taken up in 3.3.6. Roughly half the evaluation corpus has an optimizability of exactly zero (3.4). That is a property of the construction rather than a defect, but every R^2 in 4 has to be read against it (3.3.4).

3.3.4 Target Label: Optimizability

Definition

For an AIG G and synthesis algorithm S , the optimizability target is the relative node reduction $y(G)$ of (2.24), formulated as a regression task, and is not restated here. The bound motivates the terminal sigmoid in the prediction head (3.1): the model cannot predict outside the achievable range.

Note what this does and does not measure. It is a *relative* quantity, so it is comparable across circuits of different sizes, which is what makes a single model over a heterogeneous corpus meaningful, but it is also therefore harder to predict than an absolute count, since the model must infer both what will be removed and what the graph started with. It is an area proxy and says nothing about delay or power (1.2.6).

Label Distribution

The target is heavily concentrated near zero (mean 0.020, standard deviation 0.053, and just under half of all graphs at exactly zero), and what determines it is the synthesis script a graph came from rather than the circuit or the tier. Both facts constrain how the results can be read, and both are set out with the measurements in 3.3.6 and 3.3.6 rather than repeated here.

3.3.5 Graph Representation

Node and Edge Features

Each AIG is represented as an attributed graph carrying one-hot node types (constant, primary input, AND, primary output), a one-hot edge inversion flag (regular connection or inverted signal), and a topological depth encoding. The AIGER files are read and traversed with AIGVERSE [Walter, 2025], an interface to the same EPFL logic synthesis libraries [Soeken et al., 2018].

Positional Encoding Precomputation

Topological levels are computed once, offline, and stored on the graph; the model reads the stored attribute rather than recomputing it. This matters for more than speed. Levels are computed on the *unreduced* graph and then carried through reduction, so a partitioned or sparsified graph retains the depth information of the circuit it came from even where the edges that establish that depth have been cut. A runtime encoding would instead describe the reduced graph’s own topology, which is a different and, for this task, less informative quantity.

3.3.6 Dataset Analysis

This subsection characterises the corpus the model is actually trained and evaluated on: its scale, its structure, the shape of the target, and how the two evaluation splits differ from one another. Several decisions made elsewhere in this chapter rest on the distributions reported here rather than on argument: the graph-scale bound of 3.4.2, the node-budget batching of 3.4.2, the terminal sigmoid of 3.1. Every R^2 in 4 has to be read against them.

One scope caveat applies throughout, and it is not a small one. Except for 3.8, which covers every stored graph, what follows is measured on the *evaluation splits*, 96,430 test graphs over six designs and 80,525 validation graphs over five, rather than over the full corpus of 3.3.3. The splits are design-disjoint samples of the same construction, so the structural statistics can be read as representative; the label statistics cannot, because which designs land in a split changes them substantially (3.3.6).

Scale and Structure

Graph size spans nearly three orders of magnitude within the evaluation splits alone: a median of 15,977 nodes against a maximum of 232,715 (3.1, 3.5). This is the distribution that forces node-budget batching rather than a fixed graph count per batch (3.4.2); at a fixed batch size the largest graphs in the corpus are four orders of magnitude more expensive than the smallest.

Structurally the graphs are far more uniform than they are in size (3.2, 3.7). AND gates and the constant account for 82.1% of nodes on average and the primary-input/output interface for the remaining 17.9%, and the graphs are sparse and near-tree-like at 1.75 edges per node. The AND-gate share is not a descriptive statistic here but a mechanism: the `and_gate_only` sparsification of 3.2.2 drops exactly the interface nodes, so 82.1% is its node retention, and the spread on that share (63.3% to 99.9%

across graphs) is why its compression varies between circuits rather than being a dial that can be set.

Measured Target Distribution

The target is extremely concentrated near zero (3.3, 3.5). Pooled over the test split the mean optimizability is 0.020 with a standard deviation of 0.053, and 73.4% of graphs sit below $y = 0.005$. Just under half, 48.8%, are *exactly* zero: Orchestrate removes nothing at all.

Two consequences follow, and both shape how the results chapter can be read. First, R^2 is a ratio against the target’s own variance, and that variance is small, so a respectable RMSE is entirely compatible with a near-zero or negative R^2 ; the baseline diagnosis of 3.1.6 turns on precisely this. Second, with a point mass of this size at zero, a constant predictor is a strong baseline, and any model has to beat it on the minority of graphs where something is actually removable.

Where that point mass sits is not uniform (3.4). It concentrates in the smallest circuits, which are already minimal, and in the synthetic designs: the design 16384 contributes 16,080 test graphs, a sixth of the split, with a mean target of 0.0000, a maximum of 0.0003 and 92.4% of graphs at exactly zero. Orchestrate is at a fixed point on it, for the reasons given in 3.3.1. The joint distribution of size and label (3.5) is the entire signal available to the size-only baseline of 4.1.1, and it shows why that baseline is not trivial: graph size is largely a proxy for “is this optimizable at all”.

Determinants of the Target

The strongest single determinant of the target is not the circuit but the synthesis script that produced the graph (3.6). A tier-1 graph has already been optimized once, and how much Orchestrate can still remove depends overwhelmingly on which script ran first. Pooled over both evaluation splits, Syn4-derived graphs have a mean optimizability of 0.046 with only 11.7% at exactly zero, while Deepsyn- and C2RS-derived graphs average 0.0022 and 0.0011 with 54.4% and 52.1% at zero. That is a twenty- to forty-fold difference in the target, driven by a property of the label pipeline rather than of the circuit.

This is also what explains an apparent anomaly in 3.5: tier-1 graphs have a *higher* mean target (0.023) than tier-0 graphs (0.010), although 3.3.3 argues that a graph which has already been through a full synthesis script should have less left to remove. The elevation is entirely attributable to one of tier 1’s three sources. On the test split the three source means are 0.0015 (C2RS), 0.0025 (Deepsyn) and 0.0654 (Syn4), whose average is 0.023, exactly the tier-1 figure. Excluding the Syn4 third, tier 1 averages 0.0020, five times *below* tier 0. The tier is therefore the weaker cut of the data, and the source script should be reported alongside it wherever tier-wise results appear.

Composition of the Evaluation Splits

Under a design-level split the test set is a handful of whole circuits rather than a random sample (3.7), so every pooled test metric is an average over six designs and their particular mix.

With one of those six contributing a sixth of the graphs at a constant-zero target, the pooled figures are sensitive to composition in a way that a random split would have hidden.

Validation and test are correspondingly not exchangeable (3.8, 3.6). They are disjoint circuit sets, and they differ on every axis that matters: 26.2% of validation graphs sit at $y = 0$ against 48.8% of test, the test maximum target is more than twice the validation maximum (0.274 against 0.114), and the validation graphs are larger at the median (29,399 nodes against 15,977). The divergence is sharpest for the Syn4 subset, where not one validation graph has a target of exactly zero against 21.4% of the corresponding test graphs. This is expected under a design-level split rather than a defect, but it is the direct explanation for the validation-to-test gap in 4.37, and it bounds how far validation-based model selection can be assumed to transfer.

Graph scale across the test corpus




Figure 3.1: Node and edge count distributions across the evaluation corpus, by tier. This is the distribution behind the graph-scale bound of 3.4.2 and the node-budget batch sizes of 3.4.2.

Structural statistics of the corpus




Figure 3.2: Structural statistics recoverable without a pass over the graph cache. The AND-gate share is read off the AND-gate-only node masks, which drop exactly the primary inputs and outputs: 82.1% of nodes are AND gates or constants and 17.9% are interface. That figure *explains* that method's compression rather than merely accompanying it (3.2.2). The right panel checks that the evaluation splits are representative of the wider corpus sample on size.

Optimizability targets on the held-out test corpus




Figure 3.3: Distribution of the optimizability target on the held-out test corpus, pooled (log counts) and per tier. The label is concentrated near zero: the great majority of graphs sit below $y = 0.005$. This is the variance every R^2 in 4 is measured against, and it is why a respectable RMSE is compatible with a near-zero or negative R^2 .

Zero inflation: 48.8% of the test split has an optimizability of exactly zero

By design (red: Orchestrate is at a fixed point) circuit scale: small circuits are already minimal




Figure 3.4: Where the point mass at exactly zero sits. Roughly half the test split has an optimizability of exactly zero, concentrated in particular designs (16384 and 8192 reach a maximum y of 0.0003) and in the smallest circuits, which are already minimal. This is the figure the size-only baseline of 4.1.1 has to be read against: graph size is largely a proxy for “is this optimizable at all”.




Figure 3.5: Joint density of graph size and optimizability. This is the entire signal available to the size-only baseline of 4.1.1.




Figure 3.6: What the target mostly measures. A tier-1 graph is a base circuit already optimized by one of three scripts, and the label is what Orchestrate can still remove after that. Orchestrate finds almost nothing left on a C2RS- or Deepsyn-derived graph (52–54% score exactly zero) and a great deal on a Syn4 one (11.7% zero, mean y 0.046). The source script, not the circuit, is the strongest single determinant of the target.




Figure 3.7: How the design-disjoint test set is composed. Under a design-level split the test set is a handful of whole circuits rather than a random sample, so every pooled test metric is an average over these designs and their mix.

Validation and test are disjoint circuit sets, not exchangeable samples




Figure 3.8: Validation and test are disjoint circuit sets, not exchangeable samples: 26% of validation graphs sit at $y = 0$ against 49% of test, and the test maximum y is more than twice the validation maximum. This is expected under a design-level split (3.4.2) and is the direct explanation for the validation-to-test gap in 4.37; it also bounds how much validation-based model selection can be assumed to transfer.

Table 3.5: Held-out test corpus. The label is concentrated near zero, which is the denominator every R^2 in this chapter is measured against.

Tier	Graphs	Designs	Median nodes	Max nodes	Mean y	SD y	$y < 0.005$
tier0	24067	6	3559	232715	0.010	0.025	0.717
tier1	72363	6	15977	196764	0.023	0.059	0.739
pooled	96430	6	15977	232715	0.020	0.053	0.734

Table 3.6: Composition of the evaluation corpus over the two stored tiers. Validation and test differ substantially in label and size distribution, which is expected under a design-level split and is the direct explanation for the validation-to-test gap.

Split	Designs	Graphs	tier0	tier1	Median nodes	Mean y	$y = 0$	Max y
val	5	80525	20108	60417	29399	0.008	0.262	0.114
test	6	96430	24067	72363	15977	0.020	0.488	0.274

Table 3.7: Structural corpus statistics. The AND-gate share is recovered from the AND-gate-only node masks over 10,000 graphs: that method drops exactly the primary inputs and outputs, so its node retention is the AND-and-constant share. This is what fixes that method’s compression, rather than a parameter that could be dialled.

Statistic	Mean	SD	Min	Max
AND gates and constants (share of nodes)	0.821	0.114	0.633	0.999
Primary inputs and outputs (share of nodes)	0.179	0.114	0.001	0.367
Edges per node	1.752	0.166	1.577	1.992

Table 3.8: Whole-corpus statistics by tier.

NOT YET GENERATED. Run `data/creation/corpus_tier_stats.py -latex` on the cluster, where the per-graph ML CSV lives, and write the result to `media/tables/corpus_tiers.tex`.

3.3.7 Held-Out Hyperparameter Tuning Subset

A smaller subset of 50,000 graphs, balanced between the two stored tiers, was used exclusively for the initial hyperparameter tuning and then excluded from the rest.

[TODO: make much more concise. keep full text commented out.]

3.4 Experiments

3.4.1 Research Design

The manipulated variable is the reduction method applied to the training data; the encoder, dataset, splits, label, optimizer and evaluation protocol are held constant, so that differences in outcome are attributable to the reduction. The reference point is the full-graph model of RQ1, and every reduced configuration is reported relative to it rather than in isolation.

Each question makes its own demand on that design. RQ1 requires the unreduced reference point to be trainable at all, which is what bounds the corpus scale (3.4.2). RQ2 requires every configuration to be measured under one batching budget, since memory and throughput are not comparable across budgets (3.4.2). RQ3 requires the reduction to be the only thing that changed, which is why hyperparameters are fixed at the full-graph values (3.4.2). RQ4 requires pairs matched on *achieved* compression rather than on configured parameters, so that a retention gap is not a compression gap in disguise. RQ5 requires no new training, only a second evaluation pass over the Phase 3 checkpoints. The two hypotheses under test, H1 on path contraction and H2 on the strength of the heuristics, are stated with the questions they belong to (1.2.3, 1.2.4).

Positive control. One configuration has a predetermined outcome. Colour refinement at count-cap ∞ on the exact-compression track merges only nodes the encoder provably cannot distinguish (3.1.5), so a model trained on those coarsened graphs must score identically to the full-graph baseline, on reduced and on unreduced graphs alike. Any deviation is a defect to be found rather than a finding to be reported, and without it RQ5 cannot separate “models trained on reduced graphs do not generalize” from “the evaluation path is wrong”.

The experiments. Four experiments implement the design, and they do not map one-to-one onto the five research questions (3.9). RQ4 is a re-analysis of the Phase 3 runs under a paired comparison rather than a separate experiment, which is why it costs no additional training; RQ5 reuses the Phase 3 checkpoints and adds only an evaluation pass. The ordering is a dependency chain rather than a preference: Phase 3 cannot start until Phase 2 has produced the cached reduction artifacts, and Phase 4 needs Phase 3’s checkpoints.

Phase 1 establishes both the predictive reference point and the baseline hardware cost (peak memory, epoch time) that every reduced configuration is measured against. Phase 2 requires no training and determines which methods are viable at corpus scale at all. Phase 3 places the reduced configurations against

Table 3.9: The four experimental phases, what each answers, and what each costs. Only Phases 1 and 3 consume training compute.

Phase	Answers	What is run	Training cost
1	RQ1	Full-graph model and baselines	one run per
2	RQ2	Offline reduction precompute, profiled	none
3	RQ3, RQ4	One model per reduced configuration	one run per
4	RQ5	Phase 3 checkpoints on full graphs	none (evalu

the Phase 1 reference on a Pareto front of predictive degradation versus memory and speed. Phase 4 queries models trained exclusively on reduced graphs with unreduced ones, testing structural generalization.

Encoder family, target synthesis algorithm and graph scale are held fixed, so the design ranks reduction methods under those conditions and establishes nothing about transfer beyond them (5.4).

3.4.2 Setup

Data Splitting

An 80/10/10 split by volume, partitioned *by design*: all graphs derived from a given source circuit fall in the same split. This is the strong choice and it is deliberate. Because the corpus is built by transforming and then repeatedly optimizing a small number of source designs, a random split would place a base graph in training and its own optimized descendant in test: structurally near-identical circuits in different optimization states. The model could then score well by recognising the design rather than by reading its structure. Splitting by design removes that path entirely, at the cost of making the task harder: the test set contains circuits the model has never seen in any state.

Two alternative strategies are implemented and available (split by optimization recipe, and a plain random split). They are not used for the reported results; they are run once, at the primary configuration, as the protocol-sensitivity experiment RQ1a (4.1.6). The three protocols differ in what they hold out, and each corresponds to an evaluation setting already established in the literature (Table 3.10).

Graph Scale Bounds

The corpus is bounded at 366,040 nodes per AIG. This intermediate scale is what makes RQ1 possible: a single unreduced graph fits in accelerator memory, so the full-graph baseline every other result is measured against can actually be trained. Larger circuits would remove the reference point and leave only reduced configurations to compare with each other.

The bound is a scoping decision with a cost, and it is the limitation closest to the motivation: the argument for this work invokes circuits of millions to billions of nodes, and the evaluation does not reach them (5.4.2).

Table 3.10: The three train/test protocols implemented in this work, ordered from leakiest to strictest. Only the design-disjoint protocol is used for the reported results; the other two are run once each at the same configuration to quantify how much of a reported score is design recognition rather than structural reading.

Protocol	Held out	Established as	Role here
Random (per-row)	Nothing. Individual graphs are assigned independently, so a base graph and its own optimized descendant may straddle the split.	The common default in circuit representation-learning evaluations.	Ablation
Recipe-disjoint	Whole synthesis recipes: every tier-0 step of a held-out recipe, plus its tier-1 and tier-2 descendants, across all designs. Circuits stay seen; the recipe does not.	OpenABC-D Variant 1 [Chowdhury et al., 2021]; LOSTIN transductive [Wu et al., 2022].	RQ1a
Design-disjoint	Whole source circuits: every graph derived from a held-out design, in any optimization state, is confined to one split.	OpenABC-D Variant 2 [Chowdhury et al., 2021]; LOSTIN inductive [Wu et al., 2022].	Headline

Dynamic Batching

Graphs in this corpus vary in size by orders of magnitude, so a fixed batch size in graphs gives wildly varying memory use and eventually fails on the largest circuits. Batches are instead assembled to a maximum *node* budget by greedy bin-packing, which bounds activation memory regardless of which graphs are drawn. The batch plan is computed once and cached, since replanning over a corpus this size is expensive.

Two budgets are used and they are deliberately distinct. Training uses the smaller of the two; evaluation, which stores no activations or gradients, uses a substantially larger one. Keeping them separate rather than raising a single shared value avoids invalidating the batch plan the trained checkpoints were produced under.

Constraint on efficiency measurements. The evaluation budget must be identical across all configurations, because peak memory and throughput both depend on it. Any change requires re-running every configuration; a mixed set of budgets makes the RQ2 comparison meaningless rather than merely noisy. This is stated here because it constrains how the results in 4.2 may be read and combined.

One irreducible peak remains that no budget can lower: a graph larger than the budget still forms a batch of its own, since a graph-level target cannot be split across batches. The largest circuit in the corpus therefore sets a floor on peak memory for every configuration that does not reduce node counts.

Model Configuration

An Edge-Aware Graph Convolutional Network (GCN+) is trained to predict the optimizability of a circuit when subjected to the Orchestrate synthesis script. The architecture is described in 3.1 and is held fixed across all reduction configurations; the baselines it is compared against are described in 3.1.6.

Hyperparameter Tuning

Hyperparameters were selected by an Optuna [Akiba et al., 2019] sweep on the held-out balanced subset of 3.3.7, which is then excluded from all subsequent training and evaluation. Tuning on a subset that later appeared in training or test would let the reported scores absorb information from the tuning process; excluding it entirely costs 50,000 graphs and removes the question.

Tuning was performed once, on the full-graph configuration, and the resulting values are held fixed for every reduced configuration. This is the conservative choice for the comparison being made: re-tuning per reduction method would confound “this reduction retains more signal” with “this reduction received more tuning effort”. It does mean reduced configurations are evaluated at hyperparameters chosen for a different input distribution, which if anything biases against them, and bounds the strength of any negative result about reduction.

The search covered architecture and optimization jointly, in 13 to 18 dimensions depending on which encoder a trial drew. The architectural dimensions were the encoder itself, drawn from six candidates (a graph convolutional network, two edge-conditioned isomorphism networks, a plain message-passing network, an attentional convolution and a hybrid local/global model), hidden width in $\{32, 64, 128, 256, 512\}$, message-passing depth $L \in \{2, \dots, 8\}$, jumping-knowledge mode in $\{\text{last, cat, max, sum}\}$, normalisation in $\{\text{batch, layer, graph}\}$, graph pooling in $\{\text{mean, max, sum}\}$, positional encoding in $\{\text{none, level, input paths, local shortest-path sum}\}$ with width in $\{16, 32, 64, 128\}$, and, on the encoders that have them, attention heads in $\{1, 2, 4\}$, edge-channel width in $\{32, 64, 128\}$ and update-layer depth in $\{2, 3, 4\}$. The optimization dimensions were learning rate, weight decay, dropout, the loss transition point and the batch size in graphs, the first two searched over about two decades each (10^{-4} to 6×10^{-3} and 3×10^{-5} to 10^{-2}).

The sweep ran in two stages across three parallel workers with distinct sampler seeds and a shared data split. An exploration stage requested 50 trials per worker on a 15,000-graph subsample with a one-hour cap per trial; an exploitation stage requested 20 per worker on 35,000 graphs with a four-hour cap, warm-started from the ten best exploration configurations. Of the 210 trials requested, the surviving worker logs record 94 started, 52 completed and 38 pruned, 34 of those on memory exhaustion. The tuning stage therefore hits the same bottleneck the thesis is about, and it prunes hardest on precisely the configurations that would have been most expensive to train.

Training Configuration

Smooth L1 loss with transition point $\beta = 0.01$ (3.29), set well below the default because the target occupies a narrow band

within $[0, 1]$ and the default would place nearly every residual in the quadratic branch. Optimization uses AdamW at learning rate 3×10^{-4} and weight decay 3.96×10^{-3} , with linear warmup over 10,000 steps and plateau-triggered decay thereafter, stopping early on validation loss. The remaining schedule, clipping and regularisation values are listed in A.4.

Two settings change the measurements rather than the model and are disclosed with them: the training step is compiled, which moves every timing figure, and mixed precision is used on the accelerator, which moves every memory figure. Neither varies across configurations, so the comparisons hold, but no absolute number here means anything without them.

Execution Environment and Workload Profile

All experiments ran on the Snellius national supercomputer: H100 80 GB GPU nodes for training and evaluation, and CPU-only nodes for the offline reduction precompute, which is embarrassingly parallel over graphs and does not benefit from a GPU. Software versions and the total wall-clock budget are recorded in 3.4.4.

Two measured characteristics of the workload belong here rather than there, because they explain the shape of the RQ2 results. Memory scales close to linearly with nodes and edges per batch, the encoder’s own parameters being negligible at this width and depth, which is why a node budget bounds memory effectively at all. Message passing is latency-bound on gather and scatter operations rather than compute-bound: at full device occupancy the floating-point pipelines and memory bandwidth both sit far below saturation. A larger batch therefore amortises per-batch overhead without delivering a proportional speedup, and reductions that remove *edges* should help throughput more than their node compression alone would suggest.

3.4.3 Evaluation Metrics

Four groups of metrics are reported. Accuracy and ranking are reported together and never separately: on a coarsened EDA graph, CTS-Bench [Khadka et al., 2026] observed the mean absolute error barely moving while R^2 fell below zero. Absolute error alone can look healthy after the model has lost all discriminative power, which is exactly the failure a reduction study is most at risk of missing.

Notation. Let $\mathcal{D} = \{G_i\}_{i=1}^N$ be an evaluation set, $y_i = y(G_i) \in [0, 1]$ the optimizability target of 2.1.6, f_θ the trained encoder and R the reduction operator, with $R = \text{id}$ for the unreduced baseline. Residuals are written $r_i = \hat{y}_i - y_i$, and $\bar{y} = N^{-1} \sum_i y_i$ is the target mean over $\mathcal{D}$.

Evaluation protocol. Each trained configuration is evaluated in two passes over the test set, differing only in what the same weights are given as input:

$$\hat{y}_i^{\text{match}} = f_\theta(R(G_i)), \quad (3.27)$$

$$\hat{y}_i^{\text{full}} = f_\theta(G_i). \quad (3.28)$$

The matched pass (3.27) sees graphs under the reduction the model was trained with, and the full-graph pass (3.28) sees unreduced graphs. The difference between the passes isolates the cost of the structural shift from the cost of the reduction itself.

Which measurement answers which question. Accuracy and ranking on the matched pass answer RQ1 for the unreduced configuration and RQ3 for the reduced ones, and the same figures recomputed under each split protocol answer RQ1a. Retention (2.23), offline cost and the amortisation threshold (3.40), together with peak memory, per-graph step time and throughput (3.38), answer RQ2. RQ4 is the paired gap (3.35) between a domain-informed configuration and its matched generic control, with the receptive field (3.42) as the mechanism test H1 requires. RQ5 is the cross-state drop-off (3.34) between the two passes.

Predictive Accuracy

The training objective is the smooth L1 loss with transition point $\beta = 0.01$,

$$\ell_\beta(r) = \begin{cases} r^2/(2\beta), & |r| < \beta, \\ |r| - \beta/2, & \text{otherwise,} \end{cases} \quad (3.29)$$

reported on validation and test for comparability with the training curves. Since the target occupies $[0, 1]$ while $\beta = 0.01$, all but the smallest residuals fall in the linear branch of (3.29): outside a band of one percentage point the objective behaves as a mean absolute error. Accuracy proper is reported as

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |r_i|, \quad (3.30)$$

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N r_i^2}, \quad (3.31)$$

$$R^2 = 1 - \frac{\sum_{i=1}^N r_i^2}{\sum_{i=1}^N (y_i - \bar{y})^2}. \quad (3.32)$$

MAE and RMSE are on the label’s own scale, so $\text{RMSE} = 0.05$ is five percentage points of optimizability. By the power mean inequality $\text{MAE} \leq \text{RMSE}$, with equality exactly when every $|r_i|$ is identical, so the ratio RMSE/MAE measures how much a few large errors carry the average. Equation (3.32) divides the same squared error by the target’s own spread: $R^2 = 0$ is the score of the constant predictor $\hat{y}_i = \bar{y}$, and $R^2 < 0$ a model worse than that constant. Because the denominator is N times the sample variance of y within $\mathcal{D}$, R^2 is computed for validation and test only, never per training batch, whose variance is not a meaningful reference for the spread of the target.

Comparisons derived from these metrics. Let $M(c)$ be any of (3.30), (3.31), (3.32) or (3.37) evaluated for configuration c , and id the full-graph baseline. Every relative accuracy figure in the results is one of

$$\varrho(c) = \frac{M(c)}{M(\text{id})}, \quad (3.33)$$

$$\Delta_{\text{state}}(c) = M(\hat{y}^{\text{full}}(c)) - M(\hat{y}^{\text{match}}(c)), \quad (3.34)$$

$$\Delta_{\text{pair}}(c_d, c_g) = M(c_d) - M(c_g), \quad (3.35)$$

the fraction of the baseline a reduction retains, the cross-state drop-off between the two passes of (3.27) and (3.28), and the gap within a matched pairing of a domain-informed configuration c_d against its generic control c_g . Equation (3.35) is also what RQ1a reports as protocol inflation, taken between two split protocols at one fixed configuration rather than between two reductions. A ratio is used for retention and a difference for the other two because R^2 crosses zero, where a ratio is meaningless.

Variance already carried by the corpus. A score is only evidence of structural reading to the extent it exceeds what the corpus gives away for free. For a partition of $\mathcal{D}$ into groups g , the group-mean predictor $\hat{y}_i = \bar{y}_{g(i)}$ scores

$$R_{\text{grp}}^2 = 1 - \frac{\sum_i (y_i - \bar{y}_{g(i)})^2}{\sum_i (y_i - \bar{y})^2}, \quad (3.36)$$

the share of label variance explained by group membership alone. Reporting (3.32) beside (3.36) for the source script and the optimization tier separates the part of a score attributable to the circuit from the part a lookup table would reach.

Persisted predictions and stratified re-scoring. The evaluation sweep writes every per-graph prediction to disk, not only the aggregates, so every metric above can be recomputed on any subset of the test set from the same forward pass. This matters because the pooled label is severely skewed: roughly half the test graphs have $y_i = 0$ exactly and about three quarters have $y_i < 0.005$, so a pooled R^2 is dominated by a mass of near-identical labels and can look respectable while saying nothing about the graphs a practitioner would ask about. Each accuracy claim is therefore repeated on three subsets beside the pooled figure: graphs from designs the synthesis script can still change (two designs sit at a fixed point, with largest observed optimizability 0.0003, so they contribute no variance for R^2 to explain, only noise for it to divide by); graphs with $y_i > 0$; and graphs derived from a single upstream script, which holds constant how much optimization the input received before Orchestrate saw it. The strata are applied uniformly to every configuration rather than invoked where they happen to help.

Ranking Quality

Spearman rank correlation between predicted and true optimizability across the test set, that is the Pearson correlation of the midrank vectors,

$$\rho = \frac{\text{cov}(\text{rk}(\hat{y}), \text{rk}(y))}{\sigma_{\text{rk}(\hat{y})} \sigma_{\text{rk}(y)}}. \quad (3.37)$$

Equation (3.37) is unchanged by any strictly increasing map applied to $\hat{y}$, which is why it is reported alongside the error metrics rather than in place of them: a model whose predictions have collapsed toward $\bar{y}$, the failure mode a reduction study should expect, holds ρ while RMSE and R^2 degrade. Ranking is also closer to use. A practitioner deciding where to spend synthesis effort needs the ordering of circuits, not calibrated values.

The metric ranks *circuits* under one fixed script, not *scripts* for one circuit. The latter is the script-selection task the motivation invokes and this thesis does not attempt (1.2.6).

Efficiency and Resource Metrics

Three quantities are reported: peak device memory, per-graph step time $t(G_i)$, and inference throughput

$$\tau = \frac{|\mathcal{D}|}{T(\mathcal{D})}, \quad (3.38)$$

graphs scored per second over an evaluation pass of wall-clock duration T . Epoch time is logged but not compared: the budget each run actually consumed varied, so an epoch carries the training schedule as much as the reduction, and t and τ are the timing quantities that compare. The per-graph step cost t is also what enters (3.40).

Two instruments produce these numbers, under different batching regimes, and which one produced a figure fixes what it may be compared against. Throughput and inference peak memory come from the evaluation sweep under the fixed evaluation node budget of 3.4.2, identical across configurations; comparing configurations evaluated at different budgets would be invalid rather than merely noisy. Step time and step memory come from a separate benchmark at one graph per batch and no budget at all. A node budget packs a different number of graphs into a batch for every configuration, so per-batch aggregates would average over incomparable batch compositions, whereas one graph per batch matches each reduced graph to its unreduced counterpart by graph identity, which is what the paired statistics of 3.4.3 require. The price is a controlled per-graph cost rather than an estimate of wall-clock training cost.

Memory quantities reported. Peak *allocated* memory is the high-water mark of live tensors and the closest measure of the model's actual demand. Peak *reserved* memory is what the caching allocator has taken from the driver, is never smaller, and reflects allocation history and fragmentation as much as demand. Allocated is the primary figure, but reserved sets the out-of-memory boundary and is recorded alongside it. Both are taken against a *memory floor* m_{floor} (parameters, optimizer state and device context, measured after warmup with no graph resident), so the marginal cost of the graph itself is

$$\Delta m = m_{\text{peak}} - m_{\text{floor}}. \quad (3.39)$$

Δm rather than m_{peak} is the quantity a reduction can move: m_{floor} is fixed by the architecture and identical across configurations, so savings quoted against m_{peak} are diluted toward zero by an additive constant having nothing to do with the reduction. The paired savings of (3.44) are taken on m_{peak} , so every reported memory saving is a lower bound on the saving in (3.39).

Reduction Quality Metrics

The node and edge retention ratios η_V and η_E of (2.23) are reported separately rather than collapsed into one compression figure, together with the offline wall-clock cost of computing the reduction.

The offline cost matters independently of the training saving. Let c_R be the cost of computing and storing one graph's

reduction, and t_{id} and t_R the per-graph step cost on the unreduced and the reduced graph. One epoch visits each graph once, so over n epochs reducing is cheaper than not reducing once

$$n > n^* = \frac{c_R}{t_{\text{id}} - t_R}. \quad (3.40)$$

Reporting the amortisation threshold n^* rather than a raw offline cost states the trade-off in the terms that decide it, and makes the failure case explicit: (3.40) is undefined for $t_R \geq t_{\text{id}}$, so a reduction that does not shorten a step never amortises, however cheap it is to compute. The reduced graph is stored once and reloaded, so n accumulates over every epoch of every run that reuses the artifact.

Effective receptive field. Hypothesis H1 claims that coarsening helps by contracting paths, and that claim needs a measurement of its own. For a node v , let

$$\mathcal{C}_k(v) = \{u \in \mathcal{V} : u \rightsquigarrow v \text{ along a directed path of length at most } k\} \quad (3.41)$$

be its k -hop fanin cone, and let $m_u = \mathbf{1}^\top x_u$ be the member count of node u , supplied directly by the count-based feature schema of 3.1 and equal to 1 on an unreduced graph. The metric is the member-weighted mean cone size over a node set $\mathcal{S}$ sampled from $\mathcal{V}(G)$,

$$\bar{c}_k(G) = \frac{1}{|\mathcal{S}|} \sum_{v \in \mathcal{S}} \sum_{u \in \mathcal{C}_k(v)} m_u, \quad (3.42)$$

evaluated at $k = L$, the number of encoder layers, before and after reduction. The weighting is what makes the comparison meaningful: counting super-nodes rather than their members would make $\bar{c}_L$ fall under every coarsening by construction, while weighting by m_u counts original gates in both states. H1 predicts $\bar{c}_L(R(G)) > \bar{c}_L(G)$ under summarization and $\bar{c}_L(R(G)) < \bar{c}_L(G)$ under sparsification. Unlike a DAG longest-path metric, (3.42) is well defined for every method, not only for those that guarantee an acyclic quotient.

Statistical Treatment

Every reported figure is an estimate over a finite set of test circuits and is given with a 95% percentile bootstrap interval over that set. For a statistic $\hat{\theta}$ computed on $\mathcal{D}$, with $\hat{\theta}^{*(1)}, \dots, \hat{\theta}^{*(B)}$ its values on $B = 2000$ resamples of $\mathcal{D}$ drawn with replacement, the interval is the pair of empirical quantiles

$$\left[\hat{\theta}_{(\alpha/2)}^*, \hat{\theta}_{(1-\alpha/2)}^* \right], \quad \alpha = 0.05, \quad (3.43)$$

and α is fixed at that value everywhere. The percentile form is used rather than a normal approximation because neither statistic it is applied to is close to normal. Equation (3.32) is a ratio of sums of squares over a label with a point mass at zero, and the per-graph savings below are skewed across a corpus whose graph sizes span orders of magnitude.

Efficiency comparisons are additionally paired per graph against the full-graph baseline. For each test graph the saving in an efficiency quantity q , either peak memory or step time, is taken relative to that graph's own baseline measurement,

$$\delta_i = 100 \left(1 - \frac{q_R(G_i)}{q_{\text{id}}(G_i)} \right), \quad (3.44)$$

and $\{\delta_i\}$ is summarised by its mean with (3.43) and a two-sided Wilcoxon signed-rank test of $H_0: \text{med}(\delta) = 0$. Both configurations are measured on the same graph, so the pairing in (3.44) is by construction, and the skew is why a rank test replaces a paired t -test [Demšar, 2006].

Accuracy is reported with the same interval, and what that interval covers has to be said plainly, because it is narrower than it appears. Resampling $\mathcal{D}$ quantifies sampling error over circuits, which is the error a single training run does allow to be estimated. It does not cover run-to-run variance from initialisation, data order and nondeterministic kernels, since each configuration here is trained exactly once (3.4.4). Single-run comparisons of this kind have been shown to reverse under reseeding [Reimers and Gurevych, 2017, Bouthillier et al., 2021], so a difference smaller than the training variance this design leaves unmeasured cannot be called real. RQ4 is where that bites, its expected effect being the smallest in the thesis: the defensible report there is the paired gap (3.35) with its sampling interval and an explicit statement that the seed variance is missing (5.4.1).

Intervals are the primary report and p -values are not. An interval carries both the size of an effect and the precision of the estimate, whereas a significance verdict carries neither and is routinely read as though it carried both. The one test retained is the Wilcoxon test above, where the pairing is by construction and the null of no saving is a question a reader asks directly. Its p -value is reported as a number against the single α of (3.43) rather than as a row of stars.

3.4.4 Reproducibility

Split assignment, the randomised reduction methods, and model initialisation are each seeded at a fixed value; the deterministic reduction methods produce identical output by construction. Seed values and the full experimental configuration are listed in A.4.

Seeding makes a single run repeatable, not robust. Each configuration is trained once, so differences between configurations are point estimates carrying no run-to-run variance, a limitation treated in 3.4.3.

Reduction is precomputed offline, and the artifacts behind every reported experiment are released as masks over the original graphs (3.4.4). Reproducing a result therefore does not require re-running the reduction stage.

Numerical Determinism

[TODO: this seems like overkill. never seen it in another paper feel free to prove me wrong]

Code and Data Availability

Results

Figures and tables are placed beside the analysis they support, so that each research question reads in one pass. Diagnostic material the argument does not rest on, the hyperparameter search above all, is marked as such in this chapter and moves to the appendix before submission.

No summarization configuration has been trained or measured. Every summarization row below is invented and the caption carrying it says so; where the family is absent instead, the result covers partitioning and sparsification only. Hypothesis H1 (1.2.3) is therefore asserted rather than tested, and nothing in this chapter is evidence for or against it.

4.1 Baseline Predictive Performance on Full AIGs (RQ1)

4.1.1 Comparison Against Naive Baselines

“Accurately” needs a referent. Two groups supply one (3.1.6): trivial predictors, which fix what a given R^2 is worth on this label, and published circuit models ported to this task.

The grouped comparison in 4.1 is the reference point for the whole chapter: it fixes what the encoder is worth against predictors that read nothing but the label distribution or the node and edge counts. The size-only ordinary least squares fit already ranks circuits better than the encoder does, so a claim that the encoder reads circuit *structure* rather than circuit *scale* has to clear that row before any reduction result means anything.




Figure 4.1: **PARTLY FABRICATED.** Every RQ1 baseline on the design-disjoint test split, by RMSE, R^2 and Spearman correlation. The trivial predictors, fitted on validation and scored on test, and the edge-aware GCN+ are measured; the published circuit models are invented placeholders. The size-only regressor outranks the encoder on Spearman correlation.

4.1.2 Predictive Accuracy Metrics

Smooth L1 loss, RMSE and R^2 for the edge-aware GCN+ predicting Orchestrate optimizability on a test set of full AIGs, reported together (3.4.3).




Figure 4.2: Predicted against true optimizability for the full-graph baseline on the design-disjoint test split, one point per circuit. The diagonal is perfect prediction: points below it are under-predictions, points above it over-predictions.

A parity plot is the standard first look at a regression fit, and it is here because two failures that no single error number exposes are visible in it directly. A model that has compressed toward the label mean shows as a cloud flattened onto a horizontal band instead of following the diagonal, however healthy its RMSE looks. A model with heteroscedastic error shows as a spread about the diagonal that widens with the target, which says the reported error is an average over regimes that behave differently rather than a property of the model as a whole. Both matter for a label that is 49% exactly zero.

Calibration takes the same question as 4.2 and answers it by bin rather than by point, which is the readable form when most of the mass sits at one end of the range.

The banded view in 4.4 is where the flattening of 4.3 becomes a number. The circuits the motivating use case cares about most, the ones with the most to remove, are exactly the ones the baseline predicts worst.

The decomposition in 4.5 separates the headline score into the part any group-mean predictor would get for free and the part attributable to reading circuit structure. The residual $R^2 = 0.124$ is the quantity this section defends, and it is the figure every reduction in RQ3 should be compared against rather than the pooled 0.343.




Figure 4.3: Binned reliability of the full-graph baseline: mean prediction against mean truth within bins of predicted optimizability. A model that has collapsed toward the label mean shows here as a flat curve regardless of how healthy its RMSE looks.




Figure 4.4: The full-graph baseline as a function of how much the circuit actually needs optimizing. The model tracks the truth on the two lowest bands, which hold 80% of the corpus, and then flattens: at a true optimizability above 0.20 it predicts 0.046 against a true mean of 0.226, an absolute error of 0.180. It has largely learned the label’s marginal distribution rather than the task.




Figure 4.5: How much of the full-graph baseline’s $R^2 = 0.343$ is reachable without learning anything about an individual circuit. Grouping by dataset tier carries only 1.1% of the label variance, so the tier is not a confound. Grouping by *source script* carries 25.0%: a predictor that outputs nothing but each group’s mean scores $R^2 = 0.250$, roughly three-quarters of the model’s reported figure. Held against within-group variance the model retains $R^2 = 0.124$.

4.1.3 Ranking Efficacy

Spearman correlation between predicted and true optimizability, which is the metric the script-selection use case in 1.1 actually depends on.




Figure 4.6: **FABRICATED.** Reserved slot for the ranking view of the RQ1 comparison: Spearman correlation between predicted and true optimizability for each baseline on the design-disjoint test split, with a bootstrap interval per bar. Nothing is plotted and no data stands behind this float. The numbers it would draw are the Spearman column of 4.1.

Ranking is reported separately from error throughout because the two can move in opposite directions, and 4.1 already shows them doing so: the encoder wins on RMSE and R^2 and loses on Spearman correlation to a two-parameter fit on graph size.

4.1.4 Baseline Model Outcomes

The published circuit models were adapted rather than reused, and the adaptations are stated in 3.1.6. What each comparison licenses depends on them.

The table collects all of RQ1 in one place, and its footnote records which rows are invented and which measured. Two cautions attach to it. SynthNet is the only published baseline that has produced a number at all, and it collapsed to a constant prediction on this split, so its RMSE looks strong while its R^2 is negative. Its published R^2 is not the same quantity as the one reported here, because the upstream work z-scores its targets per design and so removes between-design variance from the denominator; only the sign transfers.

Table 4.1: RQ1 baseline comparison on the design-disjoint test split. Trivial predictors are fitted on validation and scored on test. Upstream R^2 figures are NOT comparable: OpenABC-D z-scores its targets per design, which removes between-design variance from the denominator. Only the sign transfers.

Model	Tier	RMSE	MAE	R^2	Spearman
Predict validation mean	1: trivial	0.054	0.023	-0.052	–
Predict validation median	1: trivial	0.056	<u>0.020</u>	-0.131	–
Graph size only (OLS)	1: trivial	<u>0.050</u>	0.024	<u>0.101</u>	0.348
Edge-aware GCN+ (ours)	This work	0.043	0.017	0.343	<u>0.185</u>
[TODO/FAKE] SynthNet	2: published models	999999	999999	-999999	-999999
[TODO/FAKE] HOGA	2: published models	999999	999999	-999999	-999999
[TODO/FAKE] DeepGate4	2: published models	999999	999999	-999999	-999999

Rows in red are FABRICATED PLACEHOLDERS, not measurements: every number in them has been replaced by the sentinel ± 999999 . **Bold** marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings. Rows marked [TODO/FAKE] are invented. Finish the DeepGate4 and HOGA ports on the baselines/openabc-synthnet-hoga branch. SynthNet is the only published baseline that has produced a number (train_baseline_synthnet_Orchestrate, val R2 = -0.168); the HOGA run crashed at epoch 0.

4.1.5 Error Analysis

Residuals broken down by graph size, by tier and by source design, which sets the expectation for which reduction methods should be expected to hurt most.

Whether error grows with scale is the precondition for the reduction chapters: a method that acts hardest on large graphs starts from whatever position 4.7 shows, not from the pooled error.

A difference across tiers is expected, and its absence would itself be a finding about the generation pipeline rather than about the model. Against 4.5, which puts only 1.1% of the label variance on the tier, the tier is not the axis that structures this corpus.

The figure and the table give the same result in two forms, and it is the most consequential result in RQ1: a pooled score over a design-level split is an average over a handful of whole circuits, not over a random sample, so the pooled number describes the mix of designs at least as much as it describes the model. That is the observation RQ1a (4.1.6) turns into a measurement.




Figure 4.7: Absolute error of the full-graph baseline against graph size (median and interquartile band), with the number of graphs per size bin alongside so that a wide band on a thin bin is not read as a trend.

Table 4.2: Per-design metrics on the design-disjoint test set. A pooled score over a design-level split is an average over this handful of whole circuits, not over a random sample.

Design	Graphs	Mean y	RMSE	MAE	R^2	Spearman
256	16040	0.001	0.001	0.001	-0.548	0.493
16384	16080	0.000	0.001	0.001	-747.953	0.043
picosoc	16045	0.007	0.010	0.006	-0.657	0.265
c1355	16102	0.006	0.026	0.006	-0.014	0.331
c6288	16052	0.044	0.055	0.038	0.347	0.539
square	16111	0.062	0.085	0.049	0.091	0.884




Figure 4.8: Absolute error of the full-graph baseline by dataset tier. Tiers differ by construction, since a graph already optimized once has less left to remove.




Figure 4.9: Per-design RMSE and R^2 for the full-graph baseline on the design-disjoint test set. The spread across designs is very wide and R^2 is strongly negative on some of them.

4.1.6 Protocol Sensitivity: How Much of the Score Is Design Recognition? (RQ1a)

Three train and test protocols under otherwise identical conditions: a plain random split, a recipe-disjoint split, and the design-disjoint split used for every reported result (3.4.2).

Table 4.3: RQ1a protocol sensitivity. Identical encoder, budget and seed; only the split changes. Random is the common default in circuit representation-learning evaluations, recipe-disjoint is OpenABC-D Variant 1, design-disjoint is Variant 2 and is what this thesis reports.

Protocol	RMSE	R^2	Spearman	Inflation vs. design-disjoint
[TODO/FAKE] Random (per-row)	999999	-999999	-999999	999999
[TODO/FAKE] Recipe-disjoint	999999	-999999	-999999	999999
Design-disjoint (ours)	0.043	0.343	0.185	0.000

Rows in red are FABRICATED PLACEHOLDERS, not measurements: every number in them has been replaced by the sentinel ± 999999 . Rows marked [TODO/FAKE] are invented. IN FLIGHT: the `-split_by random` and `-split_by recipe` training runs are on the cluster now (the design-disjoint run already exists). Identical encoder, budget and seed; only the split changes. When they land, run `src/test.py` for each, delete this block and rerun `analysis.make_all`.

The gap between the three rows is a direct measurement of how much apparent accuracy a permissive protocol buys, and until the two outstanding runs land, the comparison against published numbers obtained under leakier protocols cannot be made quantitative. The placeholder rows are set to plausible values only so that the figure has a shape; no conclusion may be drawn from their size.




Figure 4.10: **MOSTLY FABRICATED.** How much of the measured score is design recognition rather than structural inference. Only the design-disjoint row is measured; the random and recipe-disjoint rows are placeholders for two training runs that have not completed.

4.1.7 Cost of the Full-Graph Baseline

Three quantities characterise the cost of training on unreduced graphs, and together they form the reference against which every reduction in RQ2 is read: time per optimisation step and per epoch, peak allocated device memory, and inference throughput in graphs per second.

The swing bounds how much weight any single-seed comparison can carry (3.4.3). Every configuration in this thesis is trained exactly once, so a gap of a few percent between two configurations cannot be separated from the epoch-to-epoch movement visible here.

Consumed budget rather than allotted budget is what the RQ2 speedups are normalised against, since early stopping makes the two diverge.

Memory pressure being the dominant failure mode of the search is the hyperparameter-search face of the bottleneck this thesis is about, which is the one reason the sweep is reported in the body at all.

Table 4.4: Final configuration, as recorded in the run log of the reported run. Fields the run did not log are marked rather than omitted.

Hyperparameter	Value
Encoder	gcn
Hidden dim	128.0
Layers (message passing)	4 (not logged)
Positional encoding	level
PE dim	32.0
Pooling	mean
Head dropout	0.3
Learning rate	0.0003
Min learning rate	1e-06
Weight decay	0.00396
Warmup steps	10000.0
Scheduler patience	2.0
Node input dim	4.0
Edge attr dim	2.0
Seed	not recorded

The sensitivity view supports “which settings never produced a good trial” and not “which parameter matters most”. The table records the final configuration as logged, marking the fields the run did not log rather than omitting them; it moves to the appendix with the two figures above.

Training trajectories, 11 Orchestrate runs (early stopping on validation loss)




Figure 4.11: Validation loss and R^2 per epoch for every Orchestrate training run, coloured by reduction family. Validation R^2 swings by several units between consecutive epochs on some configurations.




Figure 4.12: Wall-clock time and epochs actually consumed per configuration, against epochs to best validation loss. Two runs of the same nominal budget are not the same amount of compute if one stopped at epoch 4 and the other at 10.

Hyperparameter search: memory pressure is the dominant failure mode

Sweep progress (6 outliers clipped) Objective by encoder

Trial outcomes (90 total)




Figure 4.13: **APPENDIX.** The hyperparameter sweep that fixed the architecture: progress per worker, objective by encoder, and trial outcomes. A large share of trials were pruned on memory exhaustion rather than on their objective. Reconstructed from the worker logs, because the study databases were lost with the scratch workspaces.

Hyperparameter sensitivity over 46 completed trials (dashed: best objective; bar: group median)




Figure 4.14: **APPENDIX.** Objective against each searched hyperparameter over the 52 completed trials. A coarse sensitivity read rather than an importance analysis: sampling concentrates on promising regions and the sampler state was lost with the scratch workspaces.

4.2 Efficiency Profile of Graph Reduction Techniques (RQ2)

4.2.1 Graph Reduction Offline Profile

How much each method removes, and what it costs to remove it. Node and edge retention are reported separately, because several methods preserve node counts exactly and a single compression figure would hide that (2.1.5).

Retention is the axis every later result is read against, and the spread across it is what makes a raw ranking of methods meaningless on its own: METIS keeps 93.7% of edges while random hashing keeps 40.2%, so their accuracies in RQ3 are not comparable as-is.

The width of these distributions is what makes the RQ4 pairing between AND-gate-only and PageRank pruning approximate rather than exact (4.4.1). A method whose compression varies with the circuit cannot be matched to a fixed ratio, only calibrated to one on average.

Offline cost is paid once per graph and reused by every epoch that reads the cached artifact, which is why it is reported as an amortisation threshold rather than as a raw number of seconds.

A reduction pays for itself once the epochs served exceed c_R/s , the offline cost divided by the per-epoch saving. Spanning forest has no such threshold at all: its measured time saving is -0.1% , so no number of epochs repays its 3.3 s per graph.

Edge cut and imbalance are the two quantities a partitioner trades against each other, and holding k fixed across all four reduces the comparison to that single plane. The pairs it defines are the cleanest domain-informed against generic comparisons in the thesis, and RQ4 uses them as such (4.4.1).

The probe is the cheapest of the outstanding summarization measurements, since it needs no training run, only an offline pass. Until it exists there is no upper bound on what exact coarsening could remove from this corpus.




Figure 4.15: **PARTLY FABRICATED.** Offline compression per method, reported separately for nodes and edges. The four partitioners keep every node by construction, as do two of the four sparsifiers. The summarization rows are invented, since no summarization configuration has been trained or measured.

Compression is a distribution, not a single ratio

Edge retention: distribution over 10,000 graphs Node retention: distribution over 10,000 graphs




Figure 4.16: Per-graph retention over 10,000 graphs. Compression is a distribution rather than a ratio: for the parameter-free methods it is a property of each circuit, and PageRank pruning's node retention is nearly a point mass at its configured 0.8 while AND-gate-only's is not.

Table 4.5: The four partitioners keep every node and differ only in which edges they cut, at the same k . That makes them the cleanest domain-informed / generic pairs in the thesis.

Partitioner	Kind	Mean k	Edge cut	Edge cut sd	Imbalance (sd nodes)	Max part. nodes	Offline (s)
Random hashing	Generic	4.662	0.598	0.162	37.891	6382	0.004
METIS	Generic	4.662	0.063	0.053	<u>30.433</u>	6370	3.196
Level slicing	Domain	4.662	0.295	0.236	0.378	6334	<u>1.151</u>
Span-weighted METIS	Domain	4.662	<u>0.077</u>	0.060	30.616	6371	3.701

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings.




Figure 4.17: **PARTLY FABRICATED.** Wall-clock cost of computing each reduction, per graph. The measured methods span two orders of magnitude, from 0.004 s for random hashing to 3.7 s for span-weighted METIS. The summarization entries are invented, since no summarization configuration has been run.




Figure 4.18: How many epochs over the corpus a cached reduction must serve before its offline cost is repaid by the training time it saves. This is the form in which 3.4.3 asks offline cost to be reported.




Figure 4.19: Edge cut against partition imbalance for the four partitioners at the same k . Since all four keep every node and differ only in which edges they cut, this plane is their entire comparison.




Figure 4.20: **ENTIRELY FABRICATED. NO SUMMARIZATION DATA EXISTS.** Node retention under colour refinement at depths $d = 1 \dots 4$. Every corpus graph is structurally hashed, so depth 1 should find almost nothing; how much survives beyond one hop bounds what the exact summarization track can compress and is a reportable statistic about AIGs in its own right.

4.2.2 Training Memory Dynamics and Compute Efficiency

Memory is reported as peak allocated device memory per configuration, in absolute terms and as a fraction of device capacity. Compute efficiency is reported as device utilisation alongside memory-bandwidth activity, which together distinguish a configuration that is compute-bound from one that is waiting on memory.

Peak memory tracks *node* retention, not edge retention. Activation memory in message passing scales with the number of nodes times the number of layers times the hidden width, whereas removing an edge only shortens the gather and scatter that feed those activations, which is transient work rather than stored state. The measured savings follow that exactly. Random edge dropout and spanning forest hold $|\mathcal{V}(R(G))|/|\mathcal{V}(G)| = 1$ and save 1.5% and 2.1% of mean peak memory while discarding 30.3% and 41.9% of edges respectively, and the four partitioners, which also keep every node, save between 0.3% and 3.1%. The two methods that actually remove nodes are the only ones with a double-digit saving: PageRank pruning at 0.800 node retention saves 14.5%, and AND-gate-only at 0.821 saves 11.3%. Reading the maximum instead of the mean flattens even that difference, because the largest circuit forms its own batch and floors every configuration alike. Sparsification, in short, is not a memory intervention on this encoder; it is a bandwidth intervention that happens to leave the memory bill intact.

Reporting a single peak per configuration would conflate the reduction’s effect with which graphs happened to land in the peak batch, which is why the cost profile is given against size rather than as a scalar.




Figure 4.21: Peak allocated device memory per configuration at three order statistics. The mean, p95 and maximum disagree by an order of magnitude: a graph larger than the batch budget forms a batch of its own, so the maximum is floored by the largest circuit in the corpus no matter what the reduction did to the average one.




Figure 4.22: Peak memory and step time as a function of graph size. In absolute terms every method traces the same memory curve, because memory tracks the graph rather than the reduction, so the two ratio panels are where any difference between methods is visible at all.

4.2.3 Throughput and Training Speedup

Step time and epoch time against the full-graph baseline, and inference throughput in graphs per second over the same reference.

The mismatch is large enough to invert the ranking. Spanning forest removes 41.9% of edges and returns no time at all, while METIS removes 6.3% and returns 24.3%, because the partitioners cut the batch into pieces the kernels schedule better rather than simply giving them less to do.

Savings are not normally distributed, so bare means would be misleading and the paired bootstrap is reported instead. Spanning forest's step-time saving sits inside the noise despite removing 42% of edges, which is the clearest single statement of the gap between compression and speed.




Figure 4.23: Mean step time and speedup against the full-graph baseline. Speedup under-delivers relative to compression, because the kernels are latency-bound on message-passing gather and scatter (3.4.2), so removing work does not translate proportionally into wall-clock time.




Figure 4.24: Per-graph peak-memory and step-time savings against the baseline, paired by graph id, with 95% percentile bootstrap intervals. A red p -value marks a Wilcoxon signed-rank test that fails to separate the saving from noise.

Table 4.6: Per-graph savings against the full-graph baseline, paired by graph id, with percentile bootstrap intervals and a Wilcoxon signed-rank test on the paired deltas. Savings are not normally distributed, so bare means would be misleading.

Method	Paired graphs	VRAM saving (%)	95% CI low	95% CI high	Time saving (%)	95% CI low	95% CI high	Wilcoxon p (time)
PageRank pruning	1000	14.492	14.121	14.842	19.145	18.081	20.203	0.000
AND-gate only	1000	<u>11.257</u>	10.702	11.794	28.893	27.627	30.054	0.000
Random hashing	1000	3.100	2.984	3.216	20.996	19.785	22.157	0.000
Spanning forest	1000	2.142	2.080	2.202	-0.101	-1.031	0.826	0.549
Random edge dropout	1000	1.489	1.451	1.526	0.229	-0.574	1.012	0.001
Level slicing	1000	1.487	1.383	1.596	18.291	17.322	19.300	0.000
Span-weighted METIS	1000	0.357	0.336	0.378	22.412	21.396	23.435	0.000
METIS	1000	0.279	0.262	0.297	<u>24.284</u>	23.237	25.303	0.000

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings.

4.2.4 Efficiency Ranking Across Methods

The same measurements as above, cut the other way: one row per method, ranked on what it buys. Generic against domain-informed is a column rather than the organising axis, because it is one attribute of a method.

Read down the columns and no method wins on all of them. AND-gate-only returns the largest time saving at 28.9% and the second-largest memory saving at 11.3% for 0.099 s of offline work per graph, which makes it the best-value reduction measured here on efficiency alone; RQ3 then asks what it costs in accuracy. The summarization rows are invented and carry no efficiency measurement at all, so the ranking covers partitioning and sparsification only.

Table 4.7: RQ2 efficiency profile. Node and edge retention are reported separately because several methods preserve node counts exactly. Savings are per-graph means paired against the full-graph baseline.

Method	Family	Kind	Node ret.	Edge ret.	Offline (s/graph)	Step time (s)	VRAM mean (MB)	VRAM saving (%)	Time saving (%)
Random hashing	Partition	Generic	1.000	0.402	0.004	0.026	1002.217	3.100	20.996
METIS	Partition	Generic	1.000	0.937	3.196	0.027	1056.507	0.279	<u>24.284</u>
Level slicing	Partition	Domain	1.000	0.705	1.151	0.029	1019.120	1.487	18.291
Span-weighted METIS	Partition	Domain	1.000	0.923	3.701	0.027	1055.088	0.357	22.412
Random edge dropout	Sparsification	Generic	1.000	0.697	<u>0.064</u>	0.040	1040.969	1.489	0.229
Spanning forest	Sparsification	Generic	1.000	0.581	3.309	0.041	1031.674	2.142	-0.101
PageRank pruning	Sparsification	Generic	0.800	0.627	1.739	0.029	850.969	14.492	19.145
AND-gate only	Sparsification	Domain	0.821	0.733	0.099	0.026	<u>892.163</u>	<u>11.257</u>	28.893
[TODO/FAKE] Random within-type	Summarization	Generic	999999	999999	999999	999999	999999	-	-
[TODO/FAKE] ConvMatch	Summarization	Generic	999999	999999	999999	999999	999999	-	-
[TODO/FAKE] Cone coarsening	Summarization	Domain	999999	999999	999999	999999	999999	-	-
[TODO/FAKE] MFFC coarsening	Summarization	Domain	999999	999999	999999	999999	999999	-	-
[TODO/FAKE] Colour refinement (exact)	Summarization	Domain	999999	999999	999999	999999	999999	-	-

Rows in red are FABRICATED PLACEHOLDERS, not measurements: every number in them has been replaced by the sentinel ± 999999 . **Bold** marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings. Rows marked [TODO/FAKE] are invented. No summarization configuration has been trained or measured. Needs: `sbatch src/shell/precompute_summarization.sh` for the offline pass, then one training run and one `src/test.py` pass per method. Random within-type merging is not implemented at all.

4.3 Predictive Retention and Performance Trade-offs (RQ3)

4.3.1 Accuracy Degradation

How much accuracy survives training and testing under the same reduction, against the RQ1 full-graph baseline. RMSE, R^2 and Spearman correlation are reported together throughout, because a reduction can hold mean error steady while destroying explained variance (3.4.3).

No reduction beats the full graph on explained variance. Random edge dropout retains the most, $R^2 = 0.215$ against the baseline's 0.343, and the two METIS variants fall to -3.062 and 0.008 . Ranking behaves differently: PageRank pruning reaches a Spearman correlation of 0.367 against the baseline's 0.185, so five of the eight measured reductions order circuits *better* than the unreduced model does while explaining less of their variance. That dissociation is the RQ3 result, and it is the reason the two metrics are never reported apart.

Holding error steady while losing all explanatory power is what a model does when it collapses onto the label's marginal distribution, and on a corpus that is 49% exactly zero that collapse is cheap.

PageRank pruning is the exception, scoring $R^2 = 0.204$ and $\rho = 0.521$ on tier 0, better than the unreduced baseline on that tier. A reduction outperforming the full graph on a subgroup is the first sign that the reduction is acting as a regulariser rather than as a loss of information.

The partitioners degrade catastrophically on the C2RS-derived subgroup, METIS reaching $R^2 = -736$, where the labels sit near zero and it predicts high optimizability regardless. A pooled score cannot show that, and the source script is the axis on which it hides.




Figure 4.25: **PARTLY FABRICATED.** RMSE, R^2 and Spearman correlation under the reduction each model was trained with, against the full-graph baseline. The three metrics disagree about the ranking. The summarization rows are invented, since no summarization configuration has been trained.

Table 4.8: RQ3 matched-state accuracy: models trained and tested under the same reduction. RMSE, R^2 and Spearman are reported together throughout, because a reduction can hold mean error steady while destroying explained variance.

Method	Family	Kind	Smooth L1	RMSE	R^2	Spearman	R^2 retained
Full graph	Baseline	Generic	0.014	0.043	0.343	0.185	1.000
Random hashing	Partition	Generic	0.038	0.075	-0.981	-0.183	-2.861
METIS	Partition	Generic	0.055	0.107	-3.062	0.237	-8.932
Level slicing	Partition	Domain	0.018	0.057	-0.142	0.119	-0.414
Span-weighted METIS	Partition	Domain	0.034	0.053	0.008	0.263	0.024
Random edge dropout	Sparsification	Generic	<u>0.016</u>	<u>0.047</u>	<u>0.215</u>	0.194	<u>0.627</u>
Spanning forest	Sparsification	Generic	0.016	0.049	0.138	<u>0.282</u>	0.404
PageRank pruning	Sparsification	Generic	0.019	0.053	-0.018	0.367	-0.051
AND-gate only	Sparsification	Domain	0.018	0.055	-0.064	-0.025	-0.188
[TODO/FAKE] Random within-type	Summarization	Generic	-	999999	-999999	-999999	-999999
[TODO/FAKE] ConvMatch	Summarization	Generic	-	999999	-999999	-999999	-999999
[TODO/FAKE] Cone coarsening	Summarization	Domain	-	999999	-999999	-999999	-999999
[TODO/FAKE] MFCC coarsening	Summarization	Domain	-	999999	-999999	-999999	-999999
[TODO/FAKE] Colour refinement (exact)	Summarization	Domain	-	999999	-999999	-999999	-999999

Rows in red are FABRICATED PLACEHOLDERS, not measurements: every number in them has been replaced by the sentinel ± 999999 . **Bold** marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings. Rows marked [TODO/FAKE] are invented. No summarization configuration has been trained or measured. Needs: sbatch src/shell/precompute_summarization.sh for the offline pass, then one training run and one src/test.py pass per method. Random within-type merging is not implemented at all.

Table 4.9: Matched-state accuracy partitioned two ways: by dataset tier, and by the synthesis script that produced the graph. Both partition the same test split, so the columns within each cut are disjoint. The source script separates the corpus far more sharply than the tier does.

Method	R^2 C2RS	R^2 Deepsyn	R^2 Syn4	R^2 tier0	R^2 tier0 base	R^2 tier1	ρ C2RS	ρ Deepsyn	ρ Syn4	ρ tier0	ρ tier0 base	ρ tier1
Full graph	-1.039	-0.097	<u>0.143</u>	<u>-0.038</u>	<u>-0.038</u>	0.358	-0.153	0.167	0.475	0.044	0.044	0.234
Random hashing	-127.354	-28.352	-0.766	-4.012	-4.012	-0.823	-0.433	-0.319	0.024	-0.298	-0.298	-0.156
METIS	-736.099	-164.691	0.954	-20.880	-20.880	-2.046	0.030	<u>0.208</u>	<u>0.758</u>	<u>0.225</u>	<u>0.225</u>	0.241
Level slicing	-0.111	-0.068	-0.558	-0.172	-0.172	-0.154	0.102	0.193	0.140	0.021	0.021	0.149
Span-weighted METIS	-55.986	-14.553	0.074	-1.116	-1.116	0.064	-0.235	0.030	0.643	-0.050	-0.050	0.370
Random edge dropout	-0.649	0.071	-0.045	-0.065	-0.065	<u>0.223</u>	<u>0.125</u>	0.118	0.610	0.026	0.026	0.239
Spanning forest	-0.342	<u>-0.006</u>	-0.153	-0.108	-0.108	0.143	-0.133	0.289	0.786	0.021	0.021	<u>0.368</u>
PageRank pruning	-3.507	-0.626	-0.390	0.204	0.204	-0.043	0.418	0.194	0.496	0.521	0.521	0.312
AND-gate only	<u>-0.178</u>	-0.070	-0.444	-0.169	-0.169	-0.071	-0.427	-0.154	0.506	-0.504	-0.504	0.129

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings.

A small RMSE does not imply explained variance
(R^2 clipped at -1)




Figure 4.26: RMSE against R^2 , one point per configuration. Several configurations hold RMSE close to the baseline while R^2 falls below zero, which is the same dissociation [Khadka et al. \[2026\]](#) report on a coarsened EDA graph.

Matched-state accuracy by dataset tier
tier 0 is a base graph; tier 1 has already been through one synthesis script




Figure 4.27: Every configuration split across the two stored tiers. The full-graph baseline scores $R^2 = -0.038$ on tier-0 base graphs against $+0.358$ on tier-1: its absolute error on tier 0 is smaller (RMSE 0.025 against 0.047), but tier-0 labels have so much less spread that the model fails to beat their mean.




Figure 4.28: The same partition taken by the synthesis script that produced each graph, which separates the corpus far more sharply than the tier does (4.5). Every configuration degrades from Syn4-derived graphs to C2RS-derived ones.

4.3.2 Ranking Preservation

Spearman correlation under reduction. Ranking can survive when absolute error degrades, which is the property the script-selection use case actually needs.

Two conclusions hold across all four strata: no reduction beats the full graph on R^2 , and sparsification beats partitioning. Two do not. Random hashing's R^2 moves from -0.98 to -0.06 once the fixed-point designs are excluded, and METIS is worst overall yet best on Syn4-derived graphs, which is the signature of a model that collapsed onto predicting high optimizability rather than of a failure of partitioning.

Neither of those two is visible in R^2 alone, and both are large enough to change which method a practitioner would pick. RQ5 takes the cross-state result up as a question in its own right (4.5.2).

**Robustness of the matched-state ranking to the label skew
(a spread along a row means the conclusion depends on which graphs are scored)**




Figure 4.29: The matched-state result re-scored on four subsets of the same test split, from the persisted per-graph predictions and with no retraining. The strata answer different questions rather than being more or less trustworthy: excluding the two fixed-point designs selects on a design property, restricting to $y > 0$ selects on the outcome variable so its R^2 is not comparable to the pooled column, and the Syn4 column is the easiest stratum.

Table 4.10: Matched-state accuracy re-scored on four subsets of the same test split, from the persisted per-graph predictions. The pooled label is 49% exactly zero, so a pooled R^2 is carried by a minority of graphs; a conclusion that survives all four columns is a property of the reduction rather than of the label distribution.

Method	R^2 All test graphs	R^2 Designs Orchestra can touch	R^2 Graphs with $y > 0$	R^2 Syn4-derived graphs only	ρ All test graphs	ρ Designs Orchestra can touch	ρ Graphs with $y > 0$	ρ Syn4-derived graphs only
Random edge dropout	0.215	0.192	<u>0.097</u>	-0.045	0.194	0.225	0.427	0.610
Spanning forest	<u>0.138</u>	<u>0.113</u>	0.007	-0.153	<u>0.282</u>	<u>0.300</u>	0.524	0.786
Span-weighted METIS	0.008	-0.021	0.125	<u>0.074</u>	0.263	0.040	0.364	0.643
PageRank pruning	-0.018	-0.047	-0.155	-0.390	0.367	0.179	0.305	0.496
AND-gate only	-0.064	-0.095	-0.230	-0.444	-0.025	0.179	0.312	0.506
Level slicing	-0.142	-0.175	-0.320	-0.558	0.548	0.548	<u>0.484</u>	0.140
Random hashing	-0.981	-0.061	-0.276	-0.766	-0.183	0.134	0.433	0.024
METIS	-3.062	-3.181	-1.997	0.954	0.237	0.086	0.276	<u>0.758</u>

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings. Strata: all test graphs; excluding the two designs on which Orchestra is at a fixed point (16384, 8192, max $y = 0.0003$ over 32,000 graphs); graphs with a non-zero target; and graphs derived from Syn4, the one source script that leaves substantial optimizability behind.

**Robustness of the cross-state (full graphs) ranking to the label skew
(a spread along a row means the conclusion depends on which graphs are scored)**




Figure 4.30: The same stratified treatment applied to the full-graph evaluation pass, in which each model is queried on unreduced graphs. Ranking separates sharply from error: level slicing reaches a Spearman correlation of 0.601 on full graphs, above the full-graph baseline's 0.185, while AND-gate-only reaches -0.467 and orders circuits backwards.

4.3.3 Accuracy Loss Attributable to Reduction

What each method costs and what structural damage explains it. The comparison is made at matched compression, since otherwise “which method is best” confounds a genuinely better method with one that simply reduced less.

The plane is what makes the confound concrete: AND-gate-only keeps 73.3% of edges against spanning forest’s 58.1%, so their raw accuracies are not comparable and the formal paired version of the comparison belongs to RQ4 (4.4).

Expressing accuracy as a retained fraction rather than an absolute puts it on the same footing as the memory and time savings of RQ2, which is what the Pareto analysis below needs.

If the use case is finding circuits worth optimizing rather than scoring all of them, the pooled ranking inverts and the worst method on R^2 becomes the best one. The banded view is the only place in the chapter where a fully collapsed model is distinguishable from a merely weak one.

Hypothesis H1 predicts that summarization retains more than sparsification at matched compression, because contracting a path shortens the distance a signal must travel while removing an edge can only impede it. Testing it needs both the receptive-field metric and at least one trained summarization configuration, and neither exists, so H1 is asserted rather than tested and must not be reported as evidenced.

Accuracy against achieved compression, the only fair comparison plane




Figure 4.31: Matched-state R^2 against edge and node retention, one point per configuration. Without this plane a ranking of methods confounds the quality of a heuristic with the amount it happened to remove.

Accuracy retained under reduction (1.0 = no loss; below 0 = worse than the mean)




Figure 4.32: Each metric as a fraction of the full-graph baseline, one row per method, so that the figure lines up directly against the efficiency ranking of 4.2.4.

Table 4.11: Mean prediction and mean absolute error within bands of the true optimizability. R^2 is not reported: inside a narrow band there is almost no label variance to divide by. Compare each method's predicted mean against the true mean in the first row. A row that stays flat across the bands is a model that has collapsed to a constant.

Method	pred $y = 0$	pred (0, .01]	pred (.01, .05]	pred (.05, .10]	pred (.10, .20]	pred > .20	MAE $y = 0$	MAE (0, .01]	MAE (.01, .05]	MAE (.05, .10]	MAE (.10, .20]	MAE > .20
(true mean in band)	0.000	0.002	0.021	0.065	0.161	0.226	—	—	—	—	—	—
Full graph	0.004	0.003	0.005	0.006	0.052	0.046	0.004	0.003	0.016	0.059	0.109	0.180
Random hashing	0.042	0.007	0.002	0.003	0.008	0.012	0.042	0.007	0.019	0.062	0.153	0.214
METIS	0.047	0.075	0.100	0.109	0.171	0.180	0.047	0.073	0.093	0.083	0.020	0.046
Level slicing	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.002	0.021	0.065	0.161	0.226
Span-weighted METIS	0.030	0.032	0.036	0.016	0.056	0.052	0.030	0.029	0.027	0.049	0.105	0.174
Random edge dropout	0.004	0.005	0.006	0.006	0.023	0.047	0.004	0.004	0.015	0.059	0.138	0.179
Spanning forest	0.002	0.002	0.004	0.003	0.024	0.029	0.002	0.003	0.017	0.062	0.137	0.197
PageRank pruning	0.007	0.008	0.008	0.020	0.010	0.013	0.007	0.007	0.013	0.045	0.151	0.213
AND-gate only	0.002	0.001	0.000	0.000	0.006	0.009	0.002	0.002	0.021	0.065	0.155	0.217




Figure 4.33: Every configuration as a function of how optimizable the circuit truly is. Level slicing predicts 0.0000 in every band, a fully collapsed model that the pooled metrics do not expose. METIS over-predicts everywhere, which is why its pooled R^2 is -3.06 , yet that same miscalibration makes it the *most* accurate configuration on the highly optimizable circuits: MAE 0.046 above $y = 0.20$ against the unreduced baseline's 0.180.




Figure 4.34: **ENTIRELY FABRICATED. NO SUMMARIZATION DATA EXISTS.** Mean 4-hop fanin-cone size after reduction relative to before. The metric is specified in 3.4.3 and is not implemented.




Figure 4.35: **ENTIRELY FABRICATED. NO SUMMARIZATION DATA EXISTS.** The form the H1 test would take: if coarsening helps by contracting paths, accuracy should rise with the receptive-field ratio. Both axes are invented.

4.3.4 The Pareto Front (Trade-Off Analysis)

Accuracy against the memory and speed gains of RQ2, on one pair of axes, which is where RQ2 and RQ3 are answered jointly rather than separately.

Table 4.12: Accuracy against efficiency, ranked by matched-state R^2 .

Method	R^2	Spearman	VRAM saving (%)	Time saving (%)
Full graph	0.343	0.185	0.000	0.000
Random edge dropout	<u>0.215</u>	0.194	1.489	0.229
Spanning forest	0.138	<u>0.282</u>	2.142	-0.101
Span-weighted METIS	0.008	0.263	0.357	22.412
PageRank pruning	-0.018	0.367	14.492	19.145
AND-gate only	-0.064	-0.025	<u>11.257</u>	28.893
Level slicing	-0.142	0.119	1.487	18.291
Random hashing	-0.981	-0.183	3.100	20.996
METIS	-3.062	0.237	0.279	<u>24.284</u>

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings.

The baseline being on the frontier is the outcome to read first: on this corpus, at this scale, no reduction is worth its accuracy cost if explained variance is the objective. The argument for reduction has to be made either on a subgroup, on ranking, or at a scale where the full graph does not fit at all.

Every configuration on the frontier above was selected on validation, so the gap is a direct bound on how much of that selection survives the move to held-out designs.




Figure 4.36: Accuracy against efficiency, with the dominance frontier drawn. The full-graph baseline sits on the same axes at zero saving, so “no reduction” competes directly against every reduction rather than sitting outside the comparison.




Figure 4.37: The same configurations scored on validation and on test. Under a design-level split these are different circuits rather than two samples of one population, so the size of this gap bounds how much of the validation-selected model choice transfers.

4.4 Value of Domain-Informed Adaptation (RQ4)

4.4.1 Matched-Compression Pairings

4.4.2 Retention Gap at Matched Compression

4.4.3 Structural Interpretation

4.4.4 Adequacy of the Heuristics Tested




Figure 4.38: Domain-informed minus generic within each pairing. The first two pairings are matched by construction (same k , all nodes kept, only the cut rule differs); the third is matched by calibration; the fourth is *not* matched and is hatched accordingly. Read every bar against 4.39.




Figure 4.39: How closely matched each RQ4 pairing actually is. The spanning-forest / random-edge-dropout pair is far apart because the configured drop rate of 0.3 leaves 69.7% of edges against spanning forest’s 58.1%; raising the rate to approximately 0.419 would match them.




Figure 4.40: **ERROR BARS FABRICATED.** The measured RQ4 gaps with an invented ± 0.05 band. Every configuration in this thesis is trained exactly once, so there is no run-to-run variance to judge these gaps against. The figure exists to make that visible rather than to report an interval (4.4.2).

4.5 Cross-Structural Generalization (RQ5)

4.5.1 Reduced-to-Full Inference Accuracy

Smooth L1 Loss, RMSE, and R^2 metrics from testing a model (which was trained exclusively on reduced AIGs via partitioning, sparsification, or summarization) directly on normal, full-sized AIGs.

4.5.2 Zero-Shot Ranking

Spearman correlation.

4.5.3 Matched-State vs. Cross-State Drop-Off

Comparison measuring the performance drop-off between matched-state testing (Reduced $\rightarrow$ Reduced) vs. cross-state testing (Reduced $\rightarrow$ Full).

4.5.4 Generalization by Reduction Family

4.5.5 Inference Cost

Unpaired comparison over all 8 measured reductions (bar = group mean)




Figure 4.41: Every domain-informed method against every generic one on the same axes. Weaker evidence than the pairings, since the two groups did not remove the same amount, but it answers the question a reader asks first.

Table 4.13: RQ4 pairings. Δ is domain-informed minus generic, so a positive ΔR^2 favours the domain heuristic. The fourth pairing is NOT matched: random edge dropout at its configured rate keeps 69.7% of edges against spanning forest’s 58.1%, so its gap conflates the heuristic with how much each arm removed.

Domain-informed	Generic	Matched	Edge ret. (dom.)	Edge ret. (gen.)	Δ RMSE	ΔR^2	Δ Spearman
Span-weighted METIS	METIS	by construction	0.923	0.937	-0.054	3.071	0.026
Level slicing	Random hashing	by construction	0.705	0.402	-0.018	0.839	0.302
AND-gate only	PageRank pruning	by calibration	0.733	0.627	0.001	-0.047	-0.393
Spanning forest	Random edge dropout	NOT MATCHED	0.581	0.697	0.002	-0.077	0.088

Each configuration is trained ONCE. These are point estimates with no run-to-run variance behind them, and RQ4’s expected effect is small: a gap of a few percent cannot be distinguished from noise.

**Matched-state (open) vs. full-graph (filled) evaluation of the same weights.
Dashed: the full-graph baseline. A red connector means the shift hurt.**




Figure 4.42: The same weights evaluated under the reduction they were trained with and on unreduced graphs. For every method here a full graph needs no conversion, since it already is a valid input, which is what makes this a direct test rather than a comparison across two input formats (3.2.3).




Figure 4.43: Full-graph score minus matched-state score, by method and family. A positive bar means the model scores *better* on graphs it never trained on, which would say the reduction was hurting the model rather than the evaluation.




Figure 4.44: Matched-state against full-graph R^2 . The diagonal is perfect transfer; below it the structural shift cost accuracy, above it the reduction was the thing doing damage.

Cost of serving a query, by evaluation state




Figure 4.45: Throughput and peak memory for a query, under each evaluation state. Inference stores neither gradients nor activations, so the memory argument that forced the reduction at training time need not bind here at all, which is the practical claim of 1.2.5.




Figure 4.46: **MOSTLY FABRICATED.** The “infer full on modest hardware” half of RQ5. Every surviving inference CSV records device=cuda; the CPU bars are placeholders for a src/shell/test_cpu.sh pass that has not been run.




Figure 4.47: **ENTIRELY FABRICATED.** Colour refinement at count-cap ∞ merges only nodes the encoder provably cannot distinguish, so a model trained on those graphs *must* land on the dashed baseline when queried on full graphs. Without this control a poor transfer elsewhere cannot be told apart from a broken evaluation path, which is why the slot is held open rather than dropped.

4.6 Summary of Findings

Table 4.14: RQ5: the same weights evaluated under the reduction they were trained with and on unreduced graphs. For every method here a full graph needs no conversion, it already is a valid input. The exact colour-refinement track is the exception and has not been run.

Method	Family	R^2 matched	R^2 full	ΔR^2	ρ matched	ρ full	Throughput full (g/s)
Full graph	Baseline	–	0.343	–	–	0.185	74.650
Random hashing	Partition	-0.981	-0.071	0.910	-0.183	0.145	74.669
METIS	Partition	-3.062	-2.663	0.399	0.237	0.225	74.668
Level slicing	Partition	-0.142	-0.142	0.000	0.119	0.601	74.690
Span-weighted METIS	Partition	0.008	-0.352	-0.360	0.263	-0.132	74.690
Random edge dropout	Sparsification	0.215	<u>0.247</u>	0.032	0.194	0.326	74.658
Spanning forest	Sparsification	<u>0.138</u>	0.038	-0.100	<u>0.282</u>	<u>0.411</u>	74.721
PageRank pruning	Sparsification	-0.018	-0.114	-0.097	0.367	0.004	74.721
AND-gate only	Sparsification	-0.064	-0.144	-0.080	-0.025	-0.467	74.684

Bold marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings.




Figure 4.48: **ENTIRELY FABRICATED.** The summarization family in the compression/accuracy plane and its offline cost. No summarization configuration has been trained or measured, and random within-type merging, the instrument that makes RQ4’s third pairing kind answerable at all (3.2.3), is not implemented.




Figure 4.49: **ENTIRELY FABRICATED.** Node retention under colour refinement at depths $d = 1 \dots 4$. Every corpus graph is structurally hashed, so depth 1 should find almost nothing; how much survives beyond one hop is a reportable statistic about AIGs in its own right and bounds what the exact track can compress.




Figure 4.50: **ENTIRELY FABRICATED.** Mean 4-hop fanin-cone size after reduction relative to before. The metric is specified in 3.4.3 and *not implemented*; until it is, H1 is asserted rather than tested and must not be reported as evidenced.

4.6.1 Consolidated Summary




Figure 4.51: **ENTIRELY FABRICATED.** The form the H1 test would take: if coarsening helps by contracting paths, accuracy should rise with the receptive-field ratio. Both axes are invented.

Table 4.15: Every reduction method in one row: what it removes, what it costs offline, what it saves in training, what accuracy survives, and whether the model transfers to full graphs. Generic vs. domain-informed is a column, not a split, because it is one attribute of a method rather than the organising axis.

Method	Family	Kind	Node ret.	Edge ret.	Offline (s)	VRAM sav. (%)	Time sav. (%)	RMSE	R^2	ρ	R^2 cross-state
Full graph	Baseline	Generic	–	–	–	–	–	0.043	0.343	0.185	0.343
Random hashing	Partition	Generic	1.000	0.402	0.004	3.100	20.996	0.075	-0.981	-0.183	-0.071
METIS	Partition	Generic	1.000	0.937	3.196	0.279	<u>24.284</u>	0.107	-3.062	0.237	-2.663
Level slicing	Partition	Domain	1.000	0.705	1.151	1.487	18.291	0.057	-0.142	0.119	-0.142
Span-weighted METIS	Partition	Domain	1.000	0.923	3.701	0.357	22.412	0.053	0.008	0.263	-0.352
Random edge dropout	Sparsification	Generic	1.000	0.697	<u>0.064</u>	1.489	0.229	<u>0.047</u>	<u>0.215</u>	0.194	<u>0.247</u>
Spanning forest	Sparsification	Generic	1.000	0.581	3.309	2.142	-0.101	0.049	0.138	<u>0.282</u>	0.038
PageRank pruning	Sparsification	Generic	0.800	0.627	1.739	14.492	19.145	0.053	-0.018	0.367	-0.114
AND-gate only	Sparsification	Domain	0.821	0.733	0.099	<u>11.257</u>	28.893	0.055	-0.064	-0.025	-0.144
[TODO/FAKE] Random within-type	Summarization	Generic	999999	999999	999999	–	–	999999	-999999	-999999	-999999
[TODO/FAKE] ConvMatch	Summarization	Generic	999999	999999	999999	–	–	999999	-999999	-999999	-999999
[TODO/FAKE] Cone coarsening	Summarization	Domain	999999	999999	999999	–	–	999999	-999999	-999999	-999999
[TODO/FAKE] MFFC coarsening	Summarization	Domain	999999	999999	999999	–	–	999999	-999999	-999999	-999999
[TODO/FAKE] Colour refinement (exact)	Summarization	Domain	999999	999999	999999	–	–	999999	-999999	-999999	-999999

Rows in red are FABRICATED PLACEHOLDERS, not measurements: every number in them has been replaced by the sentinel ± 999999 . **Bold** marks the best value in a column and underline the second best, over measured rows only. Lower is better for error, memory, time and offline cost; higher is better for R^2 , correlation and savings. Rows marked [TODO/FAKE] are invented. No summarization configuration has been trained or measured. Needs: `sbatch src/shell/precompute_summarization.sh` for the offline pass, then one training run and one `src/test.py` pass per method. Random within-type merging is not implemented at all.

Discussion

5.1 Comparison with Related Work

5.1.1 Positioning Against Optimizability Prediction Work

Published optimizability figures cannot be placed beside the ones reported here until the protocols that produced them are matched. Most sit in the leakier cells of 2.2.1, on random or recipe-disjoint splits, whereas every headline result in this thesis is design-disjoint. RQ1a is the translation between the two: the gap between the three protocols at one fixed configuration is the correction that would have to be applied before a published number and a number from this work could share a column (4.1.6). That measurement does not exist yet, so the comparison below stays qualitative.

SynthNet’s collapse to a constant prediction is the same observation seen from the model side (4.1.4). A model developed and reported under a protocol that keeps source designs visible carries no obligation to survive one that removes them. The collapse is therefore what the protocol audit anticipates rather than evidence of a faulty port.

5.1.2 Positioning Against Graph Reduction Work

The exact-compression result this work depends on is not proved here. Bollen et al. [2023] establish that coarsening by colour refinement to depth d leaves the output of a d -layer message-passing network unchanged. What this thesis adds is the set of conditions under which that theorem applies to an AIG encoder, stated with the reduction families in 2.1.5 and instantiated for the architecture in 3.1.5. Those conditions constrain the model and the input schema. Establishing them is an argument, not a measurement, so it belongs where the encoder is defined rather than in the results.

The empirical half is a positive control. If the conditions hold, the exact configuration must reproduce the full-graph baseline to within numerical tolerance, both under matched-state evaluation and when queried on unreduced graphs (3.4.1). A deviation is a defect to be located rather than a finding about coarsening, which is why the control is reported ahead of everything else it underwrites (4.3.1, 4.5.1).

What the proof licenses stops at that one configuration. It says nothing about the approximate summarizers, whose merges are not refinement classes, and nothing about depths below the layer count, where refinement merges nodes the encoder can still tell apart. Every accuracy claim about the approximate track is consequently empirical. The proof marks how much compression is available for free; it does not extend that guarantee to the family around it.

5.1.3 Where Direct Comparison Is Not Possible

No published work applies graph reduction to whole AIGs as input reduction for a learned synthesis predictor, which is the gap documented in 2.2.6. There is therefore no matched number for RQ2 through RQ4 to be compared against. The reduction literature reports retention on citation and social graphs under node classification, a setting whose accuracy figures do not transfer to graph-level regression over deep directed acyclic graphs. Comparison is consequently structural rather than numerical: which family behaves as its own literature predicts, and which does not.

5.2 Interpretation of Findings

5.2.1 Feasibility of Optimizability Prediction

A random split is expected to score higher than the design-disjoint protocol used throughout, for a mechanical reason rather than a substantive one. The corpus is built by transforming and repeatedly optimizing a small number of source designs, so assigning graphs independently places a base graph in training and its own optimized descendants in test (3.4.2). The model can then recover the label by recognising which design it is looking at. That route is closed at deployment, where the circuit is new, so a higher random-split score is leakage. It is not evidence that AIGs derived from one design are interchangeable.

The stronger reading of the same outcome, that the graphs are simply similar enough for the split to make no difference, is harder to dismiss, and it cannot be settled by structural distance alone. Neither direction of the implication between structure and function holds for AIGs. Two AIGs computing the same Boolean function can differ arbitrarily in structure, which is exactly what a synthesis algorithm does to one, and two structurally close AIGs can compute different functions, since a single inverted edge changes the output. No structural metric therefore establishes that a test graph is or is not a training circuit in a weaker disguise, and “similar enough” has no structural definition to be proved against.

Two measurements settle the question in the form it can be settled. The first is a collision rate between partitions: the fraction of test graphs whose structural hash, or whose colour-refinement signature at the encoder’s depth, matches a training graph exactly. Near zero under the design-disjoint protocol and materially above zero under the random one would locate the inflation in identifiable duplicate structure rather than in diffuse similarity, which is a claim about specific graphs and is checkable. The second is the per-design breakdown already specified under RQ1a (4.1.6), which separates a uniform error across held-out designs from one carried by a few. Neither has been run.

5.2.2 What Reduction Buys and What It Costs

5.2.3 Why Domain-Informed Adaptation Did or Did Not Help

5.2.4 Cross-State Transfer and Deployment

5.2.5 Unexpected Results

The widest effect in the results is not a difference between reduction methods. It is the spread across source designs. Per-design R^2 on the design-disjoint test set runs from positive to strongly negative (4.1.5), and grouping the corpus by the script that produced each graph accounts for a quarter of the label variance on its own (4.1). Two things follow. Optimizability under a fixed script is heterogeneous enough across designs that a pooled score describes the mix of designs in the test set as much as it describes the model. Any comparison between configurations that reports one pooled number consequently hides most of the behaviour, which is why the per-design and per-source breakdowns are the primary reading here rather than a robustness check appended to it.

That heterogeneity is also a partial explanation for why synthesis practice is empirical. If the response of a circuit to a fixed optimization script varies this much between designs, running the script is a rational way to find out what it does, and the value of a predictor lies in shortening that search rather than in replacing it.

What the spread does not support is the claim that AIGs are inherently unpredictable. Each configuration was trained once (3.4.3), so no run-to-run variance estimate exists, and genuine heterogeneity between designs cannot be separated from training noise on this evidence. Validation R^2 moves by several units between consecutive epochs on some configurations, which is the magnitude the between-design differences would have to be judged against. Seeds are what would turn the observation into a claim.

5.3 Practical Implications

5.4 Limitations

The limitations below were imposed by the design or discovered during the work. The choices made in advance are stated separately in 1.2.6 and are not repeated here.

5.4.1 Reproducibility

Every stochastic component of training is seeded, so an individual run repeats exactly. That is a weaker guarantee than it appears. Each configuration was trained once, so no run-to-run variance is available and every accuracy difference reported in Chapter 4 is a point estimate. This bears hardest on RQ4, where the expected effect is small: a gap of a few percentage points between a domain-informed method and its control cannot be distinguished from seed noise on the evidence collected here. The efficiency comparisons are better placed, being paired per

graph with confidence intervals and a signed-rank test, but accuracy has no equivalent treatment.

Three further dependencies limit exact reproduction. Timing and memory measurements are specific to the cluster hardware and to the enabled compilation and mixed-precision settings, so they should be read as relative comparisons rather than as portable absolutes. The labels themselves depend on the version of the synthesis tool used to generate them; a different version would produce a different corpus, and the version was not pinned. The corpus is also not uniformly regenerable. One of the four generation scripts is stochastic, and its seed was drawn per graph and never recorded (A.1), so the quarter of the tier-1 set it produced could be reconstructed only in distribution, not exactly. Those graphs are training inputs rather than label sources, so the targets this thesis reports are unaffected, but the corpus itself is reproducible only in part.

5.4.2 Scalability

This is the limitation closest to the motivation and it should be stated without hedging. The argument for this work rests on industrial circuits of millions to billions of nodes. The evaluation runs on a corpus bounded at roughly a third of a million nodes per graph, because RQ1 requires a full-graph baseline and a baseline that does not fit in memory cannot be measured. The design that makes the study interpretable is the same design that keeps it away from the scale it argues about.

What does extrapolate: the relative ordering of reduction methods by offline cost, and the qualitative shape of the memory and throughput trade-offs, since these follow from properties of the algorithms and of message passing rather than from the corpus size. What does not: the absolute compression figures, since several methods' compression is a property of circuit structure rather than a parameter, and larger industrial circuits may have different redundancy profiles; and any claim that a method which retains accuracy here would retain it at a scale where the model sees a proportionally smaller fraction of each graph.

5.4.3 Generalizability

Three axes are held fixed, each narrowing what the conclusions cover.

One synthesis algorithm. Results characterise optimizability under Orchestrate. Graphs optimized by the other three scripts exist and are unused. If the reduction rankings hold across all four, the finding is about AIG structure; if they do not, it is about Orchestrate, and nothing here can tell the two apart.

One encoder. A four-layer edge-aware GCN+ was used throughout. Depth is not incidental to these results: methods were coupled to it by design, with refinement depth set equal to the number of layers. A deeper model would merge less under colour refinement and would have a different receptive-field profile, so the ranking could shift.

One target. Node reduction is the only quantity modelled. Delay and power are the objectives synthesis actually trades against, and nothing here speaks to them.

5.4.4 Reliability and Validity

Construct validity. Node reduction is a proxy for area, and area is one of three objectives. A circuit whose node count barely moves may still have been substantially improved in delay. The label measures what the tool removed, not how much better the result is, and “optimizability” should be read in that narrower sense throughout.

Internal validity. Splitting by design was chosen to prevent a base graph and its own optimized descendants from straddling the train/test boundary. Whether it fully succeeds is testable rather than assumed, and RQ1a is that test: the gap between the design-disjoint protocol and a random split measures directly how much of the score is design recognition rather than structural inference. Only the design-disjoint run exists, so that gap is currently unmeasured and no inflation factor can be quoted anywhere in this thesis (5.2.1). Should the per-design degradation prove concentrated in a few designs rather than uniform, a second design fold separates design difficulty from model error, at the cost of one further training run.

Construct validity of the comparison itself. Hyperparameters were tuned once on the full-graph configuration and held fixed. This protects the comparison from confounding retention with tuning effort, but it means reduced configurations run at settings chosen for a different input distribution. That biases against them, and therefore bounds how strongly any negative result about reduction can be stated.

5.4.5 Limitations of the Reduction Methods Tested

The reduction methods are not a representative sample of their literatures, and the summarization family in particular was scoped down deliberately: it runs one domain-specific method, one general GNN-aware method and one exact method, against a random control. The classical, guarantee-bearing spectral branch of the coarsening literature is not among them, so this thesis reports no trained result for it.

The domain-informed adaptations are simple. Edge-span weighting, level slicing, gate-role filtering and reconvergence-bounded merging are direct encodings of structural intuitions, not sophisticated methods. If they fail to beat their generic counterparts, that bounds these heuristics, not the idea of domain-aware reduction. That is a distinction 4.4.4 carries, and it should not be quietly dropped when the conclusion is written.

Finally, the boundary between structural compression and functional optimization is argued rather than measured. The claim that merging structurally equivalent nodes does not leak the label while merging functionally equivalent nodes would be well founded, but the negative control that would demonstrate it empirically, running FRAIG-style functional reduction [Mishchenko et al., 2005] as a summarizer and observing the label leak, was not run. The super-node content ablation is in the same position: member statistics are computed but not yet consumed by the encoder, so the finding in the literature that super-node content matters is cited rather than tested here.

5.5 Outlook

5.6 Ethical Considerations

Compute and energy cost. The work has a real environmental footprint: generating a corpus of several million AIGs required sustained CPU cluster time, and the training sweeps required substantial GPU time on top. The tension is worth stating rather than eliding: a study whose purpose is to reduce the compute cost of training spent a large amount of compute establishing how. The defence is that the cost is incurred once and the artifacts are stored and reusable, so a result showing which reductions preserve the signal saves compute for every subsequent user; but that is a claim about future savings against a certain present cost.

Benchmark provenance and licensing. The corpus is derived from published benchmark circuits, and their licences govern what may be redistributed. Every non-synthetic netlist was received through the OpenABC-D distribution, so the copies used are governed by its BSD 3-Clause terms, with the EPFL suite’s MIT license applying in addition (3.3.1). A dataset release must therefore carry both attribution notices. The ISCAS-85 and MCNC circuits attach no formal terms; their inclusion rests on the same decades-long redistribution practice that OpenABC-D itself relies on.

Accessibility, and who benefits. Lowering the memory required to train circuit models cuts in two directions. It widens access, since research that currently requires large multi-device clusters could run on modest hardware, which matters for groups without industrial compute budgets. It also lowers the cost of capability for those who already have that budget. The honest position is that this work makes an existing capability cheaper rather than creating a new one, and that the widening effect is the more likely of the two to matter at this scale.

Dual use. Logic synthesis tooling is general-purpose infrastructure. Improvements to it apply to whatever circuits are being designed, and this work does nothing specific to any application domain. No mitigations are proposed because none are specific enough to be meaningful; the consideration is recorded because omitting it would be a worse answer than naming it.

Conclusion

6.1 Answers to the Research Questions

Each answer below is stated with the limitation that bounds it, so that no claim travels further than the evidence behind it.

6.1.1 RQ1: Task Feasibility and Baseline

[TODO: state the measured baseline against the three baseline groups] Whatever accuracy the full-graph model reaches, it is accuracy on a proxy: node reduction stands in for area, and area is one of three objectives synthesis trades against (5.4.4). The figure is also conditional on the evaluation protocol, and the correction that would place it beside published numbers has not been measured (5.2.1).

6.1.2 RQ2: Reduction Efficiency

[TODO: state the efficiency ranking and the amortisation thresholds] The ranking is established on a corpus bounded at an intermediate scale. Relative ordering and the qualitative shape of the memory and throughput trade-offs extrapolate to larger circuits. Absolute compression does not, because several methods compress by an amount the circuit structure dictates rather than by a configured parameter (5.4.2).

6.1.3 RQ3: Predictive Retention and Trade-offs

[TODO: state which methods hold accuracy, at what compression] The retention result holds for one encoder at one depth, one synthesis algorithm and one target. Depth is not incidental: refinement depth was tied to the layer count by design, so a deeper model would merge less and could reorder the comparison (5.4.3).

6.1.4 RQ4: Value of Domain-Informed Adaptation

[TODO: state the paired gaps once the summarization pairings are trained] Each arm of each pairing is a single training run, so a gap of a few percentage points cannot be separated from seed noise (5.4.1). A null result bounds the four heuristics actually tested and not domain-aware reduction as an idea (4.4.4).

6.1.5 RQ5: Cross-State Transfer

[TODO: state the cross-state drop-off by reduction family] The transfer claim covers the methods trained here, which are not a representative sample of their literatures: the guarantee-bearing spectral branch of coarsening is absent, and the exact

track needs a conversion step before a full graph can be queried at all (5.4.5).

6.2 Contributions Revisited

6.3 Future Work

6.3.1 Extending Beyond a Single Synthesis Algorithm

The corpus already holds graphs optimized by the three synthesis scripts left unused here (1.2.6), so the extension costs training runs rather than another generation campaign. Two questions open at once. Whether the reduction ranking established under one script survives the other three decides whether the finding is about AIG structure or about that script, and nothing in this thesis can tell the two apart. Ranking *algorithms* for a given circuit, rather than circuits under a fixed algorithm, is the task the motivation names, and a single-algorithm model cannot perform it.

6.3.2 Stronger Domain-Aware Reduction

Three routes to a stronger AIG-aware method are open, in increasing order of ambition. The first is to consume super-node content. Member statistics are computed and never read by the encoder, and [Generale et al. \[2022\]](#) find that what a super-node contains is what makes coarsening work at all. The second is to give super-nodes a semantic identity by annotating each merged cone with its Negation-Permutation-Negation (NPN) equivalence class, which carries functional information about a cone without merging functionally equivalent nodes, and so stops short of the step that would leak the label. The third is to learn the reduction rather than specify it, which reintroduces precisely the label dependence that excluded condensation from this study (2.2.3).

6.3.3 Scaling to Industrial AIGs

The bound on graph size exists because RQ1 needs a full-graph baseline, and a baseline that does not fit in memory cannot be measured (1.2.6). Lifting the bound means giving up that reference point, so a successor study compares reductions against each other and needs a substitute control, such as a subcircuit small enough to admit an unreduced baseline of its own.

6.3.4 From Prediction to Script Generation

Cheap optimizability estimates are the input a script-selection agent consumes, not the selector itself (1.2.6). Closing that loop requires ranking algorithms for one circuit, so it depends on the

multi-algorithm extension above rather than standing independently of it.

Bibliography

- Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. Optuna: A next-generation hyperparameter optimization framework. In *Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2019.
- Uri Alon and Eran Yahav. On the bottleneck of graph neural networks and its practical implications. In *International Conference on Learning Representations (ICLR)*, 2021. Over-squashing. The mechanism behind the hypothesis that path-contracting coarsening can improve rather than degrade accuracy on deep graphs.
- Luca Amarú, Pierre-Emmanuel Gaillardon, and Giovanni De Micheli. The epfl combinational benchmark suite. In *Proceedings of the 24th International Workshop on Logic & Synthesis (IWLS)*, 2015. MIT licensed. Source of the eight arithmetic circuits in the corpus.
- Surender Baswana and Sandeep Sen. A simple and linear time randomized algorithm for computing sparse spanners in weighted graphs. *Random Structures & Algorithms*, 30(4):532–563, 2007. doi: 10.1002/rsa.20130.
- Per Bjesse and Arne Borålv. Dag-aware circuit compression for formal verification. In *Proceedings of the International Conference on Computer-Aided Design (ICCAD)*, pages 42–49. IEEE Computer Society / ACM, 2004. doi: 10.1109/ICCAD.2004.1382541. Origin of DAG-aware rewriting over precomputed equivalents.
- Jeroen Bollen, Jasper Steegmans, Jan Van den Bussche, and Stijn Vansummen. Learning graph neural networks using exact compression. In *Proceedings of the 6th Joint Workshop on Graph Data Management Experiences & Systems (GRADES) and Network Data Analytics (NDA)*, pages 8:1–8:9. ACM, 2023. doi: 10.1145/3594778.3594878. Collapses nodes in the same d -step colour-refinement class; provably identical GNN output. Graded variant cr_c with $c = \infty$ (colour refinement) and $c = 1$ (bisimulation).
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, Assya Trofimov, Brennan Nychporuk, Justin Szeto, Nazanin Mohammadi Sepahvand, Edward Raff, Kanika Madan, Vikram Voleti, Samira Ebrahimi Kahou, Vincent Michalski, Tal Arbel, Chris Pal, Gaël Varoquaux, and Pascal Vincent. Accounting for variance in machine learning benchmarks. In *Proceedings of the Conference on Machine Learning and Systems (MLSys)*, 2021.
- Robert K. Brayton and Alan Mishchenko. ABC: An academic industrial-strength verification tool. In *Computer Aided Verification (CAV)*, volume 6174 of *Lecture Notes in Computer Science*, pages 24–40. Springer, 2010. doi: 10.1007/978-3-642-14295-6_5. The citable record for the ABC system. The software itself is at <https://github.com/berkeley-abc/abc>.
- Franc Brglez and Hideo Fujiwara. A neutral netlist of 10 combinational benchmark circuits and a target translator in fortran. In *Proceedings of the IEEE International Symposium on Circuits and Systems (ISCAS)*, pages 677–692, 1985.
- Sergey Brin and Lawrence Page. The anatomy of a large-scale hypertextual web search engine. *Computer Networks*, 30(1-7): 107–117, 1998. doi: 10.1016/S0169-7552(98)00110-X.
- Davide Buffelli, Pietro Lió, and Fabio Vandin. Sizeshiftreg: a regularization method for improving size-generalization in graph neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. Trains so embeddings stay consistent across coarsening ratios; the size-shift framing behind RQ5.
- Jie Chen, Yousef Saad, and Zechen Zhang. Graph coarsening: From scientific computing to machine learning, 2022.
- Animesh Basak Chowdhury, Benjamin Tan, Ramesh Karri, and Siddharth Garg. Openabc-d: A large-scale dataset for machine learning guided integrated circuit synthesis. *arXiv preprint arXiv:2110.11292*, 2021. SynthNet. Graph-level labels including percentage of nodes optimized – the closest published analogue of the optimizability label used here.
- Animesh Basak Chowdhury, Shailja Thakur, Hammond Pearce, Ramesh Karri, and Siddharth Garg. Towards the imagenets of ml4eda. In *Proceedings of the IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, 2023. doi: 10.1109/ICCAD57390.2023.10323663. Invited paper. Position piece on benchmark and dataset standards for ML4EDA, motivating the framing used here.
- Giovanni De Micheli. *Synthesis and Optimization of Digital Circuits*. McGraw-Hill, 1994. Not indexed in DBLP; metadata taken from the publisher record.
- Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006.
- Chenhui Deng, Zichao Yue, Cunxi Yu, Gokce Sarar, Ryan Carey, Rajeev Jain, and Zhiru Zhang. Less is more: Hop-wise graph attention for scalable and generalizable learning on circuits. In *Proceedings of the Design Automation Conference (DAC)*, 2024. HOGA. Hop-wise attention over precomputed features; no message passing at training time.
- Charles Dickens, Eddie Huang, Aishwarya Reganti, Jiong Zhu, Karthik Subbian, and Lise Getoor. Convolution matching for efficient graph neural network training. In *Proceedings of the ACM Web Conference (WWW)*, 2024. ConvMatch / A-ConvMatch. Reference implementation: <https://github.com/amazon-science/convolution-matching>.
- Vijay Prakash Dwivedi, Ladislav Rampásek, Michael Galkin, Ali Parviz, Guy Wolf, Anh Tuan Luu, and Dominique Beaini.

- Long range graph benchmark. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2022.
- Isabella Venancia Gardner, Marcel Walter, Yukio Miyasaka, Robert Wille, and Michael Cochez. Bias by design: Diversity quantification to mitigate structural bias effects in AIG logic optimization. In *Design, Automation & Test in Europe Conference (DATE)*, pages 1–7. IEEE, 2025. doi: 10.23919/DATE64628.2025.10993009. Structural bias: an optimizer’s outcome is bounded by the structure of the AIG it is handed. Proposes AIG-specific similarity scores to quantify the diversity of a starting set.
- Alessandro Generale, Till Blume, and Michael Cochez. Scaling r-gen training with graph summarization, 2022. Trains on a summary graph and transfers weights back through a node-to-super-node mapping; ablation finds retaining super-node content critical.
- Martin Grohe, Kristian Kersting, Martin Mladenov, and Erkal Selman. Dimension reduction via colour refinement. In *Algorithms – ESA 2014, 22th Annual European Symposium*, volume 8737 of *Lecture Notes in Computer Science*, pages 505–516. Springer, 2014. doi: 10.1007/978-3-662-44777-2_42. Colour refinement computes the coarsest equitable partition – the formal basis for losslessness under a permutation-equivariant GNN.
- William L. Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems (NIPS)*, pages 1024–1034, 2017.
- Mohammad Hashemi, Shengbo Gong, Juntong Ni, Wenqi Fan, B. Aditya Prakash, and Wei Jin. A comprehensive survey on graph reduction: Sparsification, coarsening, and condensation. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence (IJCAI)*, pages 8058–8066, 2024.
- Abdelrahman Hosny, Soheil Hashemi, Mohamed Shalan, and Sherief Reda. Drills: Deep reinforcement learning for logic synthesis. In *Proceedings of the Asia and South Pacific Design Automation Conference (ASP-DAC)*, 2020. doi: 10.1109/ASP-DAC47756.2020.9045559. Representative learned-synthesis-script-generation work: an RL agent selects among the same primitive transformations at each step.
- Zengfeng Huang, Shengzhong Zhang, Chong Xi, Tang Liu, and Min Zhou. Scaling up graph neural networks via graph coarsening. In *Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2021. SCAL. Trains on the coarsened graph and infers directly with the same weights – the shared-weights precedent for RQ5.
- Wenjing Jiang, Jin Yan, and Sachin S. Sapatnekar. ML-based aig timing prediction to enhance logic optimization. In *Design, Automation and Test in Europe Conference (DATE)*, 2025. doi: 10.23919/DATE64628.2025.10993071. XGBoost AIG-timing predictor; one of the four surveyed papers reporting matched seen/unseen-design results.
- Wei Jin, Lingxiao Zhao, Shichang Zhang, Yozen Liu, Jiliang Tang, and Neil Shah. Graph condensation for graph neural networks. In *International Conference on Learning Representations (ICLR)*, 2022. GCond. Excluded here: no node correspondence, label-dependent, architecture-tied.
- Raika Karimi, Faezeh Faez, Yingxue Zhang, Xing Li, Lei Chen, Mingxuan Yuan, and Mahdi Biparva. Logic synthesis optimization with predictive self-supervision via causal transformers. *arXiv preprint arXiv:2409.10653*, 2024. doi: 10.48550/arXiv.2409.10653. LSOformer. One of the four surveyed papers reporting matched seen/unseen-design QoR results.
- George Karypis and Vipin Kumar. A fast and high quality multilevel scheme for partitioning irregular graphs. *SIAM Journal on Scientific Computing*, 20(1):359–392, 1998. doi: 10.1137/S1064827595287997.
- Barsat Khadka, Kawsher A. Roxy, and Md Rubel Ahmed. Cts-bench: Benchmarking graph coarsening trade-offs for gnns in clock tree synthesis. *arXiv preprint arXiv:2602.19330*, 2026. doi: 10.48550/arXiv.2602.19330. Nearest competitor. Post-placement gate-level netlists, clock-skew prediction; one bespoke spatial clustering heuristic. 17.2x memory / 3x speed at negative zero-shot R^2 ; calls for domain-aware coarsening. DBLP: journals/corr/abs-2602-19330, preprint only, no published venue as of verification.
- Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- Joseph B. Kruskal. On the shortest spanning subtree of a graph and the traveling salesman problem. *Proceedings of the American Mathematical Society*, 7(1):48–50, 1956. doi: 10.1090/S0002-9939-1956-0078686-7. Not indexed in DBLP; metadata taken from the publisher record.
- Andreas Kuehlmann, Viresh Paruthi, Florian Krohm, and Malay K. Ganai. Robust boolean reasoning for equivalence checking and functional property verification. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 21(12):1377–1394, 2002. doi: 10.1109/TCAD.2002.804386. Structural hashing (one-level lookup on construction) for AIGs.
- Amy Nicole Langville and Carl Dean Meyer. Deeper inside pagerank. *Internet Mathematics*, 1(3):335–380, 2003. doi: 10.1080/15427951.2004.10129091.
- Thomas Lengauer and Robert Endre Tarjan. A fast algorithm for finding dominators in a flowgraph. *ACM Transactions on Programming Languages and Systems*, 1(1):121–141, 1979. doi: 10.1145/357062.357071.
- Min Li, Sadaf Khan, Zhengyuan Shi, Naixing Wang, Huang Yu, and Qiang Xu. Deepgate: Learning neural representations of logic gates. In *Proceedings of the Design Automation Conference (DAC)*, 2022.
- Yingjie Li, Mingju Liu, Haoxing Ren, Alan Mishchenko, and Cunxi Yu. Dag-aware synthesis orchestration. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 43(12):4666–4675, 2024. doi: 10.1109/TCAD.2024.

3397052. The orchestration policy behind the target algorithm: per-node selection among rewrite, refactor and resubstitution instead of a fixed command sequence.
- Jiawei Liu, Jianwang Zhai, Mingyu Zhao, Zhe Lin, Bei Yu, and Chuan Shi. Polargate: Breaking the functionality representation bottleneck of and-inverter graph neural network. In *Proceedings of the International Conference on Computer-Aided Design (ICCAD)*, 2024. doi: 10.1145/3676536.3676834. Polarity/functionality bottleneck in AIG GNNs – supports treating inverter polarity as a first-class relation.
- Andreas Loukas. Graph reduction with spectral and cut guarantees. *Journal of Machine Learning Research*, 20:116:1–116:42, 2019. Local Variation coarsening; restricted spectral approximation guarantee.
- Yuankai Luo, Lei Shi, and Xiao-Ming Wu. Can classic gnns be strong baselines for graph-level tasks? simple architectures meet excellence. In *Proceedings of the International Conference on Machine Learning (ICML)*, volume 267 of *PMLR*, 2025. GNN+ – the encoder this work builds on. arXiv:2502.09263. Implementation: <https://github.com/LUOyk1999/GNNPlus>.
- Alan Mishchenko and Robert K. Brayton. Scalable logic synthesis using a simple circuit structure. In *Proceedings of the 15th International Workshop on Logic & Synthesis (IWLS)*, pages 15–22, 2006. Not indexed in DBLP; metadata taken from the paper. Source of the resubstitution command and the compress2rs script that C2RS aliases.
- Alan Mishchenko, Satrajit Chatterjee, Roland Jiang, and Robert K. Brayton. Fraigs: A unifying representation for logic synthesis and verification, 2005. ERL Technical Report, UC Berkeley. Functionally-reduced AIGs via SAT sweeping – the functional end of the equivalence spectrum, used here only as a leakage control.
- Alan Mishchenko, Satrajit Chatterjee, and Robert K. Brayton. Dag-aware aig rewriting: A fresh look at combinational logic synthesis. In *Proceedings of the 43rd Design Automation Conference (DAC)*, pages 532–535. ACM, 2006. doi: 10.1145/1146909.1147048. Defines the rewrite/refactor primitives and the MFFC as the unit refactoring operates on – the basis for the S6 domain argument.
- Liwei Ni, Rui Wang, Miao Liu, Xingyu Meng, Xiaozhe Lin, Junfeng Liu, Guojie Luo, Zhufei Chu, Weikang Qian, Xiaoyan Yang, Biwei Xie, Xingquan Li, and Huawei Li. Openlsdgf: An adaptive open-source dataset generation framework for machine-learning tasks in logic synthesis. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 44(10):3830–3843, 2025. doi: 10.1109/TCAD.2025.3555506. Adaptive open-source dataset generation framework for ML4EDA, positioned against the dataset built here.
- David Peleg and Alejandro A. Schäffer. Graph spanners. *Journal of Graph Theory*, 13(1):99–116, 1989. doi: 10.1002/jgt.3190130114.
- Ladislav Rampásek, Michael Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, and Dominique Beaini. Recipe for a general, powerful, scalable graph transformer. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2205.12454.
- Nils Reimers and Iryna Gurevych. Reporting score distributions makes a difference: Performance study of LSTM-networks for sequence tagging. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 338–348, 2017.
- Yu Rong, Wenbing Huang, Tingyang Xu, and Junzhou Huang. Droppedge: Towards deep graph convolutional networks on node classification. In *International Conference on Learning Representations (ICLR)*, 2020.
- Nasrin Shabani, Jia Wu, Amin Beheshti, Quan Z. Sheng, Jin Foo, Venus Haghighi, Ambreen Hanif, and Maryam Shahabikargar. A comprehensive survey on graph summarization with graph neural networks. *arXiv preprint*, 2023.
- Zhengyuan Shi and TODO. Deepgate3: Towards scalable circuit representation learning, 2024. Pooling transformer over subcircuits. Scales the model, not the input.
- Mathias Soeken, Heinz Riener, Winston Haaswijk, Eleonora Testa, Bruno Schmitt, Giulia Meuli, Fereshte Mozafari, Siang-Yun Lee, Alessandro Tempia Calvino, Dewmini Sudara Marakkalage, and Giovanni De Micheli. The epfl logic synthesis libraries, 2018. arXiv:1805.05121. mockturtle/kitty/lorina. Defines the random AIG sampling procedure behind the eight synthetic designs.
- Mahito Sugiyama and Kazuki Sato. Kron reduction for directed graphs, 2023.
- Veronika Thost and Jie Chen. Directed acyclic graph neural networks. In *International Conference on Learning Representations (ICLR)*, 2021.
- Yuan Yuan Tian, Richard A. Hankins, and Jignesh M. Patel. Efficient aggregation for graph summarization. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2008. k-SNAP. Subsumed here by graded colour refinement at $d = 1$; cited, not run.
- TODO. Rethinking reinforcement learning based logic synthesis. *arXiv preprint arXiv:2207.11437*, 2022. Joint Transformer (recipe) + GraphSAGE (circuit) QoR prediction.
- TODO. Universal graph coarsening, 2024. LSH-based, linear-time coarsening. AH-UGC is the 2025 consistent-hashing follow-up.
- Marcel Walter. aigverse: A python library for working with logic networks, synthesis, and optimization, 2025. MIT licensed. <https://github.com/marcelwa/aigverse>. Used by the ML pipeline to read and traverse AIGER files; pyproject.toml pins $\geq 0.0.27$.
- Nan Wu, Jiwon Lee, Yuan Xie, and Cong Hao. Lostin: Logic optimization via spatio-temporal information with hybrid graph models, 2022. Encodes the synthesis recipe as a super-node – a temporal use of super-nodes, orthogonal to the structural super-nodes here.

- Keyulu Xu, Chengtao Li, Yonglong Tian, Tomohiro Sonobe, Ken-ichi Kawarabayashi, and Stefanie Jegelka. Representation learning on graphs with jumping knowledge networks. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2018.
- Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? In *International Conference on Learning Representations (ICLR)*, 2019.
- Saeyang Yang. Logic synthesis and optimization benchmarks user guide: Version 3.0. Technical report, Microelectronics Center of North Carolina (MCNC), 1991.
- Muhan Zhang, Shali Jiang, Zhicheng Cui, Roman Garnett, and Yixin Chen. D-vae: A variational autoencoder for directed acyclic graphs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- Da Zheng, Chao Ma, Minjie Wang, Jinjing Zhou, Qidong Su, Xiang Song, Quan Gan, Zheng Zhang, and George Karypis. Distdgl: Distributed graph neural network training for billion-scale graphs. In *IEEE/ACM Workshop on Irregular Applications: Architectures and Algorithms (IA3)*, 2020. doi: 10.1109/IA351965.2020.00011. Graph partitioning across devices for distributed GNN training – the setting in which cutting edges is forced rather than chosen.
- Ziyang Zheng, Shan Huang, Jianyuan Zhong, Zhengyuan Shi, Guohao Dai, Ningyi Xu, and Qiang Xu. Deepgate4: Efficient and effective representation learning for circuit design at scale. In *International Conference on Learning Representations (ICLR)*. OpenReview.net, 2025. URL <https://openreview.net/forum?id=b10lRabU9W>. Sparse graph attention transformer, sub-linear memory, virtual edges over k -hop cones.

APPENDIX A

First Appendix

A.1 Synthesis Algorithm Configurations

This section records the concrete invocations behind 3.3.2 and 2.1.3. Every entry below is transcribed from the generation pipeline itself rather than from its documentation; where the two disagree, the pipeline is authoritative and the disagreement is noted.

Transformation recipes

Each recipe is 20 commands drawn independently and uniformly from `balance`, `rewrite`, `rewrite -z`, `refactor`, `refactor -z`, `resub` and `resub -z`, where `-z` enables zero-cost replacements. The 200 recipes used are the first 200 of the 1,500 released with OpenABC-D [Chowdhury et al., 2021] and are applied unmodified. The generator strips the input, output and hashing commands from each released recipe and appends a `write_aiger` after every surviving command, which is what makes each recipe prefix a stored graph.

Each released recipe continues past the 20 transformations into `dch` and then a technology-mapping tail, `map -B 0.9, topo, stime -c, buffer -c, upsize -c` and `dnsize -c`. The stripping rule does not remove that tail, so 27 write commands are emitted per recipe, but no standard-cell library is loaded and the mapped network cannot be written as an AIG, so only the 21 graphs up to and including `dch` are produced. The count assertion in the generation job, exactly 4,201 AIGs per design, is what establishes that the remaining six writes produced nothing.

Optimization invocations

Each algorithm is applied through the template in Table A.1, with the tool configuration file sourced first in every case.

Table A.1: Optimization invocations, one per algorithm, as issued by the generation pipeline.

Algorithm	Command
Orchestrate	<code>strash; orchestrate</code>
Deepsyn	<code>strash; &get; &deepsyn -S <seed> -T 20; &put</code>
Syn4	<code>strash; &get; &syn4; &put</code>
C2RS	<code>strash; c2rs</code>

Three of the four are invoked with no parameters at all, so their behaviour is the tool’s own default in each case. The pass budgets and scoring metrics named in the dataset documentation are not passed on the command line and were not set by this work. C2RS is an alias expanding to `b -1; rs -K 6 -1; rw -1; rs -K 6 -N 2 -1; rf -1; rs -K 8 -1; b -1; rs -K 8 -N 2 -1; rw -1; rs -K 10 -1; rwz -1; rs -K 10 -N 2 -1; b -1; rs -K 12 -1; rfz -1; rs -K 12 -N 2 -1; rwz -1; b -1`, which is the source of the widening cut size reported in 2.1.3.

Deepsyn is the one stochastic algorithm. Its seed is drawn from the shell random number generator once per input graph, is interpolated into the command, and is written neither to the run log nor to the metadata tables, so those graphs cannot be regenerated (5.4.1).

A.2 Reduction Method Parameters

A.3 Baseline Hyperparameters

A.4 Experimental Configuration

A.5 Extended Results